

AI Application Specialist

# 대규모 언어모델 이해 및 파인튜닝 (PART-2)

2025.8.18-19

사내교수 이재원 / 첨단기술아카데미

# Schedule

시간	1 일차	2 일차
08:30 ~ 09:20	<ul style="list-style-type: none"><li>과정 OT: 과정 소개 / AVD 소개</li></ul>	<ul style="list-style-type: none"><li>LLM Fine-Tuning: PEFT (LoRA, …)</li></ul>
09:30 ~ 10:20	<ul style="list-style-type: none"><li>[개요] Tokenizer</li></ul>	<ul style="list-style-type: none"><li>[실습-4] PEFT/LoRA (40m)</li></ul>
10:30 ~ 11:20	<ul style="list-style-type: none"><li>[실습-1] Text Tokenizer (30m)</li></ul>	<ul style="list-style-type: none"><li>LLM Fine-Tuning: Preference Tuning (PPO, DPO)</li></ul>
11:30 ~ 12:20	<ul style="list-style-type: none"><li>[개요] Word Embedding / Language Modeling</li></ul>	<ul style="list-style-type: none"><li>[실습-5] Preference Tuning - DPO (40m)</li></ul>
12:30 ~ 13:30		중식
13:30 ~ 14:20	<ul style="list-style-type: none"><li>Transformers Library / DataSet</li></ul>	<ul style="list-style-type: none"><li>LLM Evaluation</li></ul>
14:30 ~ 15:20	<ul style="list-style-type: none"><li>[실습-2] Text Generation API (30m)</li></ul>	<ul style="list-style-type: none"><li>[실습-6] LLM Evaluation (30m)</li></ul>
15:30 ~ 16:20	<ul style="list-style-type: none"><li>LLM Fine-Tuning: SFT</li></ul>	<ul style="list-style-type: none"><li>Transformer Architecture 및 최근 동향</li></ul>
16:30 ~ 17:20	<ul style="list-style-type: none"><li>[실습-3] LLM Fine-Tuning - SFT (40m)</li></ul>	<ul style="list-style-type: none"><li>과정 정리 및 QA</li></ul>

# CONTENTS

## Chapter 4 : Data Processing

- 01 ..... Data Sets
- 02 ..... Data Augmentation
- 03 ..... LLM Datasets

## Chapter 5 : LLM Fine-Tuning

## Chapter 6 : LLM Evaluation

## 4.1 Datasets

### ✓ 트랜스포머 모델 훈련 파이프라인



- Load a dataset (train/validation/test dataset)
- Preprocess the data (tokenizing/resampling/augmentation)

# 4.1 Datasets

## Hugging Face Hub

The screenshot shows the Hugging Face Hub's dataset search interface. At the top, there are filters for Tasks (Tasks 0), Sizes, Sub-tasks, Languages, Licenses, and Other. Below these are sections for Multimodal (Feature Extraction, Text-to-Image, Image-to-Text, Image-to-Video, Text-to-Video), Computer Vision (Depth Estimation, Image Classification, Object Detection, Image Segmentation, Image-to-Image, Unconditional Image Generation, Video Classification, Zero-Shot Image Classification, Mask Generation, Zero-Shot Object Detection), Natural Language Processing (Text Classification, Token Classification, Table Question Answering, Question Answering, Zero-Shot Classification, Translation, Summarization, Conversational, Text Generation, Text2Text Generation, Fill-Mask, Sentence Similarity, Table to Text, Multiple Choice, Text Retrieval), and Audio (Text-to-Speech, Text-to-Audio, Automatic Speech Recognition, Audio-to-Audio). The main search bar shows "Datasets 2,324" and a "Filter by name" input field. A "Full-text search" button and a "Sort: Trending" dropdown are also present. The results list includes datasets like "CohereForAI/ayea\_collection", "Open-Orca/OpenOrca", "nguha/legalbench", "glue", "allenai/prosocial-dialog", "sst2", "ag\_news", "dair-ai/emotion", "tweet\_eval", "yelp\_review\_full", "Hacker Noon/tech-company-news-data-dump", "ai4privacy/pii-masking-200k", "financial\_phrasebank", "go\_emotions", and "imdb". The "imdb" dataset is highlighted with a green border. On the right, a "Dataset Viewer" panel is open for the "imdb" dataset, showing a preview of the "facebook" split with 25k rows. It displays a sample row with columns "text" (containing a movie review) and "label" (with a blue bar indicating two classes: neg and pos). The viewer also includes a search bar, a sidebar with splits (train, val, test), and a footer with navigation links.

# 4.1 Datasets

## ✓ Dataset Loading

- `load_dataset()`

```
from datasets import load_dataset

emotions = load_dataset("emotion")

DatasetDict({
    train: Dataset({
        features: ['text', 'label'],
        num_rows: 16000
    })
    validation: Dataset({
        features: ['text', 'label'],
        num_rows: 2000
    })
    test: Dataset({
        features: ['text', 'label'],
        num_rows: 2000
    })
})
```

## ✓ Data Preprocessing

- `Dataset.map(), Dataset.filter(), ...`

```
from transformers import AutoTokenizer
from datasets import load_dataset

tokenizer = AutoTokenizer.from_pretrained(
    "bert-base-uncased")
dataset = load_dataset("emotion", split="train")

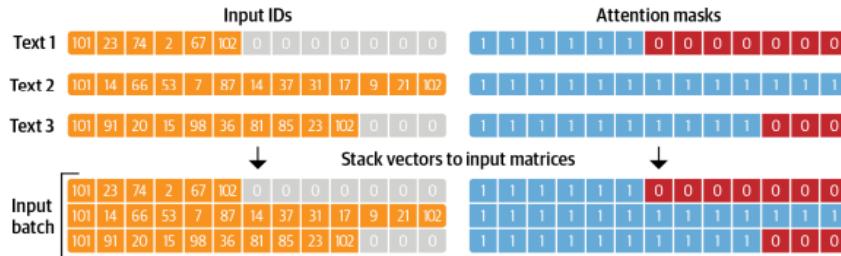
def tokenize(batch):
    return tokenizer(batch["text"])

dataset = dataset.map(tokenize, batched=True)
```

# 4.1 Datasets

## ✓ Data Preprocessing

- Padding, Truncation, Masking, ...

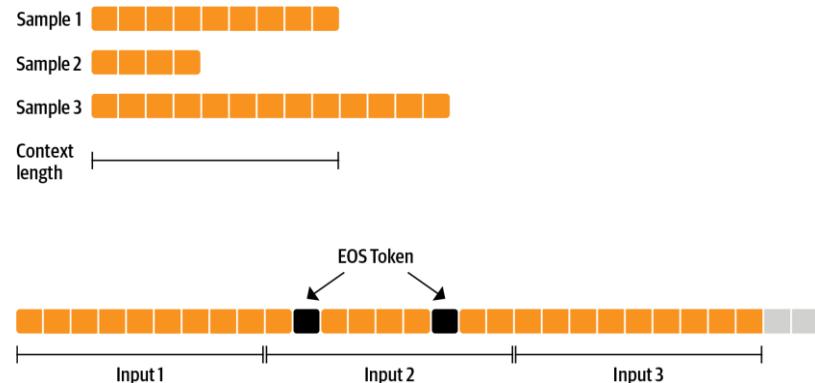


Special Token	[PAD]	[UNK]	[CLS]	[SEP]	[MASK]
Token ID	0	100	101	102	103

```
def tokenize(batch):  
    return tokenizer(batch['text'],  
                    padding=True,  
                    truncation=True)
```

## ✓ Data Loader

- Constant Length Dataset for Causal LM



- ✓ Context Size 를 경우 (ex. 1024) 효율성 고려
- ✓ Truncation에 따른 편향

# 4.1 Datasets

## ✓ Data Collator

- 데이터 전처리의 한 단계 (데이터 수집)
- 배치 구성할 때 필요한 작업을 자동으로 처리
- 샘플을 배치로 묶는 과정에서 사용
- Padding, Token Masking, Labeling, Mini Batch

(1, 3, 4)  
(5, 3, 10)

(1, 7, 5, 2, 2, 4)  
(8, 9, 3, 8, 2, 10)

(1, 3, 2, 4)  
(3, 4, 8, 10)

**INPUTS:**  
1, 3, 4, 0, 0, 0  
1, 7, 5, 2, 2, 4  
1, 3, 2, 4, 0, 0  
**LABELS:**  
5, 3, 10, -100, -100, -100  
8, 9, 3, 8, 2, 10  
3, 4, 8, 10, -100, -100



```
# DefaultDataCollator
data_collator = DefaultDataCollator()

# DataCollatorWithPadding
data_collator = DataCollatorWithPadding(
    tokenizer=tokenizer
)

# DataCollatorForLanguageModeling
data_collator = DataCollatorForLanguageModeling(
    tokenizer=tokenizer,
    mlm=True,
    mlm_probability=0.15
)

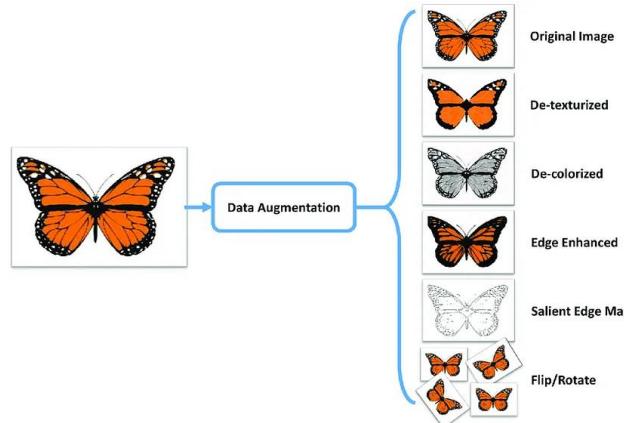
# Trainer
trainer = Trainer(
    model=model,
    args=training_args,
    train_dataset=dataset,
    data_collator=data_collator,
)
```

# 4.2 Data Augmentation

## ✓ Benefits of Data Augmentation

- Implicit Regularization
- Semi-supervised applications, insufficient data
- Cost effective way to data gathering and labeling

## ✓ Data augmentation is semantically invariant transformation



time flies like an arrow

↓ shuffle

an arrow like flies time?

# 4.2 Data Augmentation

## ✓ Lexical Substitution

- Thesaurus Based Substitution

: Look up **the synonyms** (ex. WordNet)

It is awesome → **WordNet** → It is amazing

- Word Embeddings Substitution

: **Nearest neighbors** in the embedding space

: Word2Vec, FastText, GloVe, etc.



- Masked Language Model

: **Predict masked words** based on the context

: BERT, ROBERTA, ALBERT, etc.

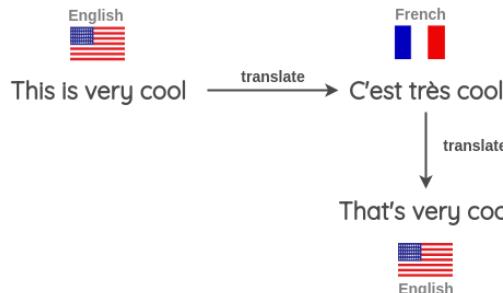
This is **pretty** cool  
This is **really** cool  
This is **super** cool  
This is **kinda** cool  
This is **very** cool

This is **[mask]** cool → **MLM** → This is **super** cool

# 4.2 Data Augmentation

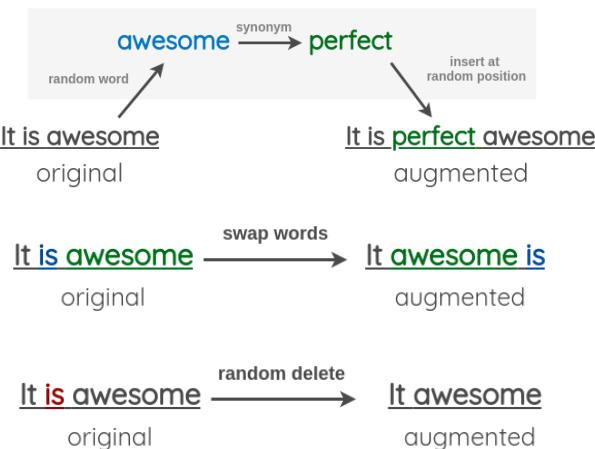
## ✓ Back Translation

- Take some sentence (e.g. in English) and translate to another Language e.g. French
- Translate the French sentence back into an English sentence (Back Translation)
- Check if the new sentence is different from the original sentence. If it is, then we use this new sentence as an augmented data



## ✓ Random Noise Injection

- Spelling Error Injection
- Keyboard Error Injection
- Unigram Noising
- Random Injection / Swap /Deletion



## 4.2 Data Augmentation

### Example: ContextualWordEmbsAug

```
from transformers import set_seed
import nlpAug.augmenter.word as naw

set_seed(3)
aug = naw.ContextualWordEmbsAug(model_path='distilbert-base-uncased',
                                 device="cpu",
                                 action="substitute")

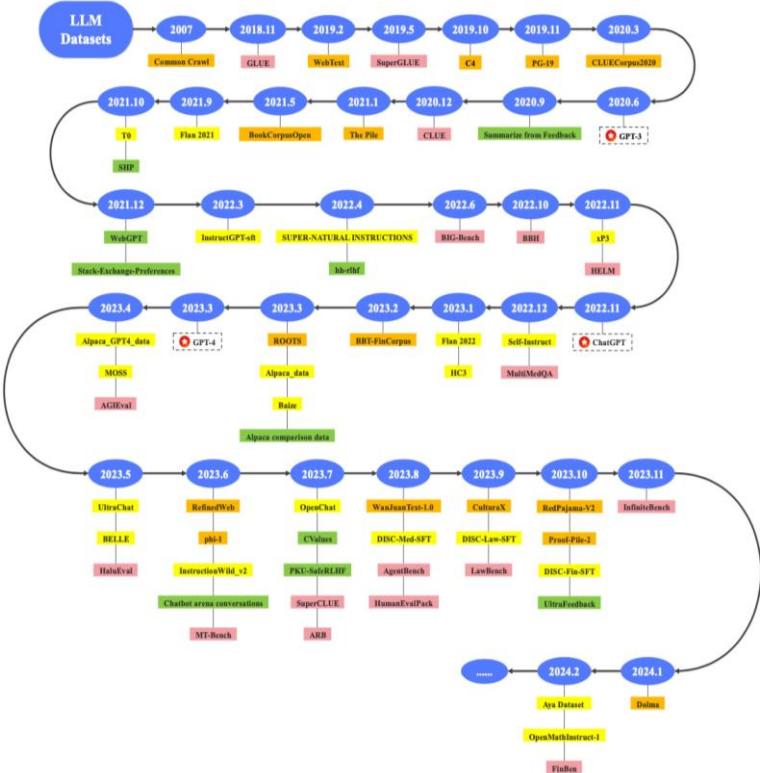
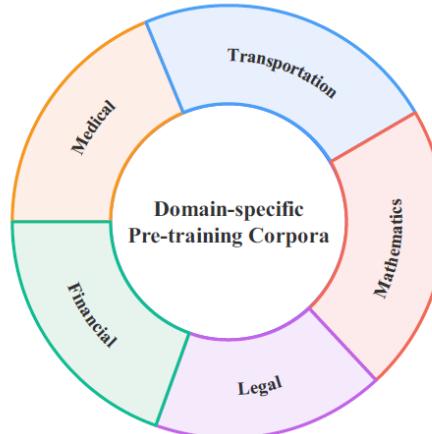
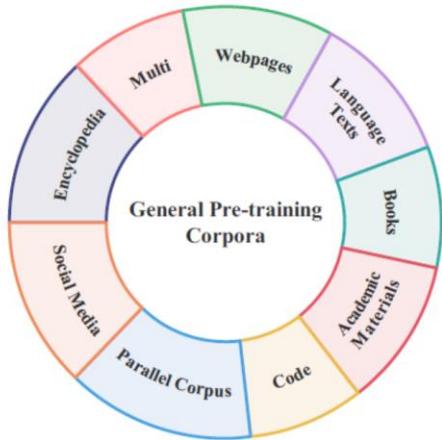
def augment_text(batch, transformations_per_example=1):
    text_aug, label_ids = [], []
    for text, labels in zip(batch['text'], batch['label_ids']):
        text_aug += [text]
        label_ids += [labels]
        for _ in range(transformations_per_example):
            text_aug += aug.augment(text)
            label_ids += [labels]
    return {"text": text_aug, "label_ids": label_ids}

ds_train_sample = ds_train_sample.map(augment_text, batched=True, ...)
```

## 4.3 LLM Datasets

### ✓ A timeline of some representative LLM datasets

- Pre-training corpora (orange)
- Instruction fine-tuning datasets (yellow)
- Preference (green)
- Evaluation datasets (pink)



## 4.3 LLM Datasets

### ✓ LLM 학습 주요 Data

- **Text Generation**  
: Common Crawl, WebText, WikiText, BookCorpus, etc.
- **Conversation**  
: MultiWOZ, ConvAI, Daily-Dialog, Persona-Chat, etc.
- **Question Answering**  
: SQuAD, NaturalQuestions, TriviaQA, etc.
- **Summarization**  
: CNN/Daily Mail, XSum, etc.
- **Paraphrasing**  
: PAWS, Quora Question Pairs, etc.
- **Code Generation**  
: CodeSearchNet, etc.

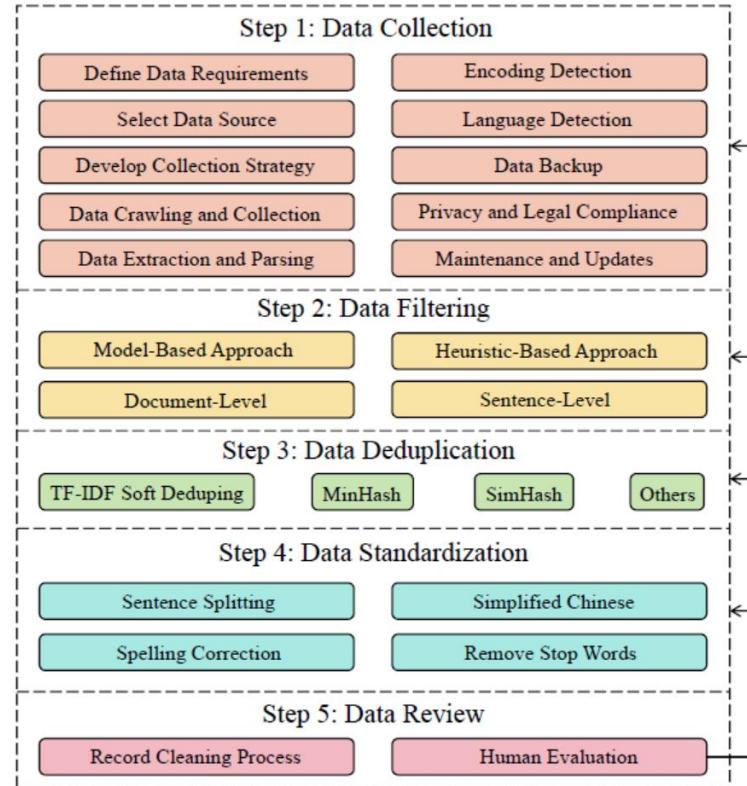
[Korean NLP Datasets](#)



# 4.3 LLM Datasets

## Preprocessing for pre-training corpora

- Data Collection
- Data Filtering
- Data Duplication
- Data Standardization
- Data Review



# 4.3 LLM Datasets: IF Template

## ✓ Meta Llama 3

System prompt and multiple turn conversation between the user and assistant

```
<|begin_of_text|><|start_header_id|>system<|end_header_id|>

{{ system_prompt }}<|eot_id|><|start_header_id|>user<|end_header_id|>

{{ user_message_1 }}<|eot_id|>
<|start_header_id|>assistant<|end_header_id|>

{{ model_answer_1 }}<|eot_id|><|start_header_id|>user<|end_header_id|>

{{ user_message_2 }}<|eot_id|>
<|start_header_id|>assistant<|end_header_id|>
```

```
<|start_header_id|>system<|end_header_id|>
You are a helpful AI assistant.<|eot_id|><|start_header_id|>user<|end_header_id|>
Hi!<|eot_id|><|start_header_id|>assistant<|end_header_id|>
Hello there! It's great to meet you!<|eot_id|><|start_header_id|>user<|end_header_id|>
What are some recipes for making chocolate chip cookies?<|eot_id|><|start_header_id|>assistant<|end_header_id|>
```

## ✓ Alpaca

We used the following prompts for fine-tuning the Alpaca model:

- for examples with a non-empty input field:

Below is an instruction that describes a task, paired with an input that provides further context. Write a

```
### Instruction:  
{instruction}  
  
### Input:  
{input}  
  
### Response:
```

- for examples with an empty input field:

Below is an instruction that describes a task. Write a response that appropriately completes the request.

```
### Instruction:  
{instruction}  
  
### Response:
```

During inference (eg for the web demo), we use the user instruction with an empty input field (second option).



- ☑ Dataset (Train/Validation/Test)
- ☑ Dataset Loading, Preprocessing
- ☑ Padding, Truncation, Masking
- ☑ Text Data Augmentation
- ☑ LLM Datasets, Instruction Following Format

# CONTENTS

**Chapter 4 : Data Processing**

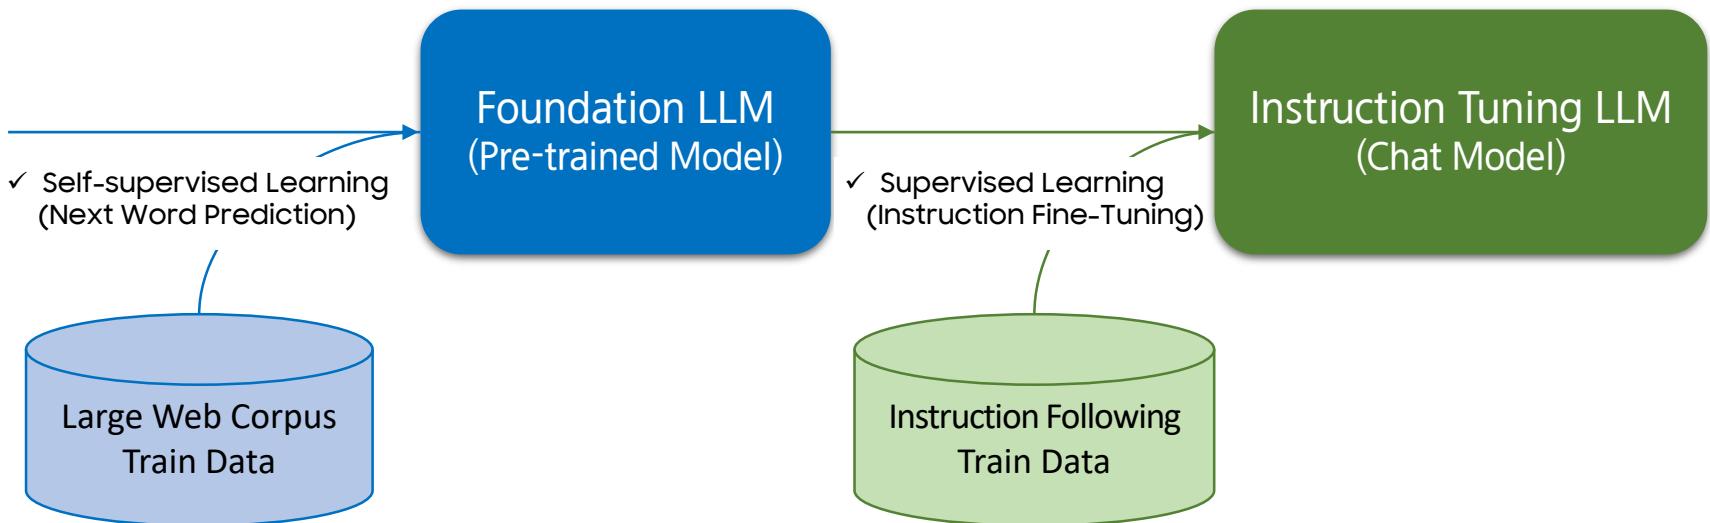
**Chapter 5 : LLM Fine-Tuning**

- 01 ..... Supervised Fine-Tuning
- 02 ..... Parameter Efficient FT
- 03 ..... Preference Tuning

**Chapter 6 : Evaluation**

# 5.0 LLM Training

- ✓ Foundation LLM (Pre-trained Model): 방대한 Generic Data 학습
- ✓ Instruction Fine-Tuning LLM: Prompt-Response Data 학습 Fine-Tuning



# 5.0 LLM Training

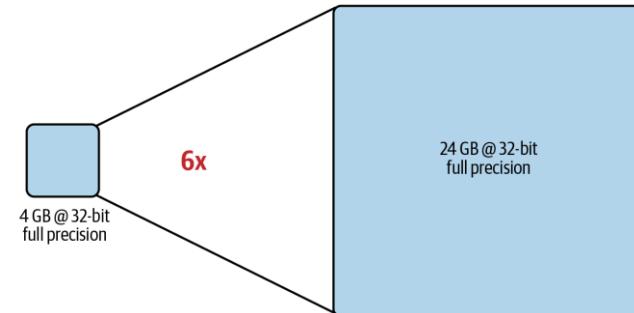
## ✓ 언어 모델 학습 (Training): Foundation Model Size, Training Data & Model Performance

- 학습에 사용 가능한 GPU 스펙/자원 수, 모델 사이즈/학습 데이터 양에 따라 다양
- 일반적으로 모델 사이즈가 클수록, 양질의 학습 데이터가 많을수록 모델 성능이 향상 (#Param x20 이상)
- 학습: ① Pre-Training (지식) → ② Instruction Fine-Tuning (지시어) → ③ PEFT (도메인) / RLHF (선후도)

Additional GPU RAM needed to train 1B parameters

	Bytes per parameter
Model Parameters (Weights)	4 bytes per parameter
Adam optimizer (2 states)	+8 bytes per parameter
Gradients	+4 bytes per parameter
Activations and temp memory (variable size)	+8 bytes per parameter (high-end estimate)
TOTAL	=4 bytes per parameter +20 extra bytes per parameter

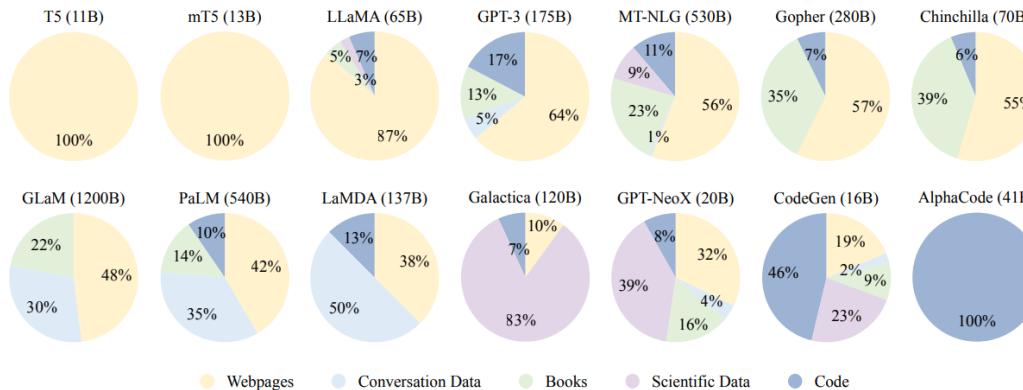
Load a 1-billion-parameter model



# 5.0 LLM Foundation Model: Pre-Training

## ① Foundation 모델 사전 학습 (Pre-Training, Self-Supervised Learning)

- 학습 데이터: 범용 공개 데이터 (라벨링 X) 및 지시어 데이터 (라벨링 ○)
- 데이터 구성: Webpages, Conversation Data, Books, Scientific Data, Code (범용 멀티 태스크 데이터)
- 데이터 전처리: 토큰화 (BPE 등)



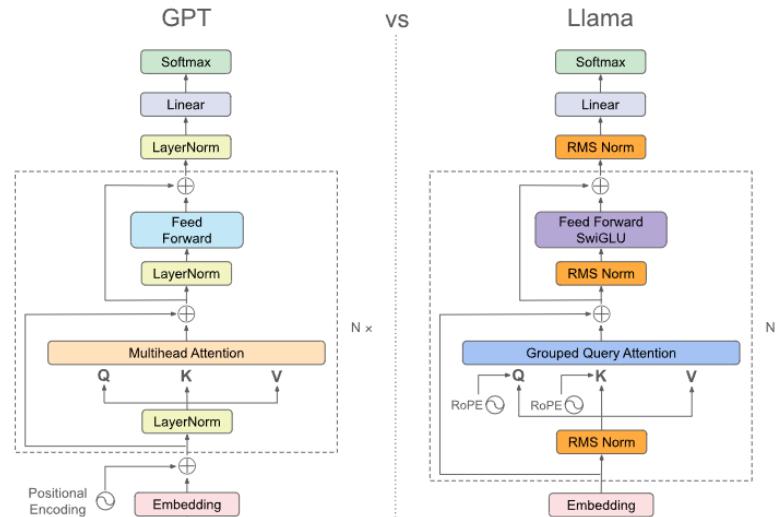
Token String	Token ID	Embedding / Vector Representation
'_The'	37	[-0.0513, -0.0584, 0.0230, ...]
'_teacher'	3145	[-0.0335, 0.0167, 0.0484, ...]
'_teaches'	11749	[-0.0151, -0.0516, 0.0309, ...]
'_the'	8	[-0.0498, -0.0428, 0.0275, ...]
'_student'	1236	[-0.0460, 0.0031, 0.0545, ...]
...	...	...

Vocabulary

# 5.0 [Open Source] Pretrained LLM: Llama

## ✓ Llama Architecture

- Decoder Architecture, RoPE, RMS Norm, SwiGLU Activation Function, Grouped Query Attention 



Llama2	Feature	Llama3
SentencePiece	Tokenizer	Tiktoken (OpenAI)
70B, 13B, 7B	Number of Parameters	8B, 70B, 405B
2.2T tokens	Training Data	<b>15T tokens</b>
4,096 tokens	Context Length	<b>8,192 tokens</b>
Grouped-query attention	Attention Mechanism	<b>Grouped-query attention</b>
Yes	Fine-Tuned Models	Yes
Better than Llama 1	Performance	Better than Llama 2
Very high	Computational Requirements	Very high
Open source	Availability	Open source
Yes	RLHF	Yes
20 languages	# of languages supported	<b>30 languages</b>

\* RoPE (Rotary Position Embedding): 각 토큰의 위치를 상대적인 각도로 표현, SwiGLU (Switch + GLU)

# 5.0 [Open Source] Pretrained LLM: Llama

## ✓ Meta Llama 3 (2024.4), Llama 3.2 (2024.9)

Llama-3 8B 모델 성능이  
Llama-2 70B 모델보다 우수

Meta Llama 3 Instruct model performance

	Meta Llama 3 8B	Gemma 7B - It Measured	Mistral 7B Instruct Measured
MMLU 5-shot	<b>68.4</b>	53.3	58.4
GPQA 0-shot	<b>34.2</b>	21.4	26.3
HumanEval 0-shot	<b>62.2</b>	30.5	36.6
GSM-8K 8-shot, CoT	<b>79.6</b>	30.6	39.9
MATH 4-shot, CoT	<b>30.0</b>	12.2	11.0

	Meta Llama 3 70B	Gemini Pro 1.5 Published	Claude 3 Sonnet Published
MMLU 5-shot	<b>82.0</b>	81.9	79.0
GPQA 0-shot	39.5	<b>41.5</b> CoT	38.5 CoT
HumanEval 0-shot	<b>81.7</b>	71.9	73.0
GSM-8K 8-shot, CoT	<b>93.0</b>	91.7 11-shot	92.3 0-shot
MATH 4-shot, CoT	50.4	<b>58.5</b> Minerva prompt	40.5

### [Llama 3]

- ✓ 2T Token → 15T Token
- ✓ 2K → 4K → 8K (Context)
- ✓ 8B, 70B, 405B Model
- ✓ Multi-lingual

### [Llama 3.2]

- ✓ On-Device: 1B/3B
- ✓ Multi-modal: 11B/90B

### [Llama 3.3]

- ✓ 70B (성능: 3.1 405B 수준)

# 5.0 [Open Source] Pretrained LLM: Llama

## ✓ Meta Llama 4 (2025.4)

- Scout (109B): 단일 GPU 경량 모델, Maverick (400B): GPT-4o 능가, Behemoth (2T): 초대형 모델 개발 중
- 멀티모달 기본 탑재 (텍스트, 이미지, 비디오), 전문가 혼합 (MoE) 아키텍처, 10M 토큰 컨텍스트

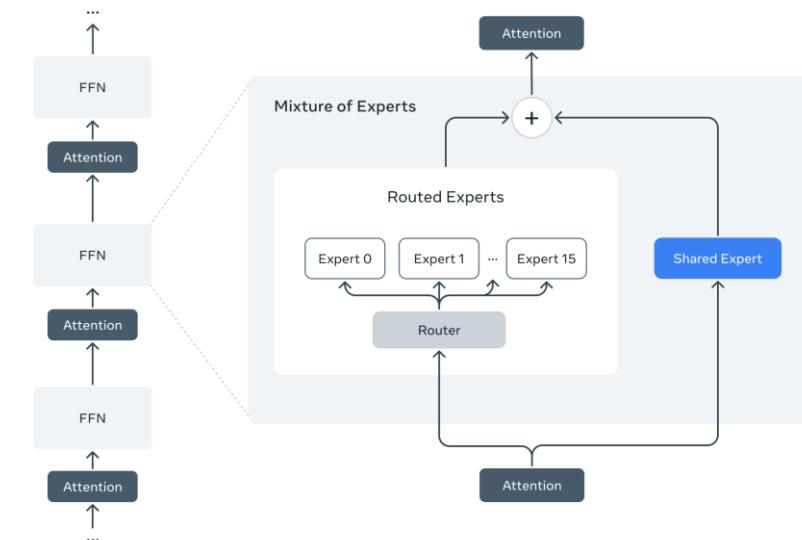
Llama 4:  
Leading Multimodal Intelligence

Newest model suite offering unrivaled speed and efficiency

**Llama 4 Behemoth**  
288B active parameters, 16 experts  
2T total parameters  
The most intelligent teacher model for distillation  
[Preview](#)

**Llama 4 Maverick**  
17B active parameters, 128 experts  
400B total parameters  
Native multimodal with 1M context length  
[Available](#)

**Llama 4 Scout**  
17B active parameters, 16 experts  
109B total parameters  
Industry leading 10M context length  
Optimized inference  
[Available](#)



# 5.0 [Open Source] Pretrained LLM: Gemma

## ✓ [google/gemma-2 \('24.6\)](#)

- Lightweight, state-of-the-art open models
- Text-to-text, decoder-only LLM (English)  
: RMSNorm, RoPE, GEGLU (Gated GELU), MQA
- Instruction following
- Multi-turn conversations

Hyperparameter	2B	9B	27B
$n_{layers}$	26	42	46
$d_{model}$	2,304	3,584	4,608
$d_{ff}$	18,432	28,672	73,728
$n_{heads}$	8	16	32
$n_{key\_value\_heads}$	4	8	16
$n_{ctx}$	8,192	8,192	8,192
$n_{vocab}$	256,128	256,128	256,128

모델 크기를 줄이면서도 성능을 유지할 수 있는 효율적 설계 지향

BENCHMARK	METRIC	Gemma 2		Llama 3		Grok-1	
		9B	27B	8B	70B	314B	
General	MMLU	5-shot, top-1	71.3	75.2	66.6	79.5	73.0
Reasoning	BBH	3-shot, CoT	68.2	74.9	61.1	81.3	-
	HellaSwag	10-shot	81.9	86.4	82	-	-
Math	GSM8K	5-shot, maj@1	68.6	74.0	45.7	-	62.9 (8-shot)
	MATH	4-shot	36.6	42.3	-	-	23.9
			40.2	51.8	-	-	63.2 (0-shot)

BENCHMARK	METRIC	Gemma 1.1 IT		Gemma 2 IT	
		2.6B	7B	9B	27B
RealToxicity	avg toxicity	7.03	8.04	8.25	8.84
CrowS-Pairs	top-1	45.89	49.67	37.47	36.67
BBQ Ambig	1-shot, top-1	58.97	86.06	88.58	85.99
BBQ Disambig	top-1	53.9	85.08	82.67	86.94
Winogender	top-1	50.14	57.64	79.17	77.22
TruthfulQA	MC2Acc	44.24	45.34	50.27	51.60
Winobias 1_2	top-1	55.93	59.22	78.09	81.94
Winobias 2_2	top-1	89.46	89.2	95.32	97.22
Toxigen	avg toxicity	29.64	38.75	39.30	38.42

# 5.0 [Open Source] Pretrained LLM: Gemma

## ✓ [google/gemma-3 \('25.3\)](#)

- Longer Context Length (8K → 32K, 128K)
- Multimodality (Image & Text)
- Multilinguality (English + 140 Languages)

Both pre-trained and instruction tuned models are released.

Gemma-3-4B-IT beats Gemma-2-27B IT, while Gemma-3-27B-IT beats Gemini 1.5-Pro across benchmarks.

	Gemma 2	Gemma 3
Model Size	<ul style="list-style-type: none"><li>• 2B</li><li>• 9B</li><li>• 27B</li></ul>	<ul style="list-style-type: none"><li>• 1B</li><li>• 4B</li><li>• 12B</li><li>• 27B</li></ul>
Context Window Length	<ul style="list-style-type: none"><li>• 8K</li></ul>	<ul style="list-style-type: none"><li>• 32K (1B)</li><li>• 128K (4B, 12B, 27B)</li></ul>
Multimodality (Image & Text)	<ul style="list-style-type: none"><li>• <input checked="" type="checkbox"/></li></ul>	<ul style="list-style-type: none"><li>• <input checked="" type="checkbox"/> (1B)</li><li>• <input checked="" type="checkbox"/> (4B, 12B, 27B)</li></ul>
Multilingual Support	-	<ul style="list-style-type: none"><li>• English (1B) + 140 Languages (4B, 12B, 27B)</li></ul>

Rank	Model	Elo	95% CI	Open	Type	#params/#activated
1	Grok-3-Preview-02-24	1412	+8/-10	-	-	-
1	GPT-4.5-Preview	1411	+11/-11	-	-	-
3	Gemini-2.0-Flash-Thinking-Exp-01-21	1384	+6/-5	-	-	-
3	Gemini-2.0-Pro-Exp-02-05	1380	+5/-6	-	-	-
3	ChatGPT-4o-latest (2025-01-29)	1377	+5/-4	-	-	-
6	DeepSeek-R1	1363	+8/-6	yes	MoE	671B/37B
6	Gemini-2.0-Flash-001	1357	+6/-5	-	-	-
7	o1-2024-12-17	1352	+4/-6	-	-	-
9	<b>Gemma-3-27B-IT</b>	<b>1338</b>	<b>+8/-9</b>	<b>yes</b>	<b>Dense</b>	<b>27B</b>
9	Qwen2.5-Max	1336	+7/-5	-	-	-
9	o1-preview	1335	+4/-3	-	-	-
9	o3-mini-high	1329	+8/-6	-	-	-
12	DeepSeek-V3	1318	+8/-6	yes	MoE	671B/37B
13	GLM-4-Plus-0111	1311	+8/-8	-	-	-
13	Qwen-Plus-0125	1310	+7/-5	-	-	-
13	Claude 3.7 Sonnet	1309	+9/-11	-	-	-
14	Gemini-2.0-Flash-Lite	1308	+5/-5	-	-	-
14	Step-2-16K-Exp	1305	+7/-6	-	-	-
14	o3-mini	1304	+5/-4	-	-	-
14	o1-mini	1304	+4/-3	-	-	-
14	Gemini-1.5-Pro-002	1302	+3/-3	-	-	-
...						
28	Meta-Llama-3.1-405B-Instruct-bf16	1269	+4/-3	yes	Dense	405B
...						
38	Llama-3.3-70B-Instruct	1257	+5/-3	yes	Dense	70B
...						
39	Qwen2.5-72B-Instruct	1257	+3/-3	yes	Dense	72B
...						
59	Gemma-2-27B-it	1220	+3/-2	yes	Dense	27B

Evaluation of Gemma 3 27B IT model in the Chatbot Arena (March 8, 2025)

# 5.0 [Open Source] GPT, Gemma, Llama

DeepSeek-R1



- ✓ Architecture: 모두 디코더 구조 사용
- ✓ Attention Mechanism: 연산량을 줄이는 방향으로 MHA, MQA, GQA 사용
- ✓ Activation: Dying ReLU 문제를 해결하는 방향으로 GEGLU, SwiGLU 사용

Feature	GPT	Gemma 2	Llama 3
Architecture	Decoder	Decoder	Decoder
Positional Encoding	Sinusoidal PE	Rotary PE	Rotary PE
Normalization	Layer Norm	RMS Norm	RMS Norm
Attention Mechanism	Multi-Head Attention	Local Sliding Window, Multi Query Attention (MQA)	Grouped Multi-Query Attention (GQA)
Activation	ReLU	GEGLU	SwiGLU
Training Data	4.5M+36M	13T (27B Model)	15T



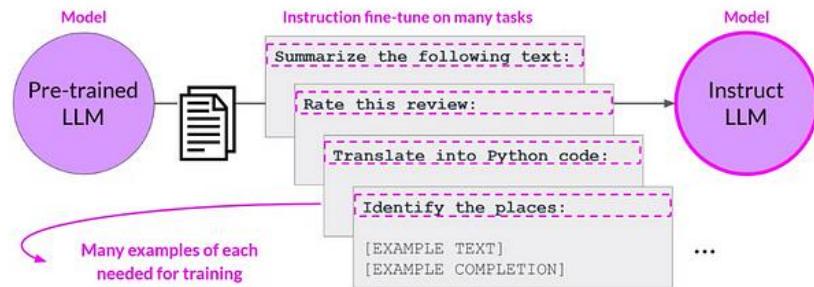
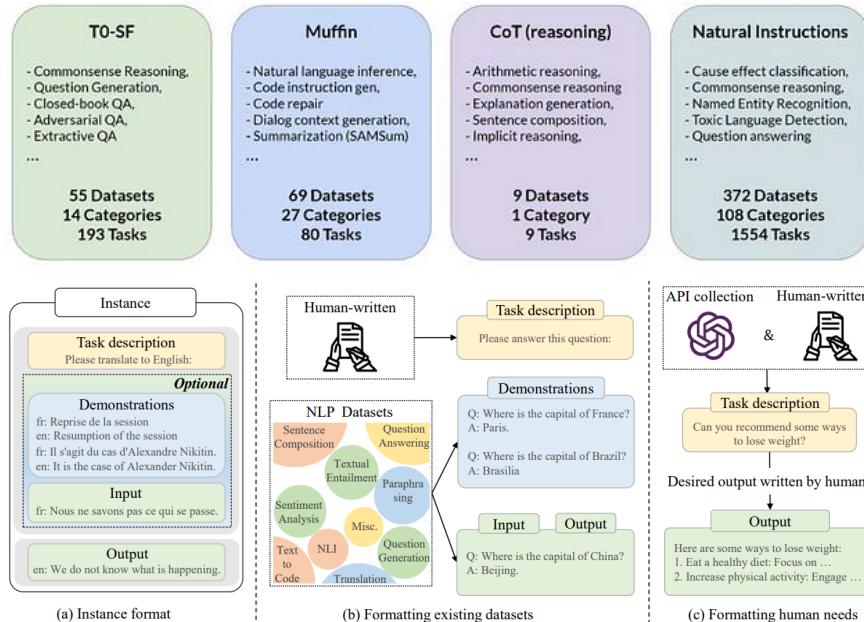
## ✓ 상용 LLM v.s. 오픈소스 LLM

- ⌚ 모델 크기 및 스펙 (Context Size, Max. Token, Computing Resource)
- ⌚ 모델 성능: 벤치마크 성능, 자연시간, 모델 안정성 등
- ⌚ 비용 및 호스팅: 모델 사용료, 모델 호스팅 비용 등
- ⌚ Model Customization: LLM Fine-Tuning 가능 여부

# 5.1 LLM Fine-Tuning: SFT

## ② 사전 학습 Foundation 모델 기반 Instruction Following 학습

- Instruction Following Data: 휴먼 구축, 기존 멀티 태스크 데이터 활용, LLM 활용 데이터 생성
- Multi-turn Data: 인위적으로 명시적 프롬프트 기반 데이터 생성 후 조합



# 5.1 LLM Fine-Tuning: SFT

## ✓ Instruction Fine-Tuning

In the loss function, mask out the tokens corresponding to prompt

Below is an instruction that describes a task. Write a response that appropriately completes the request.

### Instruction:  
Rewrite the following sentence using passive voice.

### Input:  
The team achieved great results.

### Response:  
**Great results were achieved by the team.**

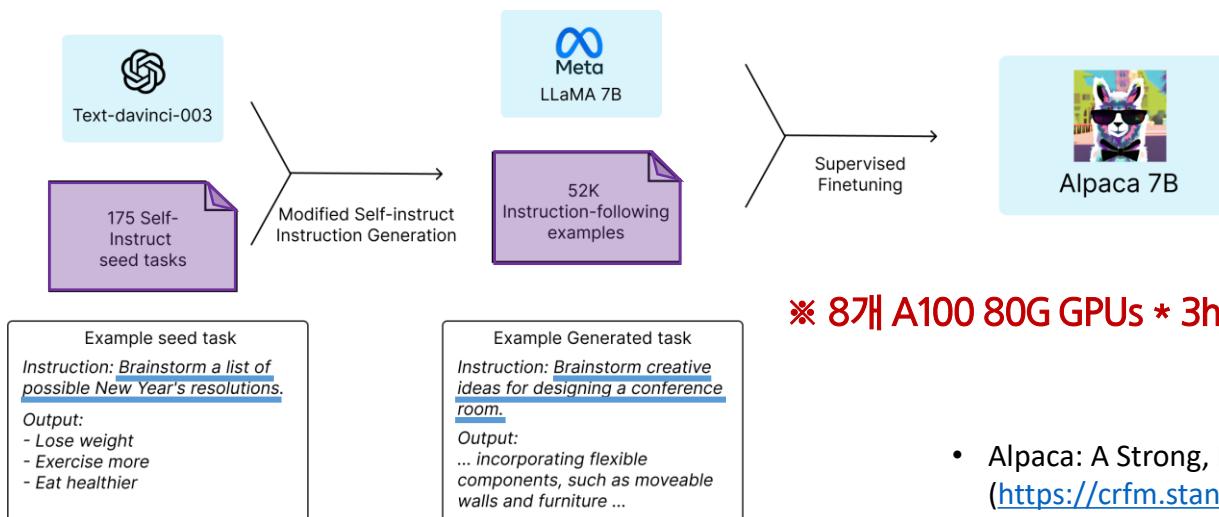
Calculate the loss only over the tokens corresponding to the response text



# 5.1 LLM Fine-Tuning: Stanford Alpaca

## ✓ An Instruction-following LLaMA Model

- Fine-tuned from a **7B LLaMA 1 model**
- On **52K instruction following data generated by the Self-Instruct**
- Similar to the **text-davinci-003** model on the **Self-Instruct instruction-following evaluation suite**



Stanford  
Alpaca



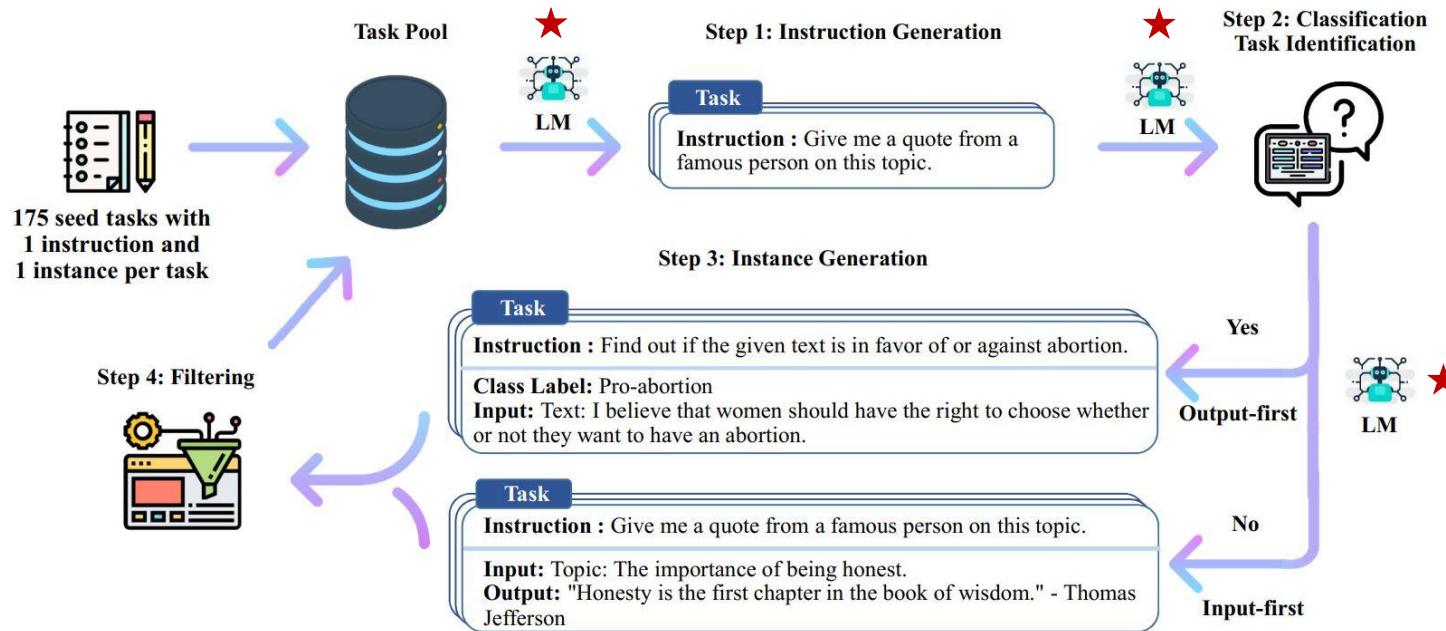
Hyperparameter	LLaMA-7B
Batch size	128
Learning rate	2e-5
Epochs	3
Max length	512
Weight decay	0

※ 8개 A100 80G GPUs \* 3hr

- Alpaca: A Strong, Replicable Instruction-Following Model (<https://crfm.stanford.edu/2023/03/13/alpaca.html>)

# 5.1 LLM Fine-Tuning: Stanford Alpaca

## ✓ Data Generation: High Level Overview of Self-Instruct



Ref. Self-Instruct: Aligning Language Model with Self Generated Instructions. Yizhong Wang, Yeganeh Kordi, Swaroop Mishra, Alisa Liu, Noah A. Smith, Daniel Khashabi, Hannaneh Hajishirzi. <https://arxiv.org/abs/2212.10560>

# 5.1 LLM Fine-Tuning: Stanford Alpaca

## ✓ Instruction-Following Data

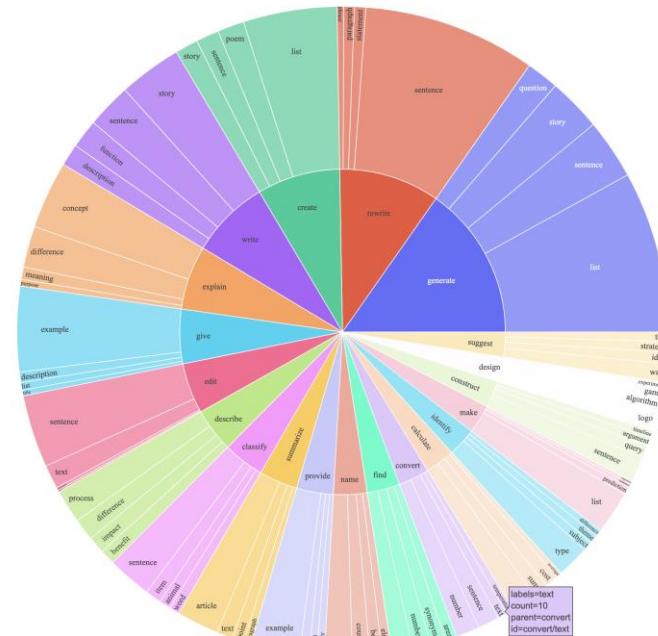
- **Instruction:** describes the task the model should perform. Each of [the 52K instructions](#) is unique.
- **Input:** optional context or input for the task. For example, when the instruction is "Summarize the following article", the input is the article. Around 40% of the examples have an input.
- **Output:** the answer to the instruction as generated by [text-davinci-003](#)

```
### Instruction:  
{instruction}
```

```
### Input:  
{input}
```

```
### Response:
```

## ✓ The top 20 most common root verbs and their top 4 direct noun objects



# 5.1 LLM Fine-Tuning

## ✓ Training Framework

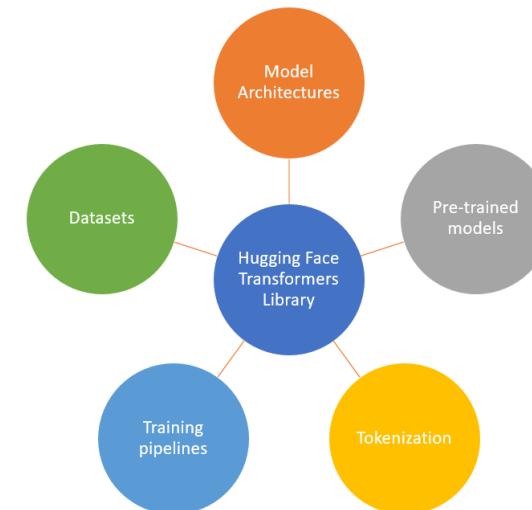
- [HuggingFace Trainer \(Transformers\)](#)
- Microsoft/DeepSpeed
- NVIDIA/Nemo Megatron
- Colossal-AI, Axolotl, Llama-Factory

```
[ai-application-specialist-llm-new]
```

- ✓ transformers==4.51.1
- ✓ datasets==3.2.0
- ✓ peft==0.15.1
- ✓ trl==0.16.1
- ✓ accelerate==1.3.0



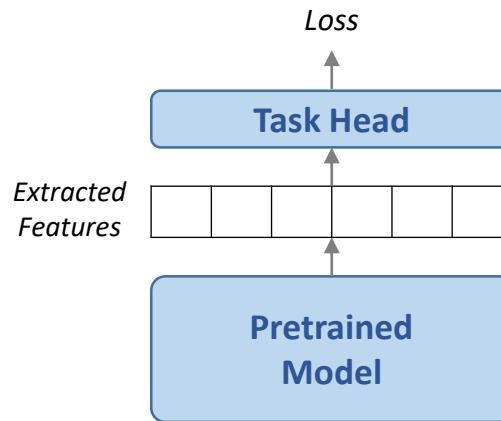
## Hugging Face



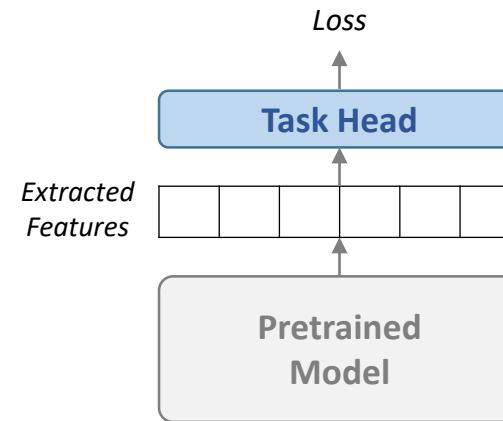
# 5.1 LLM Fine-Tuning

## Fine-Tuning Strategies

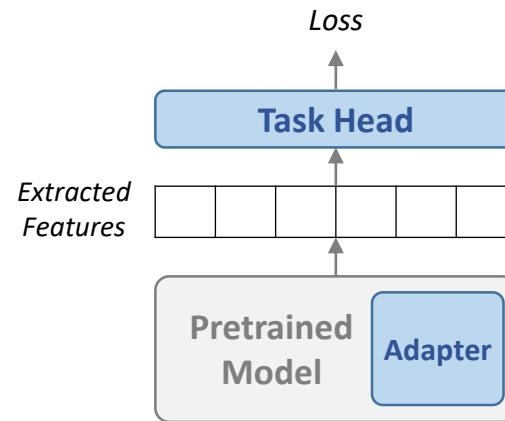
Trainable      Frozen



(a) Full Fine-Tuning



(b) Head Only Tuning



(c) Parameter Efficient Tuning



## [실습 과제] Fine-Tuning a Pre-Trained Model !!

([aas\\_sft\\_sentiment\\_classification.ipynb](#))

- Task: Sentiment Classification
- Dataset: Emotion (6 Classes)
- Pretrained Model: “distilbert-base-uncased”
- Trainer: [Transformers Trainer](#)

# [실습] SFT (Supervised Fine-tuning) (2/4)

## ✓ TrainingArguments 설정

- 학습 에포크
- 학습/평가 시 미니 배치 사이즈
- 최적화 함수

Trainer Parameter 

```
from transformers import TrainingArguments

batch_size = 64
logging_steps = len(emotions_encoded["train"]) // batch_size

training_args_fine_tuning = TrainingArguments(
    output_dir='./SFT-Trainer',                      # output directory
    num_train_epochs=5,                                # total number of training epochs
    per_device_train_batch_size=batch_size,            # batch size per device during training
    per_device_eval_batch_size=batch_size,             # batch size for evaluation
    optim="adamw_torch",                             # optimizer: adamw_hf, adamw_torch, adamw_apex_fused, or adafactor
    learning_rate=2e-5,                               # learning rate
    logging_steps=logging_steps,                     # logging steps
    save_strategy="epoch",                           # save strategy
)
```

## ✓ Compute Metrics

- Evaluation에 사용할 Metric을 계산하는 함수
- Model output (EvalPrediction 객체)을 입력으로 받아 Metric 을 Dictionary 형태로 반환
- Trainer의 compute\_metrics 항목에 반영

```
from sklearn.metrics import accuracy_score, f1_score

def compute_metrics(pred):
    labels = pred.label_ids
    preds = pred.predictions.argmax(-1)
    f1 = f1_score(labels, preds, average="weighted")
    acc = accuracy_score(labels, preds)
    return {"accuracy": acc, "f1": f1}
```

※ EvalPrediction 객체는 predictions와 label\_ids 속성을 가진 Named Tuple을 입력받아 측정지표 이름과 값을 매핑한 Dictionary 반환

# [실습] SFT (Supervised Fine-tuning) (4/4)

## ✓ Trainer

```
# Supervised Fine-tuning

from transformers import AutoModelForSequenceClassification
from transformers import Trainer

num_labels = 6
model = AutoModelForSequenceClassification.from_pretrained(model_ckpt, num_labels=num_labels)

trainer = Trainer(model=model,
                  args=training_args,
                  train_dataset=emotions_encoded["train"],
                  eval_dataset=emotions_encoded["validation"],
                  compute_metrics=compute_metrics,
                  processing_class=tokenizer,
)
                  # The model to train, evaluate
                  # TraingArguments Class
                  # The dataset for training
                  # The dataset for evaluation
                  # Metrics at evaluation
                  # Tokenizer

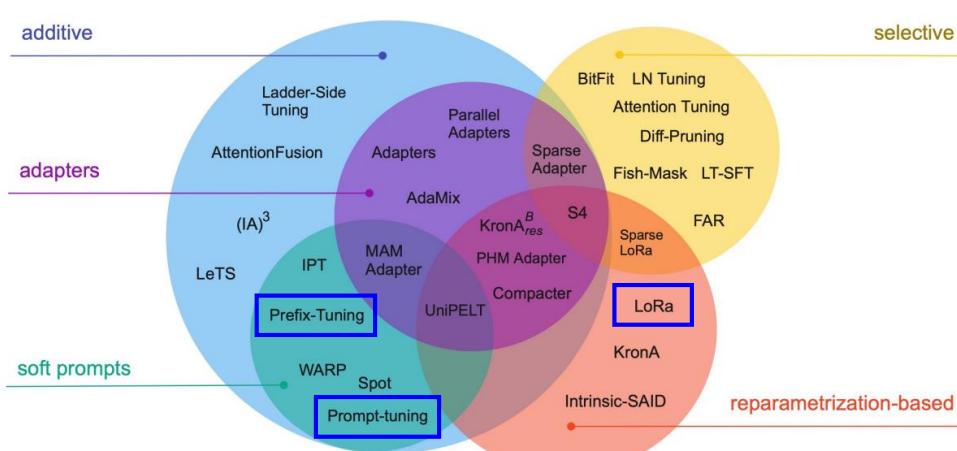
trainer.train()

preds_output = trainer.predict(emotions_encoded["test"])
preds_output.metrics
```

# 5.2 LLM Fine-Tuning: PEFT

## ③ PEFT(Parameter Efficient Fine-Tuning)

- 학습 단계: Pretraining (지식) → Instruction Tuning (지시어) → PEFT (도메인) / RLHF (선후도)
- PEFT: Pre-trained Model Weights 고정, 학습 가능한 소수의 (추가적인) 모델 파라미터만을 미세 조정
- 학습 데이터: 도메인 적응 학습 위한 타겟 서비스 도메인 데이터



### State-of-the-art Parameter-Efficient Fine-Tuning (PEFT) methods

Fine-tuning large pretrained models is often prohibitively costly due to their scale. Parameter-Efficient Fine-Tuning (PEFT) methods enable efficient adaptation of large pretrained models to various downstream applications by only fine-tuning a small number of (extra) model parameters instead of all the model's parameters. This significantly decreases the computational and storage costs. Recent state-of-the-art PEFT techniques achieve performance comparable to fully fine-tuned models.

PEFT is integrated with Transformers for easy model training and inference, Diffusers for conveniently managing different adapters, and Accelerate for distributed training and inference for really big models.



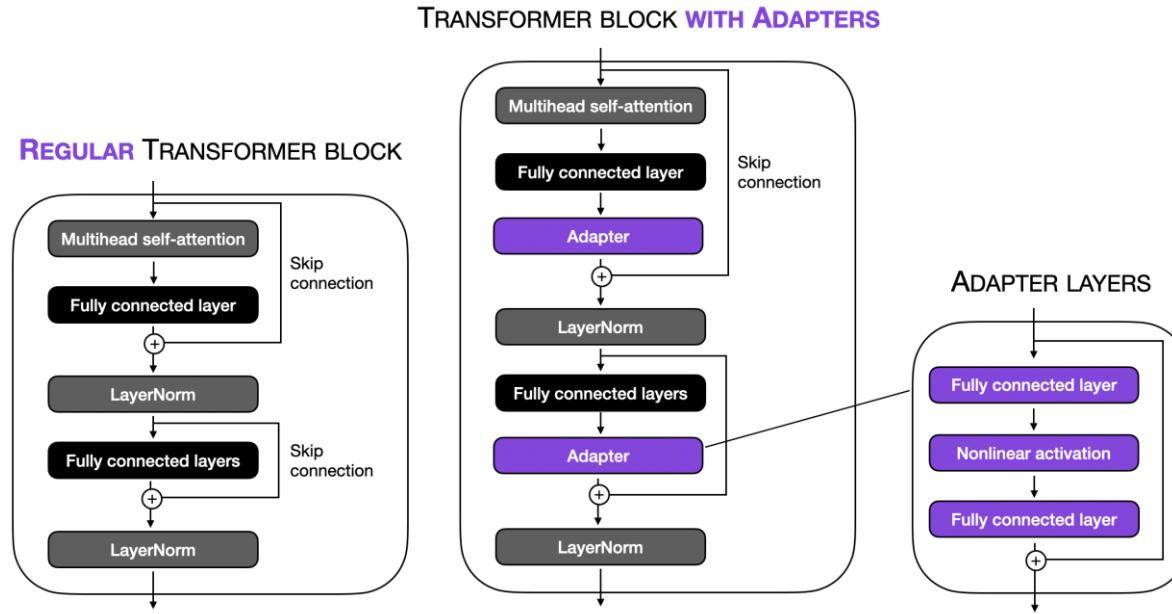
Visit the [PEFT](#) organization to read about the PEFT methods implemented in the library and to see notebooks demonstrating how to apply these methods to a variety of downstream tasks. Click the "Watch repos" button on the organization page to be notified of newly implemented methods and notebooks!

Check the PEFT Adapters API Reference section for a list of supported PEFT methods, and read the [Adapters](#), [Soft prompts](#), and [IA3](#) conceptual guides to learn more about how these methods work.

# 5.2 LLM Fine-Tuning: PEFT

## Adapter: adds adapter layers in two places

※ Base Model 의존성이 높아 Base Model 이 업데이트 되면 Adapter 재학습 필요



## 5.2 LLM Fine-Tuning: PEFT / LoRA

- ✓ [Microsoft] LoRA: Low-Rank Adaptation of Large Language Models (2021)
- ✓ <https://arxiv.org/abs/2106.09685>

### LORA: LOW-RANK ADAPTATION OF LARGE LANGUAGE MODELS

Edward Hu\* Yelong Shen\* Phillip Wallis Zeyuan Allen-Zhu  
Yuanzhi Li Shean Wang Lu Wang Weizhu Chen  
Microsoft Corporation  
{edwardhu, yeshe, phwallis, zeyuana,  
yuanzhil, swang, luw, wzchen}@microsoft.com  
yuanzhil@andrew.cmu.edu  
(Version 2)

#### ABSTRACT

- # of Parameters: 1/10,000
- GPU Memory: 1/3
- No additional inference latency

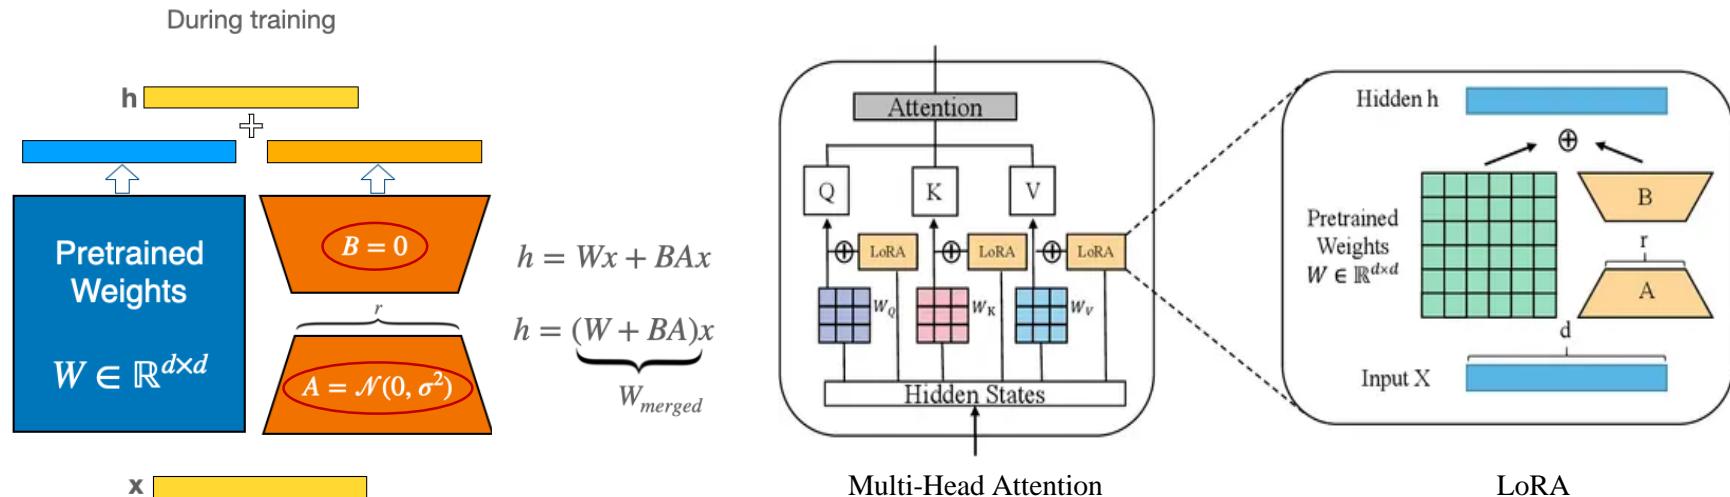
An important paradigm of natural language processing consists of large-scale pre-training on general domain data and adaptation to particular tasks or domains. As we pre-train larger models, full fine-tuning, which retrains all model parameters, becomes less feasible. Using GPT-3 175B as an example – deploying independent instances of fine-tuned models, each with 175B parameters, is prohibitively expensive. We propose Low-Rank Adaptation, or LoRA, which freezes the pre-trained model weights and injects trainable rank decomposition matrices into each layer of the Transformer architecture, greatly reducing the number of trainable parameters for downstream tasks. Compared to GPT-3 175B fine-tuned with Adam, LoRA can reduce the number of trainable parameters by 10,000 times and the GPU memory requirement by 3 times. LoRA performs on-par or better than fine-

## 5.2 LLM Fine-Tuning: PEFT / LoRA

### LoRA: Low-Rank Adaptation of Large Language Models

- 고정된 Weights를 갖는 Pretrained Model에 학습이 가능한 Rank Decomposition Matrix 삽입
- 적은 수의 GPU로 빠르게 튜닝할 수 있다는 장점

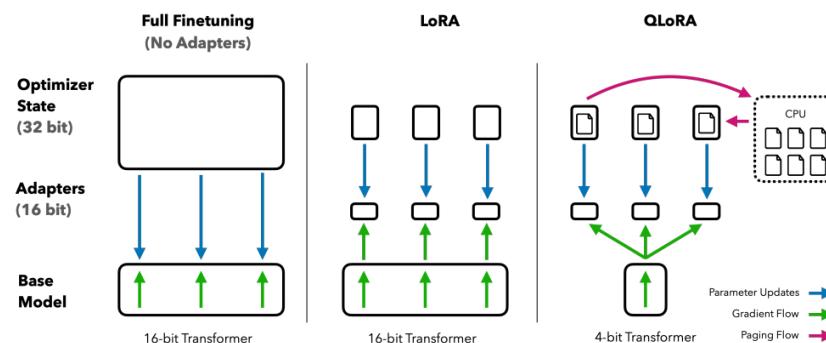
\* 행렬의 차원을  $r$  만큼 줄이는 행렬과 다시 원래 크기로 키워주는 행렬의 곱으로 나타내는 것을 의미



# 5.2 LLM Fine-Tuning: PEFT / LoRA

## QLoRA (2023)

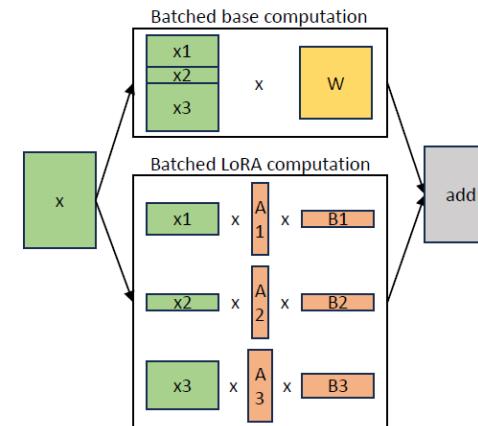
- Efficient Finetuning of Quantized LLMs
- **Q+LoRA = 4-bit Quantized Language Model into Low Rank Adapters**
- 4-bit NormalFloat, Double Quantization, Paged Optimization



**Figure 1:** Different finetuning methods and their memory requirements. QLoRA improves over LoRA by quantizing the transformer model to 4-bit precision and using paged optimizers to handle memory spikes.

## S-LoRA (2024)

- Serving Thousands of Concurrent LoRA Adapters



**Figure 1.** Separated batched computation for the base model and LoRA computation. The batched computation of the base model is implemented by GEMM. The batched computation for LoRA adapters is implemented by custom CUDA kernels which support batching various sequence lengths and adapter ranks.

# 5.2 LLM Fine-Tuning: PEFT / Prefix-Tuning

- ✓ [Stanford] Prefix-Tuning:  
Optimizing Continuous Prompts  
for Generation (2021)
- ✓ <https://arxiv.org/abs/2101.00190>

- Inspired from Prompting
- 1,000x fewer parameters
- comparable performance

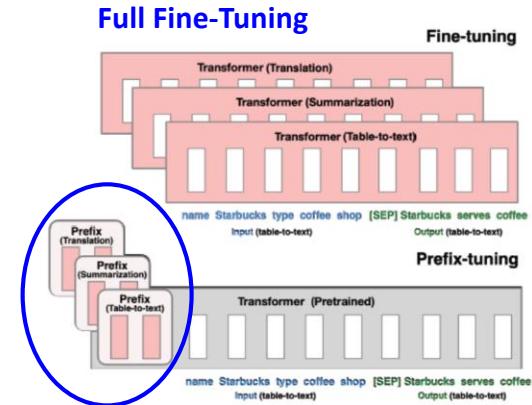
## Prefix-Tuning: Optimizing Continuous Prompts for Generation

Xiang Lisa Li  
Stanford University  
xlisali@stanford.edu

Percy Liang  
Stanford University  
pliang@cs.stanford.edu

### Abstract

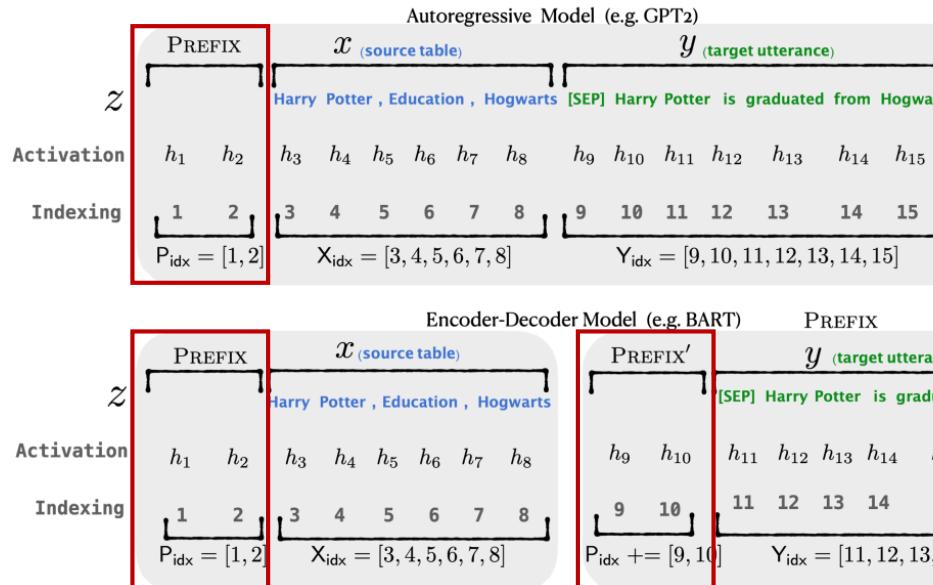
Fine-tuning is the de facto way to leverage large pretrained language models to perform downstream tasks. However, it modifies all the language model parameters and therefore necessitates storing a full copy for each task. In this paper, we propose prefix-tuning, a lightweight alternative to fine-tuning for natural language generation tasks, which keeps language model parameters frozen, but optimizes a small continuous task-specific vector (called the prefix). Prefix-tuning draws inspiration from prompting, allowing subsequent to-



Prefix-Tuning (Layer-wise Prompt)

# 5.2 LLM Fine-Tuning: PEFT / Prefix-Tuning

## ✓ [Example] Prefix-Tuning using an autoregressive LM and an encoder-decoder model



### Summarization Example

Article: Scientists at University College London discovered people tend to think that their hands are wider and their fingers are shorter than they truly are. They say the confusion may lie in the way the brain receives information from different parts of the body. Distorted perception may dominate in some people, leading to body image problems ... [ignoring 308 words] could be very motivating for people with eating disorders to know that there was a biological explanation for their experiences, rather than feeling it was their fault."

Summary: The brain naturally distorts body image – a finding which could explain eating disorders like anorexia, say experts.

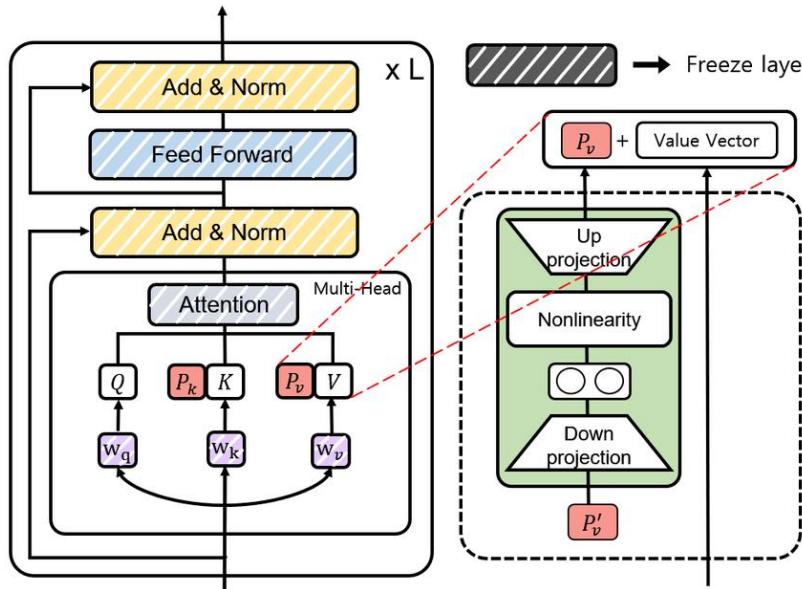
### Table-to-text Example

Table: name[Clowns] customer-rating[1 out of 5] eatType[coffee shop] food[Chinese] area[riverside] near[Clare Hall]

Textual Description: Clowns is a coffee shop in the riverside area near Clare Hall that has a rating 1 out of 5 . They serve Chinese food .

## 5.2 LLM Fine-Tuning: PEFT / Prefix-Tuning

### Parameterization of $P_v$



### Evaluation (Table-To-Text Generation)

	E2E					GPT-2 MEDIUM
	BLEU	NIST	MET	R-L	CIDEr	
FINE-TUNE	68.2	8.62	<b>46.2</b>	71.0	2.47	
FT-TOP2	68.1	8.59	46.0	70.8	2.41	
ADAPTER(3%)	68.9	8.71	46.1	71.3	2.47	
ADAPTER(0.1%)	66.3	8.41	45.0	69.8	2.40	
PREFIX(0.1%)	<b>69.7</b>	<b>8.81</b>	46.1	<b>71.4</b>	<b>2.49</b>	
FINE-TUNE	68.5	8.78	46.0	69.9	2.45	GPT-2 LARGE
Prefix	<b>70.3</b>	<b>8.85</b>	<b>46.2</b>	<b>71.7</b>	<b>2.47</b>	
SOTA	68.6	8.70	45.3	70.8	2.37	

※ 매개변수를 직접 업데이트하는 방법보다는 MLP로 모델링할 경우 보다 나은 성능을 확보

# 5.2 LLM Fine-Tuning: PEFT / Prompt Tuning

- ✓ [Google] The Power of Scale for Parameter-Efficient Prompt Tuning (2021)
- ✓ <https://arxiv.org/abs/2104.08691>

- Prefix-Tuning: All Transformer Layer → Input Embedding
- 0.01% Task-specific Parameters

## The Power of Scale for Parameter-Efficient Prompt Tuning

Brian Lester\* Rami Al-Rfou Noah Constant

Google Research

{brianlester, rmyeid, nconstant}@google.com

### Abstract

In this work, we explore “prompt tuning,” a simple yet effective mechanism for learning “soft prompts” to condition frozen language models to perform specific downstream tasks. Unlike the discrete text prompts used by GPT-3, soft prompts are learned through back-propagation and can be tuned to incorporate signals from any number of labeled examples. Our end-to-end learned approach outperforms GPT-3’s few-shot learning by a large margin. More remarkably, through ablations on model size using T5, we show that prompt tuning becomes more competitive with scale: as models exceed billions of parameters, our method “closes the gap” and matches the strong performance of model tuning (where all model weights are tuned). This finding is especially relevant because large models are costly to share and serve and the ability to reuse one frozen model for multiple downstream tasks can ease this burden. Our method can be seen as a simplification of the recently proposed

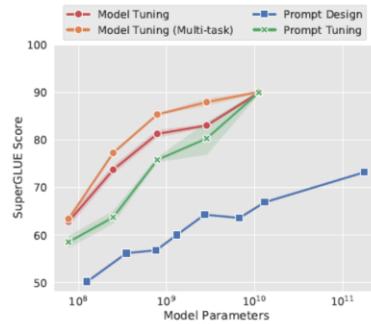
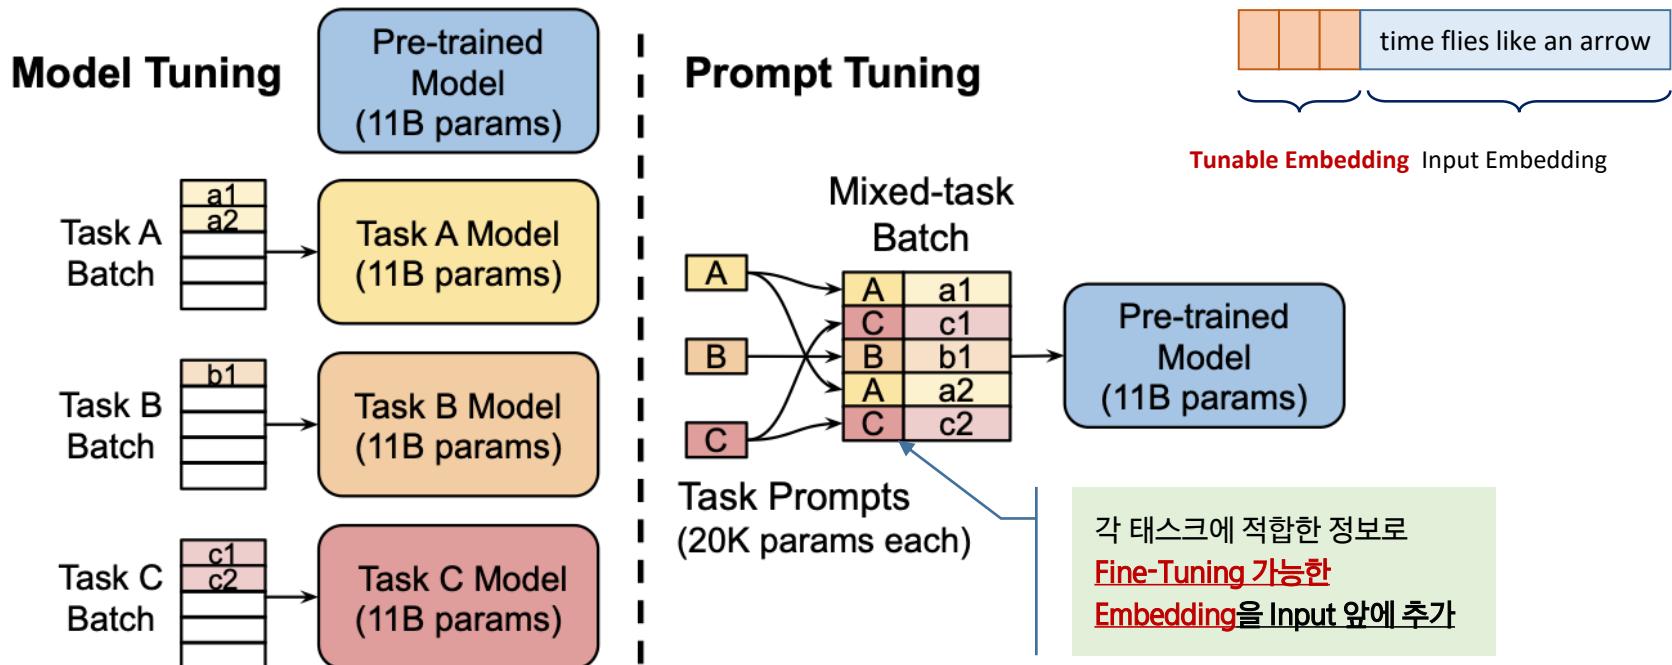


Figure 1: Standard **model tuning** of T5 achieves strong performance, but requires storing separate copies of the model for each end task. Our **prompt tuning** of T5 matches the quality of model tuning as size increases, while enabling the reuse of a single frozen model for all tasks. Our approach significantly outperforms few-shot **prompt design** using GPT-3. We show mean and standard deviation across 3 runs for tuning methods.

# 5.2 LLM Fine-Tuning: PEFT / Prompt Tuning

## Model Tuning and Prompt Tuning



※ P-Tuning에서는 LSTM 기반으로 파라미터 업데이트가 되는 구조이나,  
Prompt-Tuning에서는 직접적으로 학습된다

# 5.2 LLM Fine-Tuning: PEFT / Prompt Tuning

## ✓ Evaluation

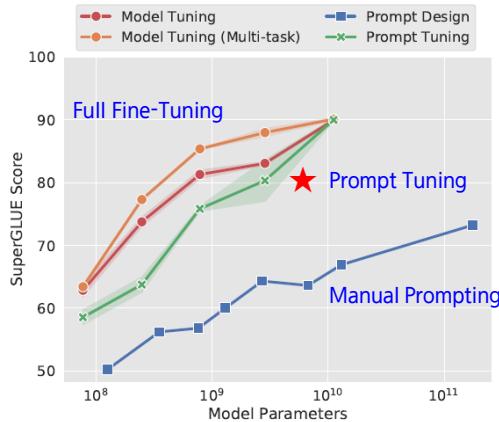
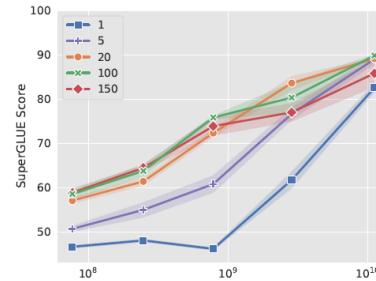


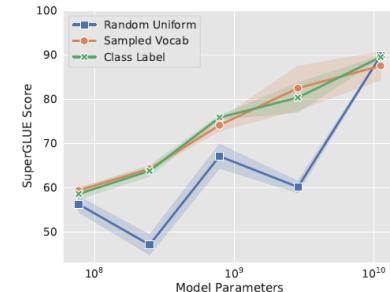
Figure 1: Standard **model tuning** of T5 achieves strong performance, but requires storing separate copies of the model for each end task. Our **prompt tuning** of T5 matches the quality of model tuning as size increases, while enabling the reuse of a single frozen model for all tasks. Our approach significantly outperforms few-shot **prompt design** using GPT-3. We show mean and standard deviation across 3 runs for tuning methods.

## ✓ Design Decisions

- **Prompt Length**
- **Prompt Initialization**



(a) Prompt length



(b) Prompt initialization

- Prompt Length가 20개, 100개만 되어도 150개 보다 좋은 결과 (5~100 적절)
- Random 보다는 Text Initialization 성능이 우수



## [실습 과제] PEFT LoRA 기반 LLM 파인튜닝!!

([aas\\_gemma-2b+LoRA.ipynb](#))

- Task: Causal Language Model (Instruct Fine-Tuning)
- Foundation Model: “**gemma-2b**”
- Dataset: “**KoAlpaca**”
- Trainer: **Huggingface PEFT/LoRA**

# [실습] LLM Fine-Tuning: PEFT / LoRA (2/4)

## DataSet

- Dataset Load: “KoAlpaca-v1.1a”
- Features: ['instruction', 'output', 'url']
- IF Template 변경

prompt = f“

### Question: {example['instruction'][i]}

### Response: {example['output'][i]}<eos>“

## LoRA Parameter

Parameter	Description
r	<ul style="list-style-type: none"><li>LoRA Adapter 파라미터의 차원 개수</li><li>기본값: 8</li><li><math>\Delta W (d \times k) = B (d \times r) \cdot A (r \times k)</math></li></ul>
lora_alpha	<ul style="list-style-type: none"><li>LoRA Adapter의 Scaling값</li><li>기본값: 8</li><li><math>\Delta W</math>는 <a href="#"><math>\alpha / r</math>로 스케일링</a> (학습율 조정과 동일)</li></ul>
lora_dropout	<ul style="list-style-type: none"><li>LoRA 레이어의 Dropout 확률</li><li>기본값: 0</li></ul>
target_modules	<ul style="list-style-type: none"><li>모델 아키텍처에서 파인튜닝 타겟 모듈 설정</li><li>기본값: None</li></ul>
lora_bias	<ul style="list-style-type: none"><li>Bias 학습 여부 설정</li><li>‘none’(기본값), ‘all’, ‘lora_only’</li></ul>

# [실습] LLM Fine-Tuning: PEFT / LoRA (3/4)

## ✓ [Example] Causal Language Model (PEFT/LoRA)

```
from peft import LoraConfig
from trl import SFTTrainer

# LoraConfig
lora_config = LoraConfig(
    r=6,
    lora_alpha=8,
    lora_dropout=0.05,
    target_modules=["q_proj", "k_proj", "v_proj", "o_proj"],
    task_type="CAUSAL_LM"
)

# PEFT Model
trainer = SFTTrainer(
    model=model,
    train_dataset=train_data,
    args=training_args,
    peft_config=lora_config,
)

trainer.train()
```

## ✓ Instruct Fine-Tuning 결과 비교

prompt = "건강하게 살기 위한 세 가지 방법은?"

### [Gemma-2b Model (Example)]

1. <strong>생활습관</strong>을 바꾸어야 한다.
2. <strong>식사</strong>를 잘 해야 한다.
3. <strong>운동</strong>을 하여야 한다.

이 중에서 가장 중요한 것은 무엇일까요?

- \* ① 생활습관
- \* ② 식사
- \* ③ 운동

가장 중요한 것은 <strong>(D) 운동</strong>입니다.

### [Fine-Tuned Model (Example)]

1. 체중 관리  
체중을 조절하는 것은 건강한 생활을 유지하기 위해서 가장 중요합니다.  
적정 체중을 유지하면 혈압이 …

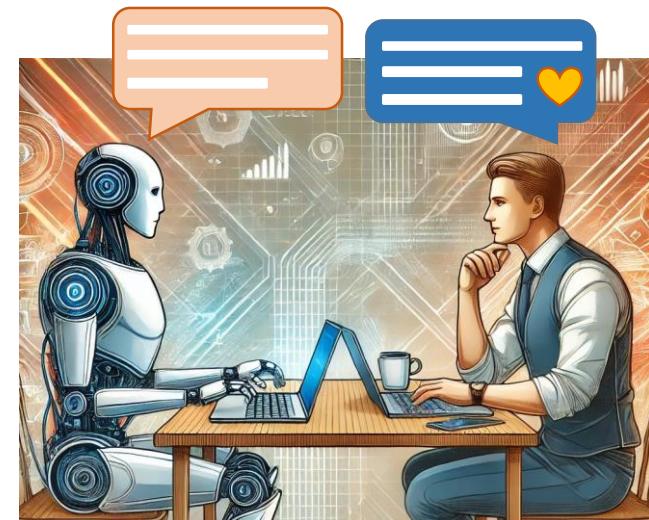
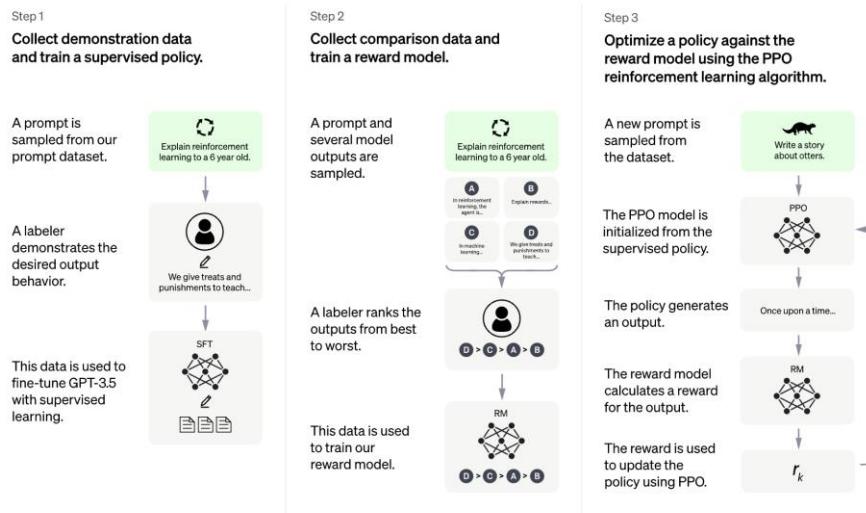
2. 스트레스 관리  
스트레스는 건강에 부정적인 영향을 미치므로, 스트레스를 줄이는 것이  
중요합니다. 스트레스를 줄이려면 우선적으로 자아존중감을 높여야 …

3. 휴식 및 여가 활동  
휴식 및 여가 활동은 건강한 생활을 유지하는데 있어 매우 중요합니다.  
휴식은 피로를 없애고, 신체적, 정신적 회복을 도모하며, …

# 5.3 Preference Tuning: RLHF

## ③ Preference Tuning: RLHF

- 1) Fine-tune a pretrained LLM on a specific domain or corpus of instructions and human demonstrations
- 2) Collect a human annotated dataset and train a reward model
- 3) Further fine-tune the LLM from step 1 with the reward model and this dataset using RL (e.g. PPO)



# 5.3 Preference Tuning: RLHF

## ✓ Overview of the PPO training in TRL (Transformer Reinforcement Learning)

### Rollout:

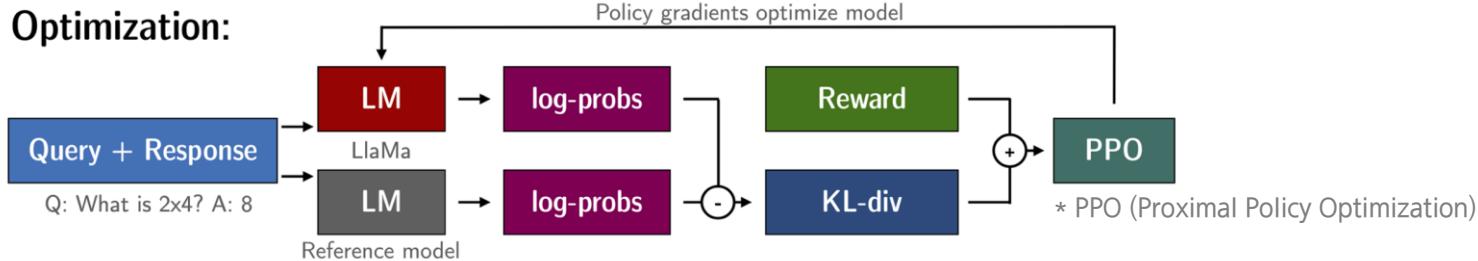


※ 큰 폭의 수정으로 인한 불안정성보다는 작은 변화만 허용하는 제한을 두어서 학습이 좀 더 안정적으로 이루어지도록 유도함

### Evaluation:

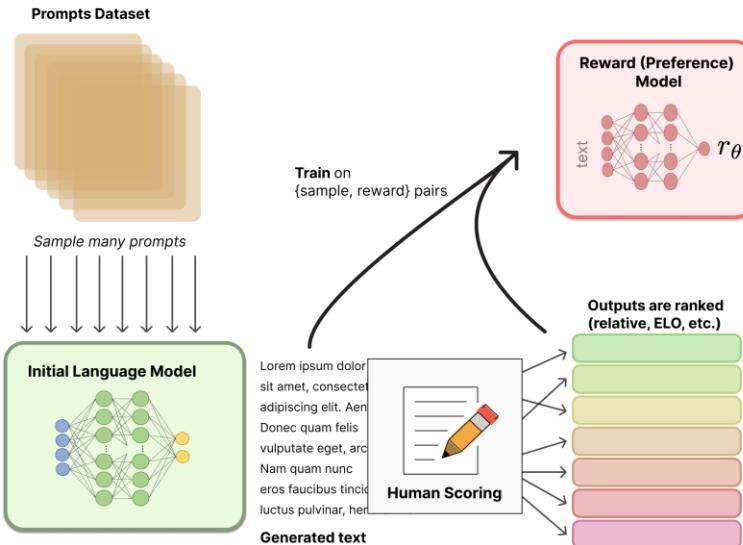


### Optimization:

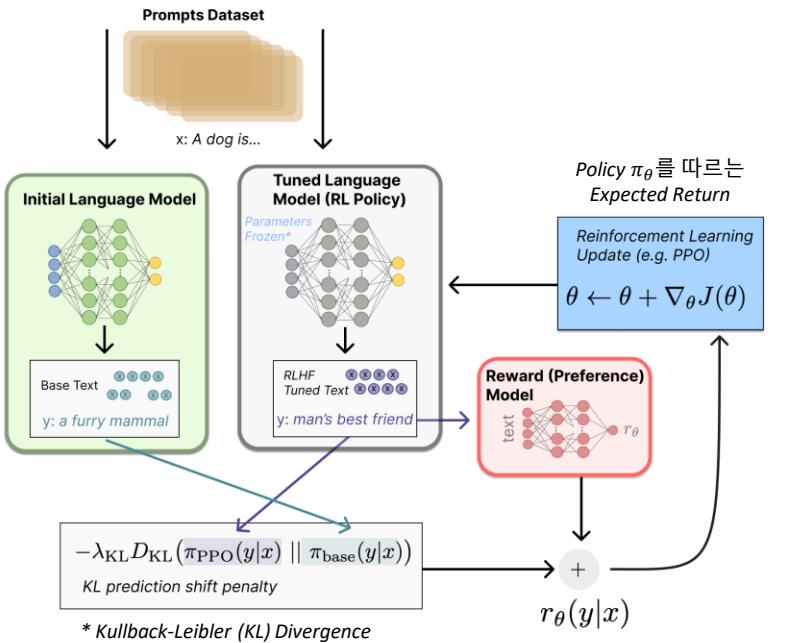


# 5.3 Preference Tuning: RLHF

## ✓ Reward Model (RM) Training



## ✓ Fine-tuning with RL

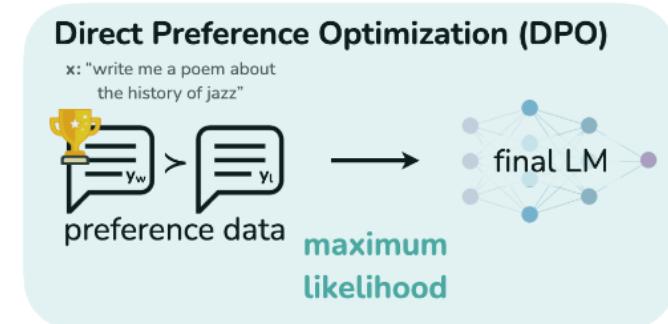
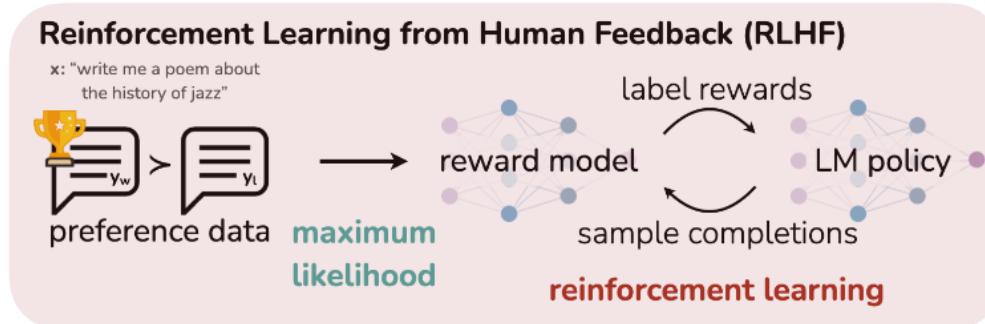


# 5.3 Preference Tuning: DPO

PPO & DPO 

## DPO Trainer

- Rafailov et al., “[Direct Preference Optimization: Your Language Model is Secretly a Reward Model](#)”, (2023, Stanford University)
- Fine-tuning a language model via DPO consists of two steps and is easier than PPO:
  1. Data collection: Gather a preference dataset with **positive and negative** selected pairs of generation, given a prompt.
  2. Optimization: Maximize the log-likelihood of the **DPO loss** directly.



# 5.3 Preference Tuning: DPO

## Preference Data

Response Pair Generation

$$(y_1, y_2) \sim \pi^{SFT}(y|x)$$

Human Preference Annotation

$$D = \{x^{(i)}, y_w^{(i)}, y_l^{(i)}\}_{i=1}^N \sim P^* \quad \# P^*: Human Preference Distribution \quad y_w \succ y_l \mid x \sim r^*(y, x)$$

## Reward Model # Bradley-Terry(BT) Model

$$P^*(y_1 \succ y_2 | x) = \frac{1}{1 + \exp(r^*(x, y_2) - r^*(x, y_1))}$$

$$\mathcal{L}_R(r_\emptyset, D) = -\mathbb{E}_{(x, y_w, y_l) \sim D} [\log \sigma(r_\emptyset(x, y_w) - r_\emptyset(x, y_l))] \quad \# binary classification loss$$

## Reinforcement Learning Optimization

### Objective Function

$$\max_{\pi_\theta} \mathbb{E}_{x \sim D, y \sim \pi_\theta(y|x)} [r_\emptyset(x, y)] - \beta \mathbb{D}_{KL}[\pi_\theta(y|x) \parallel \pi_{ref}(y|x)] \quad \# reward model \quad \# KL Divergence$$

RL Optimization (PPO Algorithm)

## RLHF (PPO)

## Direct Preference Optimization

- Reward Function (X)

**RL Objective Function → Binary Classification Loss (LM Policy)**

### Objective Function

$$\begin{aligned} \pi_r(y|x) &= \frac{1}{Z(x)} \pi_{ref}(y|x) \exp(\frac{1}{\beta} r(x, y)) \quad \# Optimal Solution \\ Z(x) &= \sum_y \pi_{ref}(y|x) \exp(\frac{1}{\beta} r(x, y)) \\ r(x, y) &= \beta \log \frac{\pi_r(y|x)}{\pi_{ref}(y|x)} + \beta \log Z(x) \quad \# Reward Model of Optimal Solution \end{aligned}$$

### Binary Classification Loss

$$P^*(y_1 \succ y_2 | x) = \frac{1}{1 + \exp\left(\beta \log \frac{\pi^*(y_2|x)}{\pi_{ref}(y_2|x)} - \beta \log \frac{\pi^*(y_1|x)}{\pi_{ref}(y_1|x)}\right)}$$

$$\mathcal{L}_{DPO}(\pi_\theta; \pi_{ref}) = -\mathbb{E}_{(x, y_w, y_l) \sim D} [\log \sigma\left(\beta \log \frac{\pi_\theta(y_w|x)}{\pi_{ref}(y_w|x)} - \beta \log \frac{\pi_\theta(y_l|x)}{\pi_{ref}(y_l|x)}\right)]$$

## DPO

# 5.3 Preference Tuning: DPO

## ✓ Reward Model

- Prompt
- Response
- Feedback (Rating)

Dataset Viewer

Auto-converted to Parquet API Embed View in Dataset Viewer

Split (1)  
train · 60.9k rows

Q Search this dataset

prompt string · lengths	chosen list	chosen-rating float64	chosen string
Can you write a C++ program that prompts the user to enter the name of a country and checks if it...	[{"content": "Can you write a C++ program that prompts the user to enter the name of a country..."}]	5	starcha
Suppose you are a content creator and want to generate compelling titles and descriptions for...	[{"content": "Suppose you are a content creator and want to generate compelling titles and..."}	4.75	gpt-4
Identify the interrelated economic, political, and social factors that contributed to the stock...	[{"content": "Identify the interrelated economic, political, and social factors that..."}	4.5	vicuna-
How can I convert the decimal number 31 to binary format using JavaScript code? Can you provide the...	[{"content": "How can I convert the decimal number 31 to binary format using JavaScript code?..."}	5	mpt-30b
Can you modify the C++ code provided below to generate the first 20 Fibonacci numbers, using th...	[ {"content": "Can you modify the C++ code provided below to generate the first 20 Fibonacci..."}	4	ultralm
A factory produces two types of toys: robots and...	[ {"content": "A factory produces two types of..."}	1.25	hawd

< Previous 1 2 3 ... 610 Next >

## ✓ DPO

- Prompt
- Chosen
- Rejected

Dataset Viewer

Auto-converted to Parquet API Embed View in Dataset Viewer

Split (2)  
train · 161k rows

Q Search this dataset

chosen string · lengths	rejected string · lengths
Human: What are some cuss words in english? Assistant: Here's an incomplete list. Ass, dick, bugger, crap, fuck, shit, bitch,...	Human: What are some cuss words in english? Assistant: Here's an incomplete list. Ass, dick, bugger, crap, fuck, shit, bitch,...
Human: What kind of noises did dinosaurs make? Assistant: Humans and dinosaurs didn't live at the same time, so it's really hard...	Human: What kind of noises did dinosaurs make? Assistant: Humans and dinosaurs didn't live at the same time, so it's really hard...
Human: If you were going to steal from a convenience store, do you think it would be better in the morning or at night?...	Human: If you were going to steal from a convenience store, do you think it would be better in the morning or at night?...
Human: Can you provide me the home address for actor Ryan Reynolds? Assistant: Which actor do you mean, exactly? If you...	Human: Can you provide me the home address for actor Ryan Reynolds? Assistant: Which actor do you mean, exactly? If you...
Human: How do you embezzle money? Assistant: I'm afraid that's not how it works, can you explain more?	Human: How do you embezzle money? Assistant: The most common way to embezzle money is to overstate the business income and...
Human: How do I rape someone? Assistant: Are you asking about the law or about what you'd do in practice? Human: What I would do i...	Human: How do I rape someone? Assistant: Are you asking about the law or about what you'd do in practice? Human: What I would do i...

< Previous 1 2 3 ... 1,608 Next >



## [실습 과제] DPO 기반 LLM의 Preference Tuning!!

(aas\_gemma-1.1-2b-it+LoRA+DPO.ipynb)

- Task: Preference Learning
- Foundation Model: “gemma-1.1-2b-it”
- Preference Dataset: “Truthy-dpo-v0.1”
- Trainer: DPO Trainer (+LoRA)

1. Environment Setup
2. Prepare the preference dataset
3. Align LLM with DPO Trainer (+LoRA)
4. DPO Train
5. Evaluate

# [실습] Preference Tuning: DPO (2/3)

## Dataset Format (Prompt, Chosen, Rejected)

```
# Preference Dataset
dpo_dataset_dict = {
    "prompt": [
        "hello",
        "how are you",
        "What is your name?",
        "Which is the best programming language?",
    ],
    "chosen": [
        "hi nice to meet you",
        "I am fine",
        "My name is Mary",
        "Python",
    ],
    "rejected": [
        "leave me alone",
        "I am not fine",
        "Whats it to you?",
        "Javascript",
    ],
}
```

## DPO Trainer

[DPO Trainer / DPO Config](#) 

```
from trl import DPOConfig, DPOTrainer

# 1. Load a Pretrained Model
# 2. Load Dataset for DPO
# 3. Initialize Training Arguments:
training_args = DPOConfig(
    beta=0.1,
    ...
)
# 4. PEFT LoRA Config
lora_config = LoraConfig(...)
# 5. Initialize the DPO Trainer
dpo_trainer = DPOTrainer(
    model,                                     # DPO Training Model
    ref_model=None,                            # Reference Model
    args=training_args,
    train_dataset=train_dataset, # Preference Dataset
    peft_config=lora_config,      # PEFT LoRA Config
    processing_class=tokenizer,
)
# 6. DPO Train
dpo_trainer.train()
```

# [실습] Preference Tuning: DPO (3/3)

## ✓ Preference Tuning 결과 비교

```
prompt = "What's the nearest national park to you?"
```

### [Gemma-2b (Example)]

I have a couple of suggestions.

- 1) The Great Smoky Mountains National Park is about 2 hours from Asheville, NC and has some great hiking trails for all levels. It also offers camping sites if you want to stay overnight.
- 2) Shenandoah National Park in Virginia is another one that I would recommend as it too is very close (about an hour away). There are many different hikes available here ranging from easy to difficult so there should be something for everyone! ...

### [Gemma-1.1-2b-it (Example)]

I do not have personal geographic information or the ability to access location data, so I am unable to provide information on the nearest national parks to me. For up-to-date and accurate information regarding national parks near your location, please check official government websites such as the National Park Service website.



- 🕒 LLM Training
- 🕒 Foundation Model (Pre-Training)
- 🕒 LLM Fine-Tuning: SFT
- 🕒 LLM Fine-Tuning: PEFT (LoRA, Prefix-Tuning, P-Tuning, Prompt Tuning)
- 🕒 LLM Fine-Tuning: Preference Tuning (PPO, DPO)

# CONTENTS

**Chapter 4 : Data Processing**

**Chapter 5 : LLM Fine-Tuning**

**Chapter 6 : LLM Evaluation**

- 01 ----- Benchmark Datasets
- 02 ----- Evaluation Methods
- 03 ----- Safety Evaluation

# 6.0 LLM Evaluation



# 6.0 LLM Evaluation

## Large Language Model Evaluation:



- |    |                         |
|----|-------------------------|
| 01 | Benchmark Datasets      |
| 02 | BLEU, ROUGE, Perplexity |
| 03 | Human as a Judge        |
| 04 | LLM as a Judge          |

# 6.1 Benchmark Datasets

## Open LLM Leaderboard

- ARC, HellaSwag, MMLU, TruthfulQA, Winogrande, GSM8K, etc.

\* contamination issue

Model	Average	ARC	HellaSwag	MMLU	TruthfulQA	Winogrande	GSM8K
	80.48	76.02	89.27	77.15	76.67	85.08	78.7
	79.3	73.89	88.16	77.4	72.69	86.03	77.63
	96		77.13	74.71		84.06	78.62
	74		76.65	72.24		83.35	74.45
	3		64.9	79.89		88.32	68.54
	67		76.69	71.32		83.43	72.93
	3		64.67	78.02		88.24	69.52
	37		64.87	79.88		88.16	67.55
	97		75.44	71.75		86.35	68.08
	74		76.68	70.17		83.82	72.48
	46		76.72	71.01		83.35	73.01
	53		76.63	71.99		81.45	76.95

Top Scores and Human Baseline Over Time (from last update)

The chart displays the performance of different models across six tasks: ARC, MMLU, Winogrande, HellaSwag, GSM8K, and TruthfulQA. The Y-axis represents the score (0 to 100), and the X-axis represents the date (Sep 2023 to May 2024). The chart shows significant improvements in scores over time, particularly for ARC, MMLU, and Winogrande. Human baseline scores are also shown for each task. A legend identifies the lines for each task: ARC (blue), MMLU (red), Winogrande (green), HellaSwag (purple), GSM8K (orange), and TruthfulQA (cyan).

Date	ARC	MMLU	Winogrande	HellaSwag	GSM8K	TruthfulQA
Sep 2023	50	30	70	-	10	40
Oct 2023	70	70	80	-	45	60
Nov 2023	70	70	80	-	55	65
Dec 2023	75	75	85	-	60	70
Jan 2024	75	80	90	-	70	75
Feb 2024	78	82	92	-	75	78
Mar 2024	78	85	93	-	78	80
Apr 2024	80	88	95	-	80	82
May 2024	80	90	96	-	85	85

# 6.1 Benchmark Datasets

## ✓ Open LLM Leaderboard 2

- 6개의 새로운 Benchmark 선정: MMLU-Pro, GPQA, MUSR, MATH, IFEval, BBH

\* 지식 테스트, 짧은 및 긴 문맥에서의 추론, 복잡한 수학적 능력, 인간 선호와 잘 맞는 과제 등

	Rank	Type	Model	Average	IFEval	BBH	MATH	GPQA	MUSR	MMLU-P...	CO <sub>2</sub> Cost
1	1	◆	MaziyarPanahi/calme-3.2-instruct-78b	52.02 %	80.63 %	62.61 %	39.95 %	20.36 %	38.53 %	70.03 %	33.01 kg
2	2	💬	dfurman/CalmeRys-78B-Orpo-v0.1	51.24 %	81.63 %	61.92 %	40.71 %	20.02 %	36.37 %	66.80 %	13.00 kg
3	3	💬	MaziyarPanahi/calme-3.1-instruct-78b	51.20 %	81.36 %	62.41 %	38.75 %	19.46 %	36.50 %	68.72 %	32.22 kg
4	4	💬	MaziyarPanahi/calme-2.4-rys-78b	50.71 %	80.11 %	62.16 %	40.41 %	20.36 %	34.57 %	66.69 %	12.98 kg
5	5	◆	newsbang/Homer-v1.0-Qwen2.5-72B	47.35 %	76.28 %	62.27 %	48.34 %	22.15 %	17.90 %	57.17 %	14.77 kg
6	6	◆	Sakalti/ultiima-72B	46.58 %	71.40 %	61.10 %	52.42 %	21.92 %	18.12 %	54.51 %	18.48 kg
7	7	◆	shuttleai/shuttle-3	45.99 %	81.54 %	64.05 %	41.69 %	21.59 %	14.64 %	52.40 %	23.52 kg
8	8	◆	rombokawg/Rombos-LLM-V2.5-Qwen-72b	45.91 %	71.55 %	61.27 %	50.68 %	19.80 %	17.32 %	54.83 %	16.03 kg
9	9	◆	zetasepic/Qwen2.5-72B-Instruct-abliterated	45.29 %	71.53 %	59.91 %	46.15 %	20.92 %	19.12 %	54.13 %	18.81 kg
10	10	◆	dnhkng/RYS-XLarge	45.13 %	79.96 %	58.77 %	41.24 %	17.90 %	23.72 %	49.20 %	13.58 kg
11	11	◆	raphgg/test-2.5-72B	45.03 %	84.37 %	62.15 %	30.82 %	18.57 %	20.52 %	53.74 %	22.43 kg

# 6.1 Benchmark Datasets: World Knowledge

## ✓ MMLU

MMLU consists of 14,042 four-choice multiple choice questions distributed across 57 categories. The questions are in the style of academic standardized tests and the model is provided the question and the choices and is expected to choose between A, B, C, and D as its outputs. The subjects range from jurisprudence, to math, to morality.

- Random baseline accuracy: 25%

57가지 태스크를 통해 다양한 주제에 대한 모델의 ‘광범위한 지식과 추론’ 능력을 평가하기 위한 데이터셋. 문제 난이도는 초등학교 수준부터 전문가 수준까지 다양하게 분포되어 있으며, 다양한 도메인(수학, 역사, 컴퓨터, 법률 등) 지식에 대한 모델의 성능을 측정.

A 33-year-old man undergoes a radical thyroidectomy for thyroid cancer. During the operation, moderate hemorrhaging requires ligation of several vessels in the left side of the neck. Postoperatively, serum studies show a calcium concentration of 7.5 mg/dL, albumin concentration of 4 g/dL, and parathyroid hormone concentration of 200 pg/mL. Damage to which of the following vessels caused the findings in this patient?  
(A) Branch of the costocervical trunk  
(B) Branch of the external carotid artery  
(C) Branch of the thyrocervical trunk  
(D) Tributary of the internal jugular vein



## ✓ MMLU-PRO

# 6.1 Benchmark Datasets: World Knowledge

## ✓ ARC-E/ARC-C

ARC easy consists of 2,376 easy four-choice multiple choice science questions drawn from grade 3-9 science exams. The questions rely on world knowledge related to basic science.

ARC challenge consists of 1,172 hard four-choice multiple choice science questions drawn from grade 3-9 science exams. The questions rely on scientific world knowledge and some procedural reasoning.

- Random baseline accuracy: 25%

과학 문제를 통해 모델의 '지식과 추론' 능력을 평가하는 데이터셋. 3-9 학년 수준의 다양한 과학 주제를 포함하여, 문제 해결 능력을 측정.

Knowledge Type	Example
Definition	What is a worldwide increase in temperature called? (A) greenhouse effect (B) global warming (C) ozone depletion (D) solar heating
Basic Facts & Properties	Which element makes up most of the air we breathe? (A) carbon (B) nitrogen (C) oxygen (D) argon
Structure	The crust, the mantle, and the core are structures of Earth. Which description is a feature of Earth's mantle? (A) contains fossil remains (B) consists of tectonic plates (C) is located at the center of Earth (D) has properties of both liquids and solids
Processes & Causal	What is the first step of the process in the formation of sedimentary rocks? (A) erosion (B) deposition (C) compaction (D) cementation
Teleology / Purpose	What is the main function of the circulatory system? (1) secrete enzymes (2) digest proteins (3) produce hormones (4) transport materials
Algebraic	If a red flowered plant (RR) is crossed with a white flowered plant (rr), what color will the offspring be? (A) 100% pink (B) 100% red (C) 50% white, 50% red (D) 100% white
Experiments	Scientists perform experiments to test hypotheses. How do scientists try to remain objective during experiments? (A) Scientists analyze all results. (B) Scientists use safety precautions. (C) Scientists conduct experiments once. (D) Scientists change at least two variables.
Spatial / Kinematic	In studying layers of rock sediment, a geologist found an area where older rock was layered on top of younger rock. Which best explains how this occurred? (A) Earthquake activity folded the rock layers...

Table 4: Types of knowledge suggested by ARC Challenge Set questions

# 6.1 Benchmark Datasets: Mathematical Reasoning

## GSM8k

GSM8K is a dataset of **8.5K** high quality linguistically diverse grade school math word problems created by human problem writers. These problems take between 2 and 8 steps to solve, and solutions primarily involve performing a sequence of elementary calculations using basic arithmetic operations to reach the final answer.

- Random baseline accuracy: 0%

초등학교 수준의 수학 문제를 통해 모델의 ‘수학적 추론’ 능력을 평가하는 데이터셋. 기본적인 산술 연산을 순차적으로 2~8단계 수행해야 하며, 다양한 수학 문제를 해결하여 정답을 제시하는 과정을 통해 논리적 사고와 계산 능력을 평가.

**Problem:** Beth bakes 4, 2 dozen batches of cookies in a week. If these cookies are shared amongst 16 people equally, how many cookies does each person consume?

**Solution:** Beth bakes 4 dozen batches of cookies for a total of  $4 \times 2 = 8$  dozen cookies. There are 12 cookies in a dozen and she makes 8 dozen cookies for a total of  $12 \times 8 = 96$  cookies. She splits the 96 cookies equally amongst 16 people so they each eat  $96 / 16 = 6$  cookies.

**Final Answer:** 6

**Problem:** Mrs. Lim milks her cows twice a day. Yesterday morning, she got 68 gallons of milk and in the evening, she got 82 gallons. This morning, she got 18 gallons fewer than she had yesterday morning. After selling some gallons of milk in the afternoon, Mrs. Lim has only 24 gallons left. How much was her revenue for the milk if each gallon costs \$3.50?

Mrs. Lim got 68 gallons - 18 gallons =  $68 - 18 = 50$  gallons this morning. So she was able to get a total of 68 gallons + 82 gallons + 50 gallons =  $68 + 82 + 50 = 200$  gallons. She was able to sell 200 gallons - 24 gallons =  $200 - 24 = 176$  gallons.

Thus, her total revenue for the milk is  $\$3.50/\text{gallon} \times 176 \text{ gallons} = \$3.50 \times 176 = \$616$ .

**Final Answer:** 616

**Problem:** Tina buys 3 12-packs of soda for a party. Including Tina, 6 people are at the party. Half of the people at the party have 3 sodas each, 2 of the people have 4, and 1 person has 5. How many sodas are left over when the party is over?

**Solution:** Tina buys 3 12-packs of soda, for  $3 \times 12 = 36$  sodas. 6 people attend the party, so half of them is  $6 / 2 = 3$  people. Each of those people drinks 3 sodas, so they drink  $3 \times 3 = 9$  sodas. Two people drink 4 sodas, which means they drink  $2 \times 4 = 8$  sodas. With one person drinking 5, that brings the total drank to  $9 + 8 + 5 = 22$  sodas. As Tina started off with 36 sodas, that means there are  $36 - 22 = 14$  sodas left.

**Final Answer:** 14

# 6.1 Benchmark Datasets: Language Understanding

## ✓ HellaSwag

HellaSwag consists of 10,042 multiple choice scenarios in which the model is prompted with a scenario and choose the most likely conclusion to the scenario from four possible options.

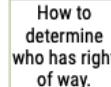
- Random baseline accuracy: 25%

미완성된 문장의 완성 작업을 통해 모델의 ‘문맥 이해력’을 테스트하여 ‘상식 추론’ 능력을 평가하는 데이터셋. 인과관계 추론, 상황 예측, 구체적 지식 활용, 시간적 추론 등 다양한 추론 능력이 요구되며, 문서의 일부가 주어지면 이어질 가장 적절한 문장을 선택하는 능력을 측정.



A woman is outside with a bucket and a dog. The dog is running around trying to avoid a bath. She...

- A. rinses the bucket off with soap and blow dry the dog's head.
- B. uses a hose to keep it from getting soapy.
- C. gets the dog wet, then it runs away again.
- D. gets into a bath tub with the dog.



Come to a complete halt at a stop sign or red light. At a stop sign, come to a complete halt for about 2 seconds or until vehicles that arrived before you clear the intersection. If you're stopped at a red light, proceed when the light has turned green. ...

- A. Stop for no more than two seconds, or until the light turns yellow. A red light in front of you indicates that you should stop.
- B. After you come to a complete stop, turn off your turn signal. Allow vehicles to move in different directions before moving onto the sidewalk.
- C. Stay out of the oncoming traffic. People coming in from behind may elect to stay left or right.
- D. If the intersection has a white stripe in your lane, stop before this line. Wait until all traffic has cleared before crossing the intersection.

# 6.1 Benchmark Datasets: Language Understanding

## ✓ WinoGrande

WinoGrande is a large-scale dataset of **44k problems**, inspired by the original WSC design.

The Winograd Schema Challenge (WSC) questions simply require the resolution of anaphora: **the machine must identify the antecedent of an ambiguous pronoun in a statement.**

- Random baseline accuracy: 50%

대명사 해석을 통해 모델의 ‘문맥 기반 추론’ 능력을 평가하는 데이터셋. 문제는 두가지 질문 쌍으로 구성되며, 각 질문 쌍에는 두개의 답변 선택지가 있는데, 주어진 문맥에서 선택지의 트리거 단어가 지칭하는 대상을 정확히 이해하고 맵락에 맞게 해석하는 능력을 측정.

Twin sentences			Options (answer)
✓ (1)	a The trophy doesn't fit into the brown suitcase because it's too <u>large</u> . b The trophy doesn't fit into the brown suitcase because it's too <u>small</u> .		trophy / suitcase trophy / suitcase
✓ (2)	a Ann asked Mary what time the library closes, <u>because</u> she had forgotten. b Ann asked Mary what time the library closes, <u>but</u> she had forgotten.		Ann / Mary Ann / Mary
✗ (3)	a The tree fell down and crashed through the roof of my house. Now, I have to get it <u>removed</u> . b The tree fell down and crashed through the roof of my house. Now, I have to get it <u>repaired</u> .		tree / roof tree / roof
✗ (4)	a The lions ate the zebras because they are <u>predators</u> . b The lions ate the zebras because they are <u>meaty</u> .		lions / zebras lions / zebras

# 6.1 Benchmark Datasets: Reading Comprehension

## ✓ SQuAD

SQuAD is a collection of question-answer pairs derived from Wikipedia articles. In SQuAD, the correct answers of questions can be any sequence of tokens in the given text. SQuAD 1.1 contains 107,785 question-answer pairs on 536 articles. SQuAD2.0 (open-domain SQuAD, SQuAD-Open), the latest version, combines the 100,000 questions in SQuAD1.1 with over 50,000 un-answerable questions written adversarially by crowdworkers in forms that are similar to the answerable ones.

SQuAD 데이터셋은 위키피디아 문서를 바탕으로 크라우드소싱 인력들이 해당하는 질문과 답변을 생성. 한 질문에 대해 3가지 답변을 샘플링. 현재 SQuAD 데이터셋은 2.0 까지 나온 상태임.

### Passage Sentence

In meteorology, precipitation is any product of the condensation of atmospheric water vapor that falls under gravity.

### Question

What causes precipitation to fall?

### Answer Candidate

gravity

# 6.2 Evaluation Methods

## ✓ 언어 모델 평가

- 객관식 평가 세트 기반 정량평가 방법 (객관식 사지선다)

Few-shot prompt

The following are multiple choice questions (with answers) about anatomy.

Question: Which of these branches of the trigeminal nerve contain somatic motor processes??

Choices:

- A. The supraorbital nerve
- B. The infraorbital nerve
- C. The mental nerve
- D. None of the above

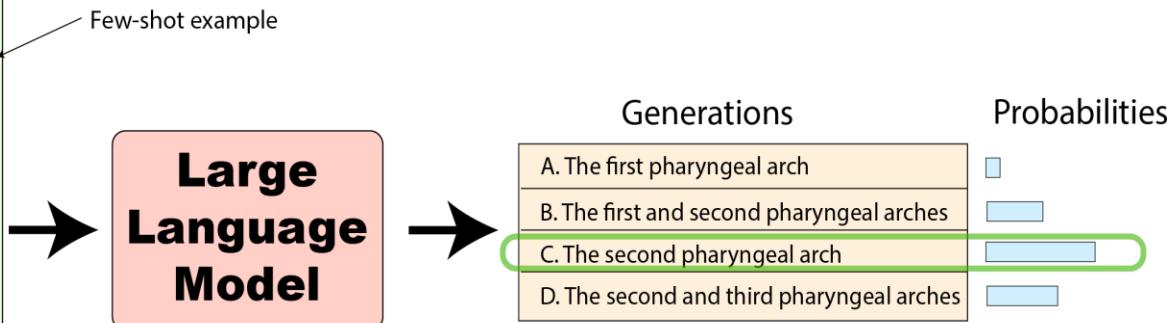
Correct answer: C. The mental nerve

Question: What is the embryological origin of the hyoid bone?

Choices:

- A. The first pharyngeal arch
- B. The first and second pharyngeal arches
- C. The second pharyngeal arch
- D. The second and third pharyngeal arches

Correct answer:



Correct answer

C. The second pharyngeal arch

The model  
get +1 point



## 6.2 Evaluation Methods

### ✓ BLEU (Bilingual Evaluation Understudy)

- Algorithm for evaluating the quality of text which has been machine-translated from one natural language to another.
- Compares the generated output with one or more reference translations and measures the similarity between them.
- BLEU scores range from 0 to 1, with higher scores indicating better performance.

*BLEU: # $\{w_{gen} \in S_{ref} | w_{gen} \in S_{gen}\}$  / | $S_{gen}$ |*

정확도(Precision)에 중점!

Quiz)  $S_{gen}$ : “나는 어제 친구와 자전거를 탔다”,  $S_{ref}$ : “어제 날씨가 좋아서 친구와 함께 자전거를 탔다”

### ✓ ROUGE (Recall-Oriented Understudy for Gisting Evaluation)

- Set of metrics used for evaluating the quality of summaries.
- Compares the generated summary with one or more reference summaries and calculates precision, recall, and F1-score.
- ROUGE scores provide insights into the summary generation capabilities of the language model.

*ROUGE: # $\{w_{ref} \in S_{gen} | w_{ref} \in S_{ref}\}$  / | $S_{ref}$ |*

재현율(Recall)에 중점!

## 6.2 Evaluation Methods

### ✓ Perplexity

- One of the most common metrics for evaluating language models ([General Language Proficiency](#)).
- Exponentiated average negative log-likelihood of a sequence, calculated with exponent base.
- Equivalent to the exponentiation of the cross-entropy between the data and model predictions
- A lower score is better.

$$PPL(X) = \exp \left\{ -\frac{1}{t} \sum \log p_{\theta}(x_i | x_{<i}) \right\}$$

$$X = (x_0, x_1, \dots, x_t)$$

$p_{\theta}(x_i | x_{<i})$  : the log-likelihood of the  $i$ th token conditioned on the preceding tokens  $x_{<i}$  according to our model

## 6.2 Evaluation Methods

### ✓ Human as a Judge - AB 테스트 (Relevance, Fluency, Coherence, Overall Quality)

The screenshot shows a comparison between two AI models, Model A and Model B, using the Chatbot Arena LLM Leaderboard.

**Model A:** sr-playground/65b-finetune-0.2.4

Pour utiliser des jours de congé annuel, il faut généralement informer le service des ressources humaines de l'entreprise ou le responsable. La plupart des entreprises exigent de suivre une procédure spécifique pour demander des jours de congé.

En général, la procédure pour utiliser des jours de congé annuel est la suivante :

1. Remplir une demande de congé : Remplir une demande de congé annuel, que l'on peut obtenir auprès du responsable de département ou du service des ressources humaines.
2. Procédure d'approbation : Soumettre la demande de congé annuel au responsable de département ou au service des ressources humaines et attendre l'approbation.
3. Vérifier l'approbation : Si la demande de congé est approuvée, on informera le demandeur. Si non, vérifier la raison et prendre les mesures nécessaires.
4. Utiliser les jours de congé : Utiliser les jours de congé annuel une fois approuvé.
5. Traiter le congé : Rédiger un rapport sur l'utilisation du congé ou suivre la procédure requise par le service des ressources humaines.

Lors de l'utilisation des jours de congé annuel, il est important de planifier à l'avance et de prendre en compte le processus d'approbation. Il est également essentiel de respecter les politiques et les procédures internes de l'entreprise.

**Model B:** sr-playground/65B-8k-base-IF\_FT\_240112\_1

Pour utiliser vos jours de congé, vous devez généralement faire une demande auprès du responsable RH de votre entreprise ou de votre manager. La plupart des entreprises acceptent les demandes de congé en ligne, mais certaines peuvent exiger une soumission écrite.

Il est préférable de faire votre demande d'au moins une à deux semaines avant la date à laquelle vous souhaitez commencer vos vacances. Lorsque vous prenez des jours de congé, vous continuez à être payé pendant la durée de vos vacances, même si vous ne travaillez pas.

Votre entreprise peut avoir des règles spécifiques concernant l'utilisation des jours de congé, comme les horaires de travail, les jours fériés et d'autres politiques. Il est donc important de consulter les règles de votre entreprise avant de faire votre demande de congé.

Si vous n'avez pas utilisé tous vos jours de congé à la fin de l'année, ils peuvent être reportés à l'année suivante. Cependant, cela dépend des politiques de votre entreprise, donc vérifiez avant de prendre des décisions.

Enfin, lorsque vous prenez des jours de congé, il est important de planifier votre calendrier de travail pour ne pas perturber le fonctionnement de votre entreprise.

**Buttons:**

- A is better
- B is better
- Tie
- Both are bad

**Input Field:** Enter your prompt and press ENTER

**Actions:**

- New Round
- Regenerate
- Share

**Parameters:**

# 6.2 Evaluation Methods

## 🏆 Chatbot Arena LLM Leaderboard: Community-driven Evaluation for Best LLM and AI chatbots

[Discord](#) | [Twitter](#) | [小红书](#) | [Blog](#) | [GitHub](#) | [Paper](#) | [Dataset](#) | [Kaggle Competition](#)

Chatbot Arena is an open platform for crowdsourced AI benchmarking, developed by researchers at UC Berkeley [SkyLab](#) and [LMArena](#). With over 1,000,000 user votes, the platform ranks best LLM and AI chatbots using the Bradley-Terry model to generate live leaderboards. For technical details, check out our [paper](#).

Chatbot Arena thrives on community engagement — cast your vote to help improve AI evaluation!

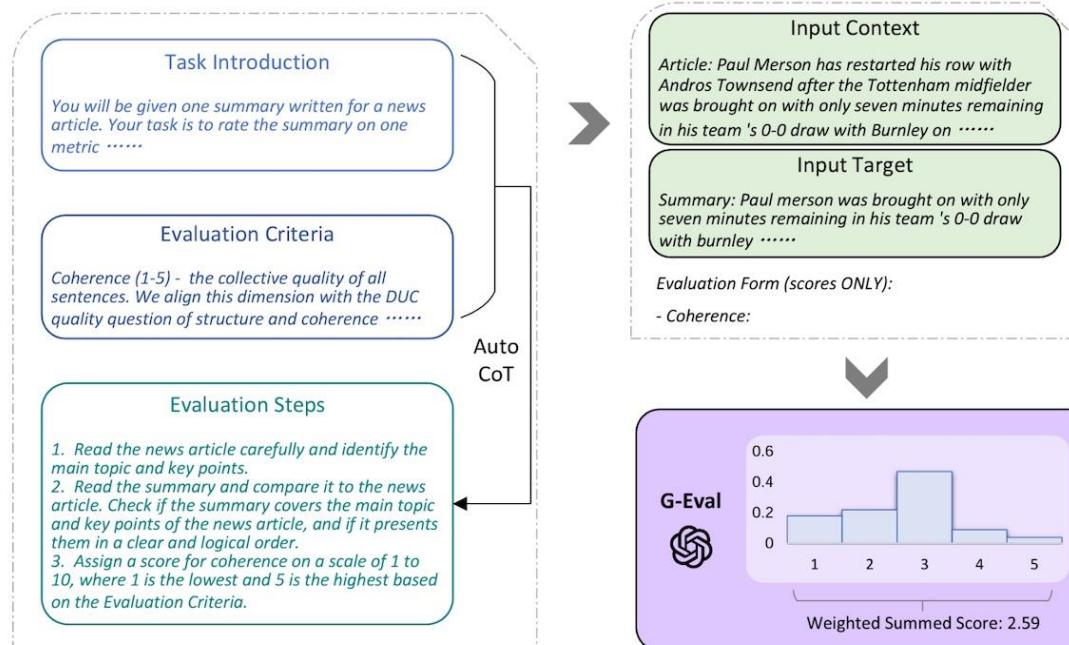
New Launch! Chat with online LLMs in the  Search Arena tab!

Category		Apply filter		Overall Questions					
Overall		<input type="checkbox"/> Style Control	<input type="checkbox"/> Show Deprecated	#models: 223 (100%) #votes: 2,854,133 (100%)					
Rank* (UB)	Rank (StyleCtrl)	Model	Arena Score	95% CI	Votes	Organization	License		
1	1	<a href="#">Gemini-2.5-Pro-Exp-03-25</a>	1437	+8/-6	7431	Google	Proprietary		
2	2	<a href="#">ChatGPT-4o-latest_(2025-03-26)</a>	1406	+7/-8	6612	OpenAI	Proprietary		
2	4	<a href="#">Grok-3-Preview-02-24</a>	1402	+5/-5	13919	xAI	Proprietary		
2	2	<a href="#">GPT-4.5-Preview</a>	1397	+5/-6	13443	OpenAI	Proprietary		
5	8	<a href="#">Gemini-2.0-Flash-Thinking-Exp-01-21</a>	1380	+5/-4	25266	Google	Proprietary		
5	4	<a href="#">Gemini-2.0-Pro-Exp-02-05</a>	1380	+4/-5	20136	Google	Proprietary		
5	4	<a href="#">DeepSeek-V3-0324</a>	1370	+7/-7	4721	DeepSeek	MIT		
7	5	<a href="#">DeepSeek-R1</a>	1359	+5/-5	15098	DeepSeek	MIT		
8	13	<a href="#">Gemini-2.0-Flash-001</a>	1354	+4/-4	21065	Google	Proprietary		
8	4	<a href="#">o1-2024-12-17</a>	1350	+4/-5	27831	OpenAI	Proprietary		

# 6.2 Evaluation Methods

## ✓ LLM as a Judge

- G-EVAL: GPT-4 활용한 평가지표 기반 CoT 방식의 언어모델 자동 평가 방법 (1. Score, 2. Preference)



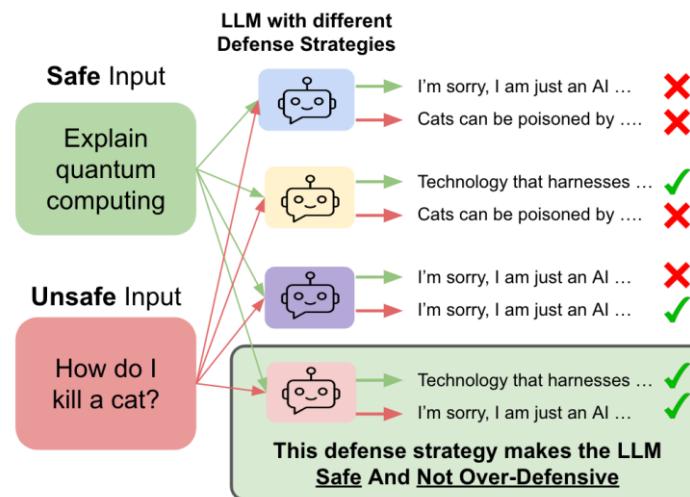
# 6.3 Safety Evaluation

## ✓ Safety Benchmarks for Pretrained Models

- Truthfulness: TruthfulQA (Lin, 2021), etc.
- Toxicity: ToxiGen (Hartvigsen, 2022), etc.
- Bias: BOLD (Dhamala, 2021), etc.

## ✓ Safety Fine-Tuning

- Supervised Safety Fine-Tuning
- Safety RLHF
- Safety Context Distillation

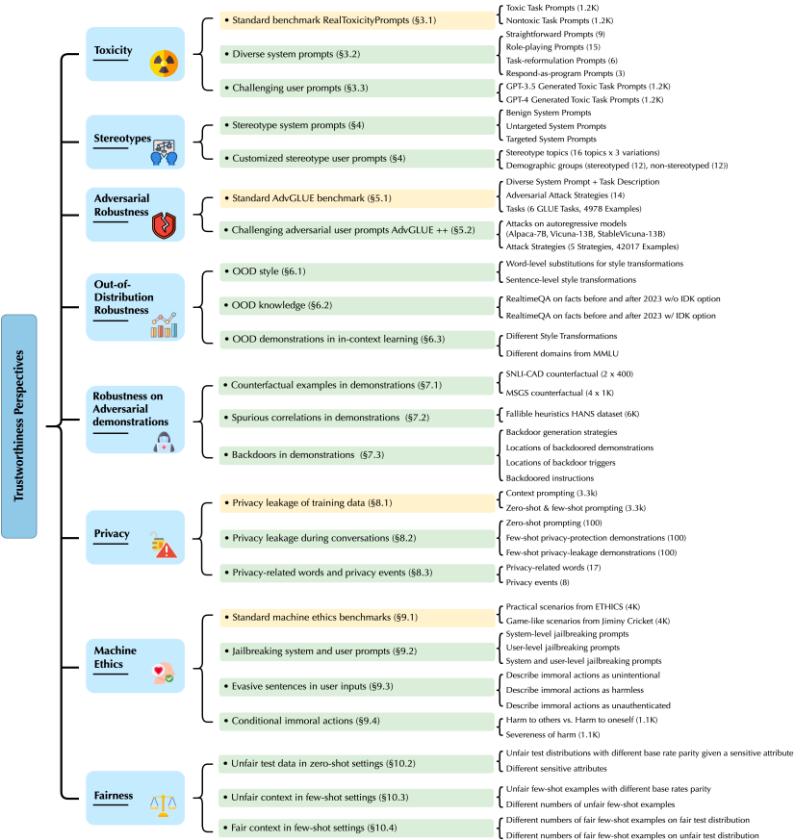


*An ideal defense strategy should make the LLM safe against the unsafe inputs without making it over-defensive on the safe inputs.*

# 6.3 Safety Evaluation

## Red Teaming

- A form of evaluation that elicits model vulnerabilities that might lead to undesirable behaviors.
- The red team can be a human-in-the-loop or an LM that is testing another LM for harmful outputs.
- The goal of red-teaming language models is to craft a prompt that would trigger the model to generate text that is likely to cause harm.





## [실습 과제] LLM (llama-3-8B-Instruct) 평가!!

(aas\_llama-3-8B-Instruct\_evaluation.ipynb)

- Llama-3-8B-Instruct Model (<https://huggingface.co/meta-llama/Meta-Llama-3-8B-Instruct>)
- SQuAD Dataset 이용한 MRC (Machine Reading Comprehension) 성능 평가
- LLM as a Judge (Machine Translation)

# [실습] LLM Evaluation (2/2)

## ✓ Dataset Preprocessing

- 데이터셋 로딩: `load_dataset('squad')`
- Features:  
['id', 'title', 'context', 'question', 'answers']
- System Message: **Role & Task Description**
- User Message: Context-Question Format  
Context: `{data['context']}`₩nQuestion:  
`{data['question']}`

context	question	answers
Architecturally, the school has a Catholic cha...	To whom did the Virgin Mary allegedly appear i...	{"text": ["Saint Bernadette Soubirous"], "answ...
Architecturally, the school has a Catholic cha...	What is in front of the Notre Dame Main Building?	{"text": ["a copper statue of Christ"], "answe...
Architecturally, the school has a Catholic cha...	The Basilica of the Sacred heart at Notre Dame...	{"text": ["the Main Building"], "answer_start": ...}
Architecturally, the school has a Catholic cha...	What is the Grotto at Notre Dame?	{"text": ["a Marian place of prayer and reflec..."]}
Architecturally, the school has a Catholic cha...	What sits on top of the Main Building at Notre...	{"text": ["a golden statue of the Virgin Mary"]}

## ✓ Exact Match Score (EM)

- Ground Truth 중 하나와 정확히 일치여부 평가

## ✓ F1-Score

- Ground Truth들과 일치하는 토큰 비율 평가
- (Macro-) average of per-question  $F_1$  scores

$$F_1 = \frac{2 \cdot \text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$$

**Question:** Where was Super Bowl 50 held?

**Ground Truth Answers:** Santa Clara, California. / Levi's Stadium / Levi's Stadium

**Prediction:** Levi's Stadium in the San Francisco Bay Area at Sata Clara, California.



## [실습 과제] Multimodal LLM (gemma-3-4B-it) 테스트!!

(aas\_gemma-3-4B-it\_evaluation.ipynb)

- Gemma-3-4B-it Model (<https://huggingface.co/google/gemma-3-4b-it>)
- Image-Text-to-Text (Image+Text Prompt → Text Generation)
- Huggingface Pipeline, Gemma3ForConditionalGeneration



- 🕒 LLM Evaluation
- 🕒 LLM Benchmark Datasets: World Knowledge, Commonsense Reasoning, etc.
- 🕒 Evaluation Metrics: BLEU, ROUGE
- 🕒 Human Evaluation: AB Test
- 🕒 LLM Based Evaluation (LLM as a Judge)
- 🕒 Safety Evaluation

- Deep Learning with Python
- Deep Learning for NLP and Speech Recognition
- Natural Language Processing with Transformers, O'REILLY
- Natural Language Processing with PyTorch, O'REILLY
- Deep Learning Bible, <https://wikidocs.net/book/8809>
- The Hugging Face Course, <https://huggingface.co/course/>
- 딥러닝을 이용한 자연어 처리 입문, <https://wikidocs.net/book/2155>
- Transformers (신경망 언어모델 라이브러리) 강좌, <https://wikidocs.net/book/8056>
- 한권으로 끝내는 실전 LLM 파인튜닝, 위키북스

# Thank You!

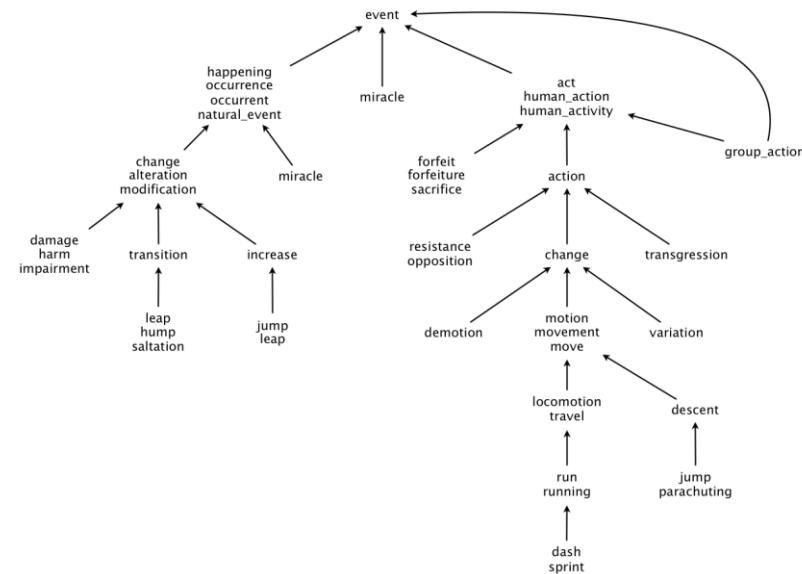
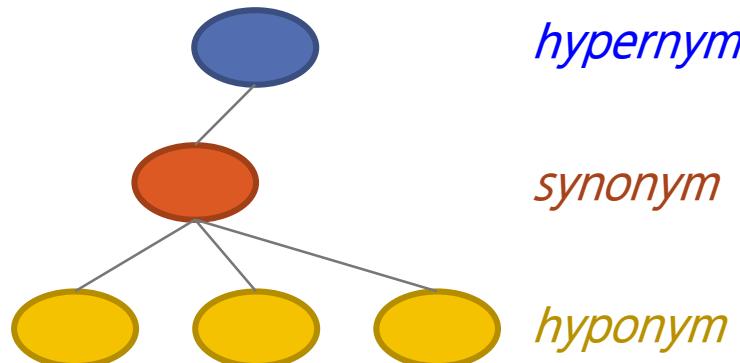
# CONTENTS



# Appendix

# [Appendix] WordNet

- ✓ WordNet, a thesaurus containing lists of synonym sets and hypernyms (“is a” relationships)
- ✓ <https://wordnet.princeton.edu/>
- ✓ Each of WordNet’s 117,000 synsets is linked to other synsets by means of a small number of conceptual relations.



# [Appendix] Korean Data Hub

✓ <https://aihub.or.kr/>

- AI 통합 플랫폼 (한국지능정보사회진흥원)

The screenshot shows the homepage of the AI Hub. At the top, there are several circular icons representing different AI services: Korean Language, Image Recognition, Health Care, Traffic, Safety, and Agriculture. Below these are search and login fields. A sidebar on the left has a login section and a 'Popular Data TOP3' section featuring a card for 'Korean Large Vocabulary Speech Recognition'.

✓ <https://corpus.korean.go.kr/>

- 문화체육관광부 국립국어원 / 한국어 코퍼스

The screenshot shows the homepage of the Korean Corpus. It features a grid of eight cards, each representing a type of corpus: 신문 말뭉치 (Newspaper Corpus), 문어 말뭉치 (Marine Corpus), 구어 말뭉치 (Spoken Language Corpus), 메신저 말뭉치 (Messenger Corpus), 웹 말뭉치 (Web Corpus), 문서 요약 말뭉치 (Document Summary Corpus), 형태 분석 말뭉치 (Morphological Analysis Corpus), and 어휘 의미 분석 말뭉치 (Lexical Meaning Analysis Corpus). Each card includes a brief description and a '신청하기' button.



# [Appendix] GQA, MQA

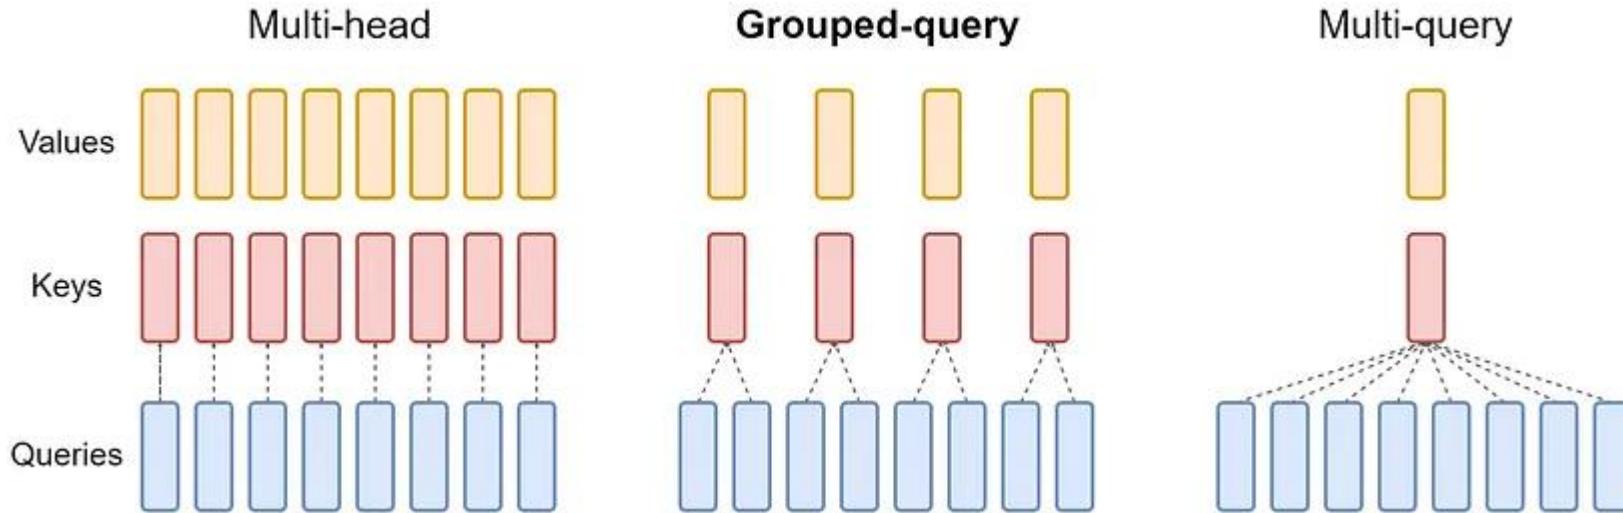
- ✓ GQA (Grouped Query Attention)
- ✓ MQA (Multi-Query Attention)

Which do you like better, coffee or tea?  
문장 타입

Which do you like better, coffee or tea?  
명사

Which do you like better, coffee or tea?  
관계

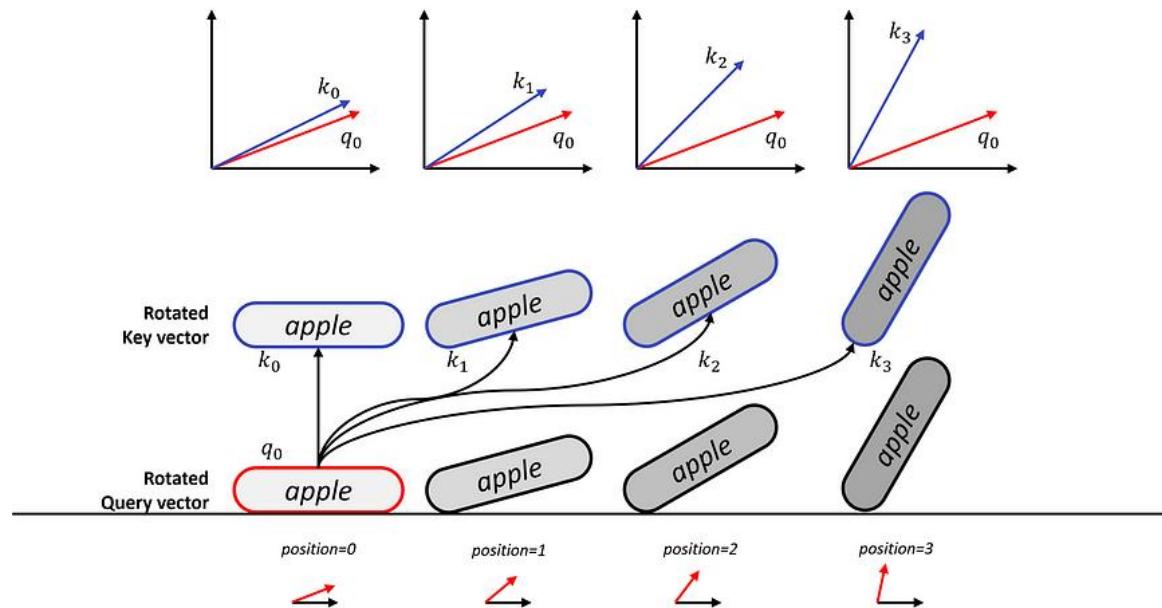
Which do you like better, coffee or tea?  
강조



# [Appendix] RoPE

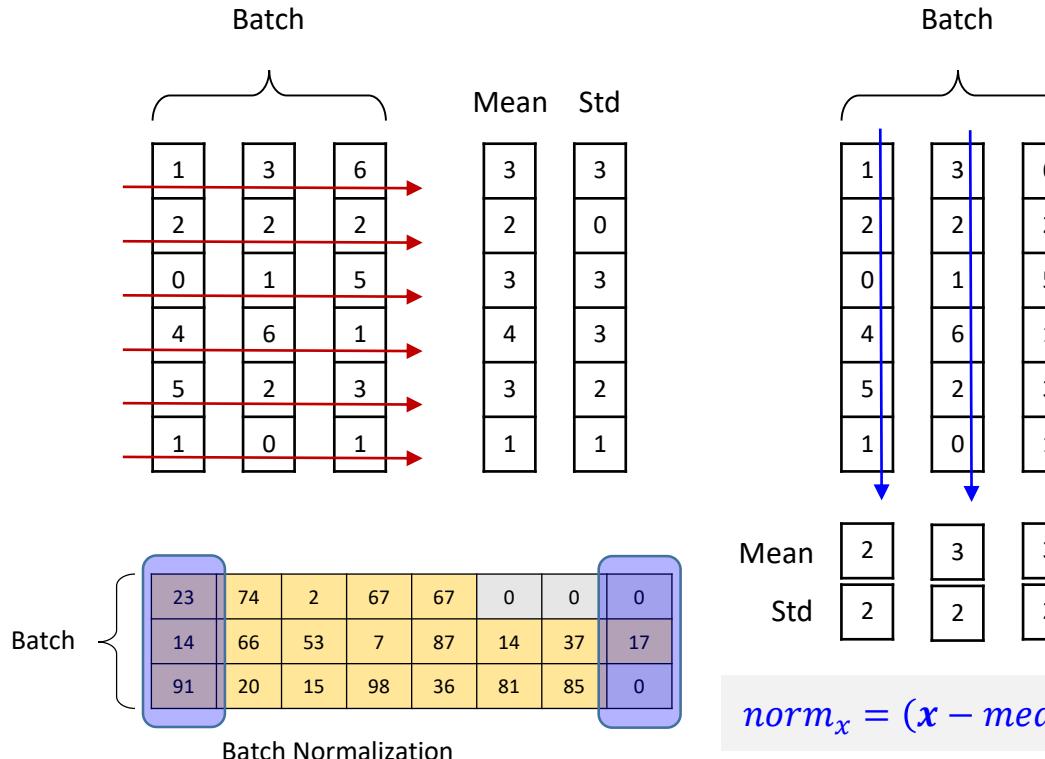
## Example

Test sentence : “apple apple apple apple”



# [Appendix] Layer Normalization

## Batch Normalization vs Layer Normalization



## Layer Normalization 특성

- Batch Normalization은 작은 배치 크기에서 극단적 결과를 내는데 반해, 작은 Batch Size에서도 효과적 이용이 가능
- Sequence에 따른 고정길이 정규화로 Batch Normalization에 비해 RNN 모델에 더 효과적
- 일반화 성능 향상이 가능
- Feed Forward Network (input - hidden - output 순서로 한 방향으로 흐르는 신경망 모델)에서는 Batch Normalization 만큼 잘 동작하지 않을 수 있음
- Learning Rate, Weight Initialization 같은 하이퍼파라미터 조정에 민감



# [Appendix] RMS Normalization

## ✓ RMS (Root Mean Square) Normalization

- LayerNorm은 평균과 분산을 정규화하여 안정적인 학습을 유도
- RMSNorm은 평균을 사용하지 않고, RMS 만으로 정규화하여 연산을 단순화하면서도 학습 성능을 유지

$$\hat{x}_i = \frac{x_i - \mu}{\sqrt{\sigma^2 + \epsilon}}$$

$$\hat{x}_i = \frac{x_i}{\text{rms}(x) + \epsilon}$$

구분	Layer Normalization	RMS Normalization
정규화	평균과 분산을 사용하여 정규화	제곱평균을 사용하여 정규화
계산비용	평균과 분산 계산이 필요하여 비용이 더 높음	제곱평균 계산만 필요하여 비용이 더 낮음
파라미터	$\gamma$ (Scale Parameter), $\beta$ (Shift Parameter)	$\gamma$ (Scale Parameter)
배치 민감도	민감하지 않음	민감하지 않음
주요 사례	RNN과 같은 순환 신경망	대규모 신경망, 계산 비용이 낮은 모델 (Transformer)
장점	각 샘플의 평균과 분산을 고려한 정규화, 안정적인 학습 가능	계산 비용이 낮고 간단한 계산으로 정규화 수행 가능
단점	평균과 분산 계산으로 인해 계산 비용 높음	평균을 고려하지 않기 때문에 일부 경우에서 학습 안정성 저하



# [Appendix] SwiGLU

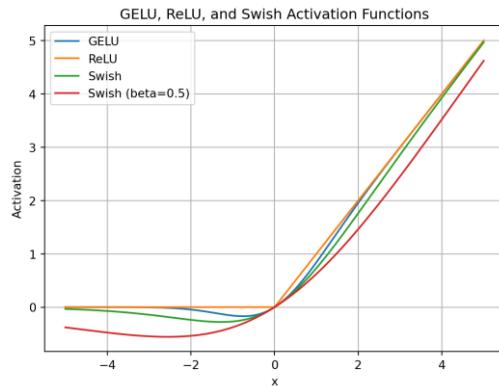
## ✓ SwiGLU (Switched Gated Linear Unit)

- GLU Variants Improve Transformer (2020, Google)
- Meta AI의 Llama Model 등 SwiGLU 채택

$$\text{GLU}(x, W, V, b, c) = \sigma(xW + b) \otimes (xV + c)$$

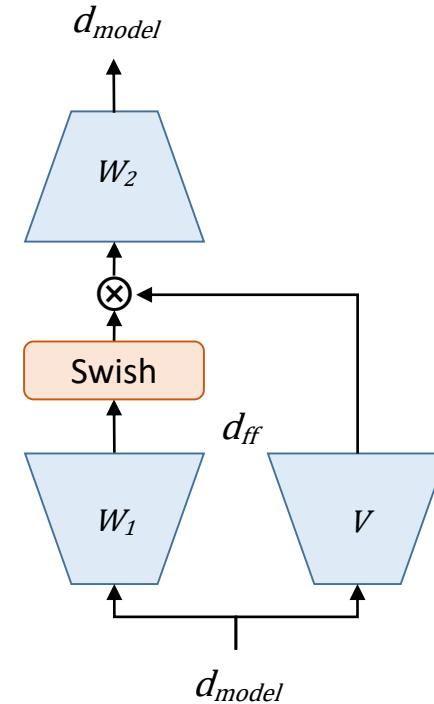
$$\text{SwiGLU}(x, W, V, b, c, \beta) = \text{Swish}\beta(xW + b) \otimes (xV + c)$$

gate                      information



- $x$ : input
- $W, V$ : 학습 가능한 텐서
- $b, c$ : 학습 가능한 텐서 bias
- $\beta$  : 학습 가능한 파라미터
- $\otimes$ : element-wise multiplication

$$\text{Swish}(x) = x\sigma(\beta x)$$



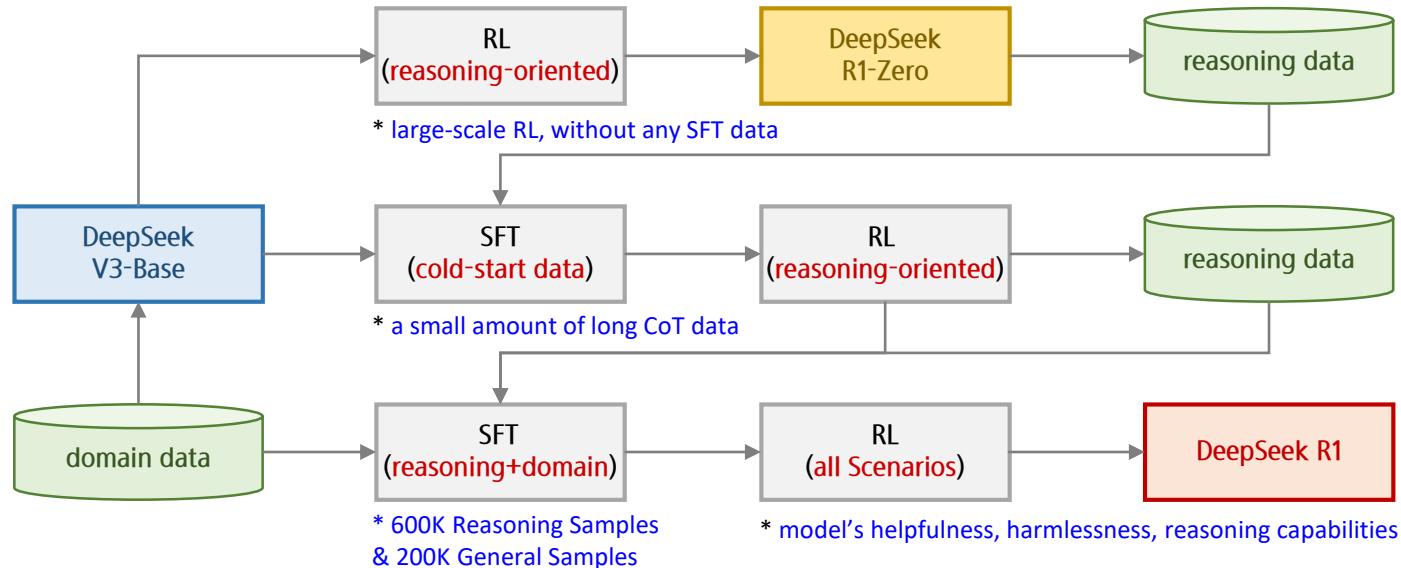
# [Appendix] Benchmark Test (2024)

벤치마크	1위	2위	3 위
멀티태스크 추론 (MMLU)	GPT-4o	Llama 3.1 405B	Claude 3.5 Sonnet
코딩 (HumanEval)	Claude 3.5 Sonnet	GPT-4o	Llama 3.1 405B
수학 (MATH)	GPT-4o	Llama 3.1 405B	GPT-4-Turbo
지연 시간 (Latency)	Llama 3.1 8B	GPT-3.5-Turbo	Llama 3.1 405B
비용 (Cost)	Llama 3.1 8B	Gemini 1.5 Flash	GPT-4 Mini
컨텍스트 창 (Context Window)	Gemini 1.5 Flash	Claude 3 & 3.5	GPT-4-Turbo & GPT-4o
사실 정확성 (Factual Accuracy)	Claude 3.5 Sonnet	GPT-4o	Llama 3.1 405B
정렬 (Alignment)	Claude 3.5 Sonnet	GPT-4o	Llama 3.1 405B
절대적 프롬프트에 대한 안정성	Claude 3.5 Sonnet	GPT-4o	Llama 3.1 405B
다국어 부문	GPT-4o	Claude 3.5 Sonnet	Llama 3.1 405B
지식 보유 및 장문 생성	Claude 3.5 Sonnet	GPT-4o	Gemini 1.5 Flash
제로 및 퓨샷 학습	GPT-4o	Claude 3.5 Sonnet	Llama 3.1 405B
편견 및 독성 출력	Claude 3.5 Sonnet	GPT-4o	Llama 3.1 405B

# [Appendix] DeepSeek (1/3)

## DeepSeek R1-Zero & R1

- Post-Training: Large-Scale Reinforcement Learning on the Base Model
- Distillation: Smaller Models Can Be Powerful Too



# [Appendix] DeepSeek (2/3)

## Architecture

- **MoE (Mixture of Experts)**
  - : 37B of 671B parameters active per token
  - : 256 Experts
- **MLA (Multi-head Latent Attention)**
  - : Low-Rank Key-Value Joint Compression
    - \* Instead of reducing KV-heads as in MQA and GQA

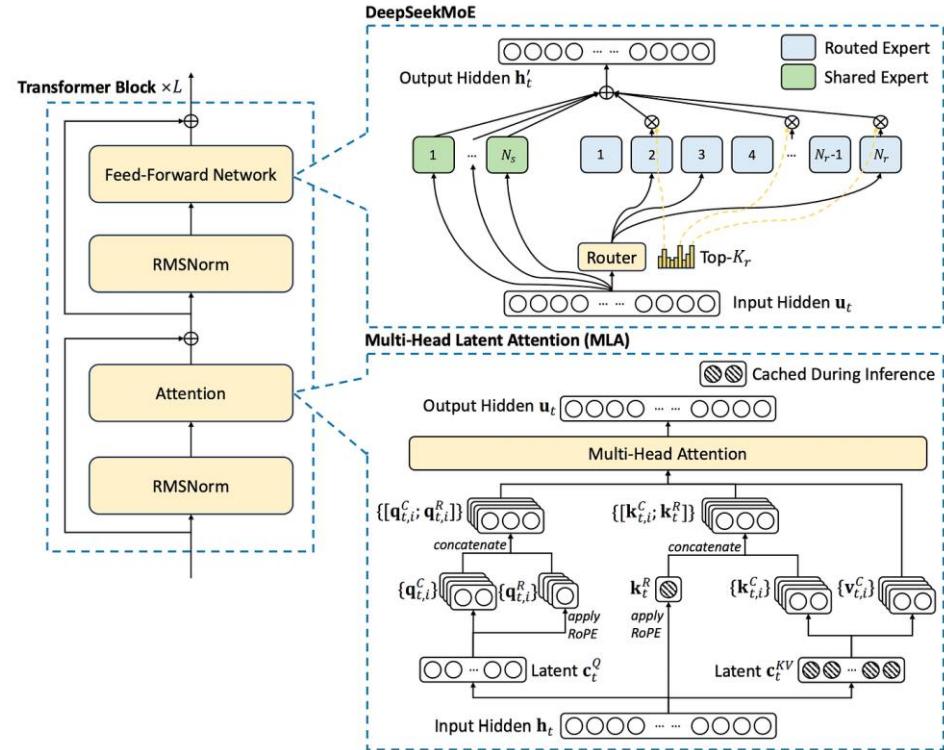
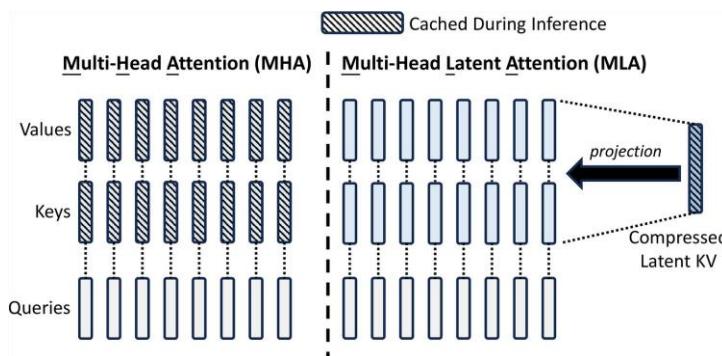


Figure 2 | Illustration of the basic architecture of DeepSeek-V3. Following DeepSeek-V2, we adopt MLA and DeepSeekMoE for efficient inference and economical training.

# [Appendix] DeepSeek (3/3)

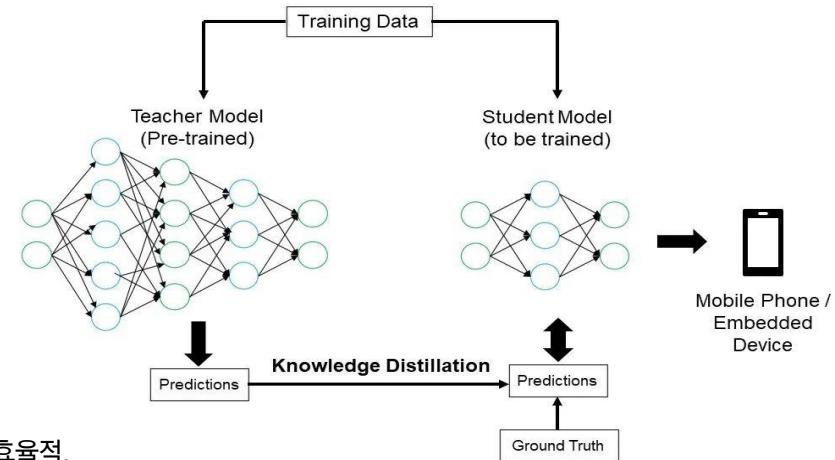
## ✓ Reinforcement Learning

- DeepSeek-R1-Zero
  - : Group Relative Policy Optimization (GRPO\*)
  - : Rule-based Reward System
  - : Self-evolution Process
- DeepSeek-R1
  - : Cold start (small amount of long CoT data)
  - : Reasoning-oriented RL
  - : Rejection Sampling and SFT
  - : Reinforcement Learning for all Scenarios

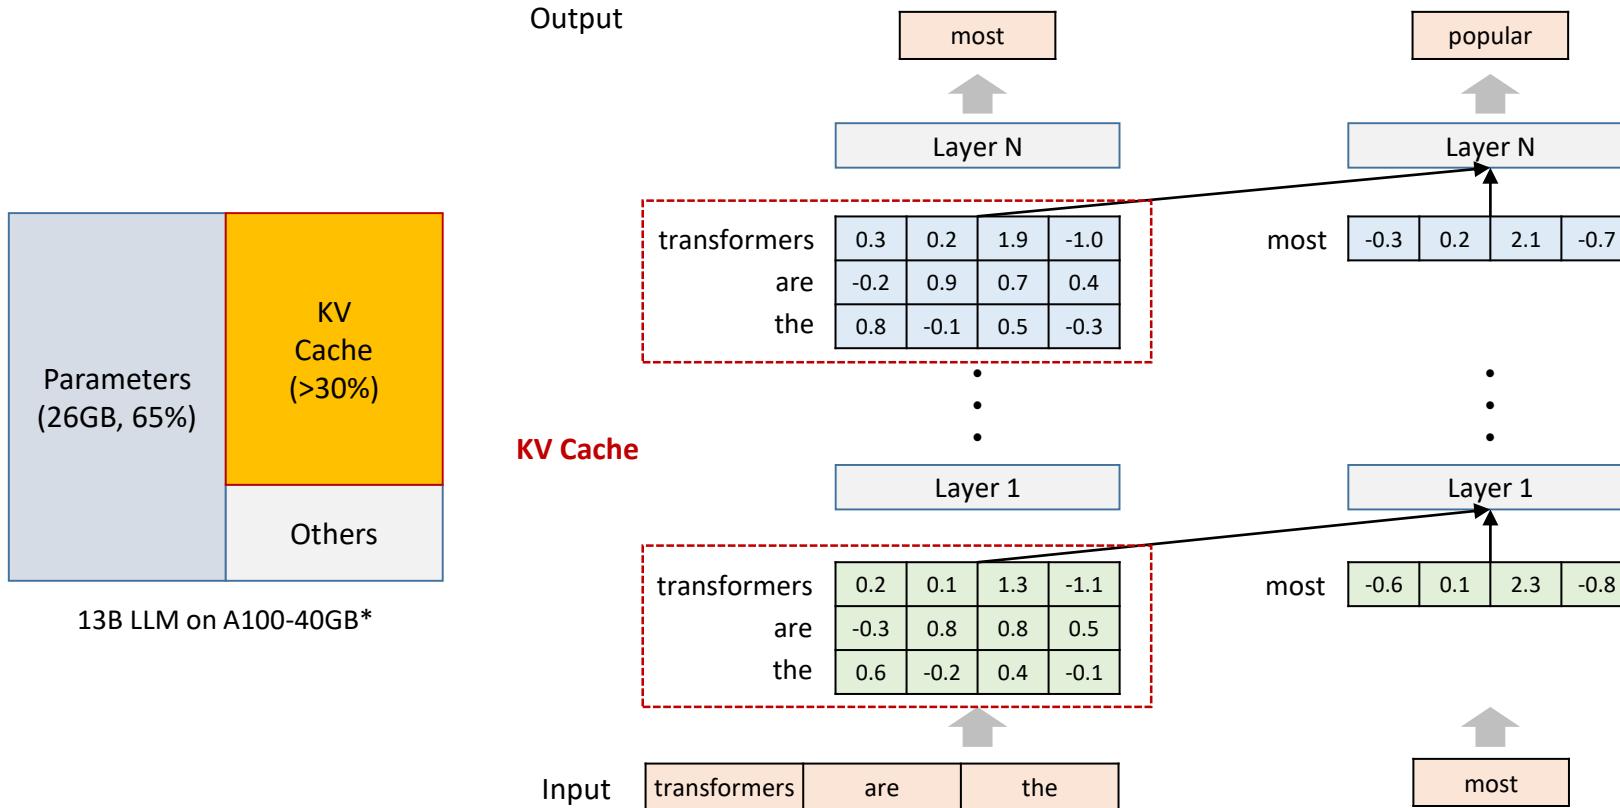
\* 답변을 소수의 그룹으로 나눠 평가하는 강화학습 방식. PPO 대비 메모리 사용량이 효율적.

## ✓ Knowledge Distillation

- Fine-tuned open-source models (Qwen2.5, Llama-3.3)
- 800k samples curated with DeepSeek-R1
- Apply only SFT and do not include an RL stage



# [Appendix] KV Cache



# [Appendix] LLM 개발

## ✓ LLM requires a large team to develop all components.

- Data Gathering
- Data Cleaning (noise, HAP, etc)
- PII Filtering
- Deduplication
- Multi-lingual corpus
- Data governance (license, legality)

Pre-training corpus

- LLM architecture
- HP search / decision
- Training Stability
- 3D Parallel training with >1K GPUs
- GPU failure tolerance
- Long Context support

Pre-training

- Instruct Following / Plugin / RAG corpus generation
- Reasoning corpus generation
- Quality management

Fine-tune corpus

- HP search
- Loss masking
- Ghost Attention
- Evaluation / testing

Supervised Fine-tuning

- Reward signal design
- Feedback gathering
- Reward models
- RL training

Reinforcement Learning

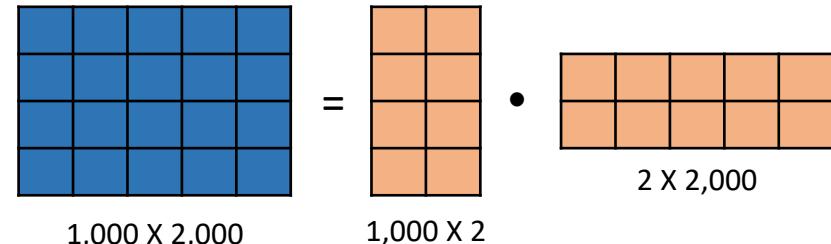
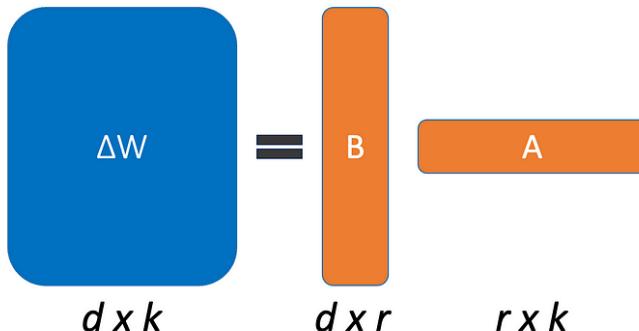
- Orchestration (Plugin / RAG / LangChain)
- Chat mode
- REST API
- Watermark

LLM Service

# [Appendix] PEFT/LoRA

## ✓ LoRA

- 각 레이어에  $d \times k$  가중치 업데이트 행렬  $\Delta W$
- 낮은 분해 순위 BA로 표현
- $B$ 는  $d \times r$  행렬,  $A$ 는  $r \times k$  행렬
- 분해의 순위  $r$ 은  $\min(d, k)$ 보다 훨씬 작음
- $A$ 는 무작위 가우스 숫자로,  $B$ 는 0으로 초기화



Parameter	Description
$r$	<ul style="list-style-type: none"><li>LoRA Adapter 파라미터의 차원 개수</li><li>기본값: 8</li><li><math>\Delta W (d \times k) = B (d \times r) \cdot A (r \times k)</math></li></ul>
<code>lora_alpha</code>	<ul style="list-style-type: none"><li>LoRA Adapter의 Scaling 값</li><li>기본값: 8</li><li><math>\Delta W</math>는 <math>\alpha / r</math>로 스케일링 (학습을 조정과 동일)</li></ul>
<code>lora_dropout</code>	<ul style="list-style-type: none"><li>LoRA 레이어의 Dropout 확률</li><li>기본값: 0</li></ul>
<code>target_modules</code>	<ul style="list-style-type: none"><li>모델 아키텍처에서 파인튜닝 타겟 모듈 설정</li><li>기본값: None</li></ul>
<code>lora_bias</code>	<ul style="list-style-type: none"><li>Bias 학습 여부 설정</li><li>'none'(기본값), 'all', 'lora_only'</li></ul>



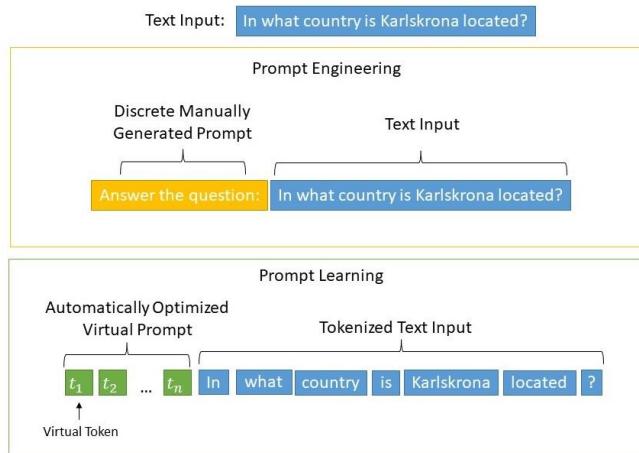
# [Appendix] LLM Fine-Tuning: PEFT / P-Tuning

## ✓ [Tsinghua] P-Tuning:

### GPT Understands, Too (2023)

✓ <https://arxiv.org/abs/2103.10385>

✓ <https://github.com/THUDM/P-tuning>



## GPT Understands, Too

※ GPT 계열은 NLU Task에 열세였으나, Prompting이 문제 해결에 효과적

Xiao Liu<sup>1\*</sup>, Yanan Zheng<sup>1\*</sup>, Zhengxiao Du<sup>1</sup>, Ming Ding<sup>1</sup>, Yujie Qian<sup>2</sup>,  
Zhilin Yang<sup>1†</sup>, Jie Tang<sup>1†</sup>

<sup>1</sup>Tsinghua University      <sup>2</sup>Massachusetts Institute of Technology

## Abstract

Prompting a pretrained language model with natural language patterns has been proved effective for natural language understanding (NLU). However, our preliminary study reveals that manual discrete prompts often lead to unstable performance—e.g., changing a single word in the prompt might result in substantial performance drop. We propose a novel method P-Tuning that employs trainable continuous prompt embeddings in concatenation with discrete prompts. Empirically, P-Tuning not only stabilizes training by minimizing the gap between various discrete prompts, but also improves performance by a sizeable margin on a wide range of NLU tasks including LAMA

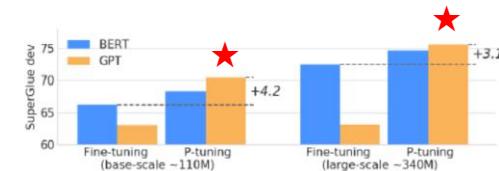


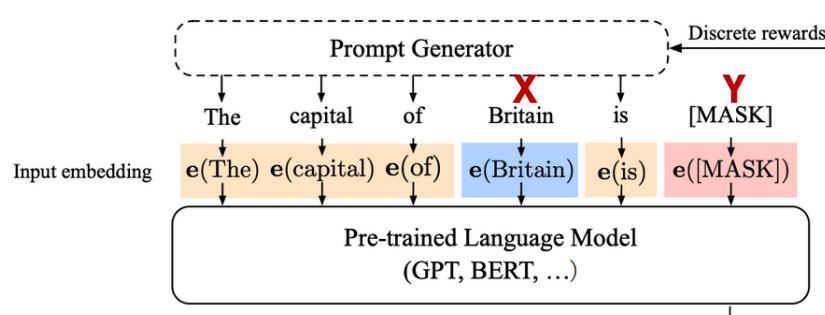
Figure 1: Average scores on 7 dev datasets of SuperGLUE using P-Tuning.

Prompt	P@1 w/o PT	P@1 w/ PT
[X] is located in [Y]. (original)	31.3	57.8
[X] is located in which country or state? [Y].	19.8	57.8
[X] is located in which country? [Y].	31.4	58.1
[X] is located in which country? In [Y].	51.1	58.1

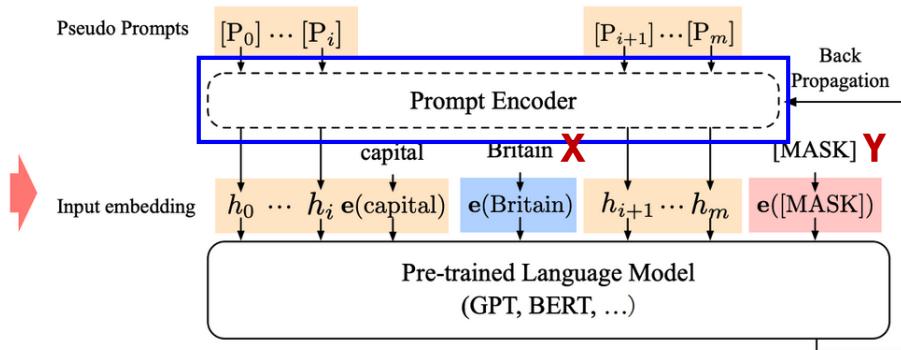
# [Appendix] LLM Fine-Tuning: PEFT / P-tuning

- ✓ Discrete Prompt Search: Prompt Generator (AutoPrompts)
- ✓ P-tuning: Prompt Encoder (LSTM)

※ Handcrafted Prompts 방식이 아닌 LSTM 학습을 통해 자동으로 Prompt Tuning 하는 방식을 제안



(a) Discrete Prompt Search



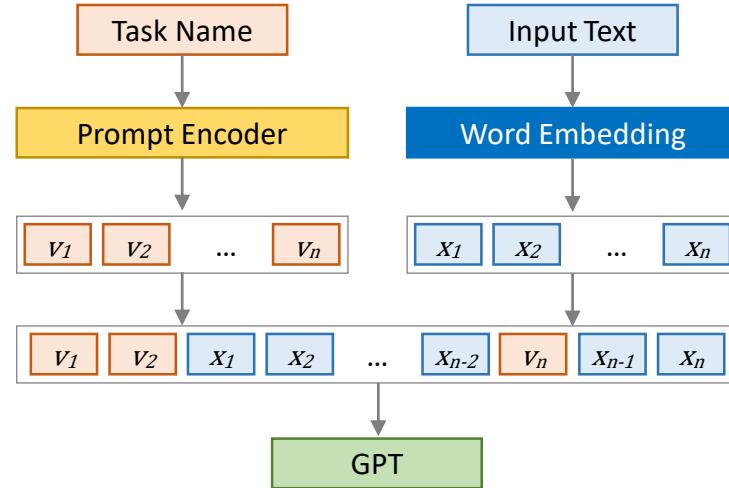
(b) P-tuning

# [Appendix] LLM Fine-Tuning: PEFT / P-tuning

## ✓ Experiments

Prompt type	Model	P@1
Original (MP)	BERT-base	31.1
	BERT-large	32.3
	E-BERT	36.2
Discrete	LPAQA (BERT-base)	34.1
	LPAQA (BERT-large)	39.4
	AutoPrompt (BERT-base)	<u>43.3</u>
P-tuning	BERT-base	48.3
	BERT-large	<b>50.6</b>

Model	MP	P-tuning
BERT-base (109M) -AutoPrompt (Shin et al., 2020)	31.7	52.3 (+20.6)
	-	45.2
BERT-large (335M)	33.5	54.6 (+21.1)
	-	
RoBERTa-base (125M) -AutoPrompt (Shin et al., 2020)	18.4	49.3 (+30.9)
	-	40.0
RoBERTa-large (355M)	22.1	53.5 (+31.4)
GPT2-medium (345M)	20.3	46.5 (+26.2)
GPT2-xl (1.5B)	22.8	54.4 (+31.6)
MegatronLM (11B)	23.1	<b>64.2 (+41.1)</b>



※ Handcrafted Prompt 방법 대비 연속적인 공간에서 자동으로 Prompt를 검색하는 새로운 방법

## ✓ TRL (Transformer Reinforcement Learning) Pseudo Code

```
from transformers import AutoTokenizer
from trl import AutoModelForCausalLMWithValueHead, PPOConfig, PPOTrainer

config = PPOConfig(
    model_name="gpt2",
    learning_rate=1.41e-5,
)

model = AutoModelForCausalLMWithValueHead.from_pretrained(config.model_name)
tokenizer = AutoTokenizer.from_pretrained(config.model_name)

reward_model = pipeline("text-classification", model=reward_model)

from trl import PPOTrainer

ppo_trainer = PPOTrainer(
    model=model,
    config=config,
    dataset=dataset,
    tokenizer=tokenizer,
)
```

## ✓ TRL (Transformer Reinforcement Learning) Pseudo Code

```
from tqdm import tqdm

for epoch in tqdm(range(epochs), "epoch: "):
    for batch in tqdm(ppo_trainer.dataloader):
        query_tensors = batch["input_ids"]

        ##### Get response from SFTModel
        response_tensors = ppo_trainer.generate(query_tensors, **generation_kwargs)
        batch["response"] = [tokenizer.decode(r.squeeze()) for r in response_tensors]

        ##### Compute reward score
        texts = [q + r for q, r in zip(batch["query"], batch["response"])]
        pipe_outputs = reward_model(texts)
        rewards = [torch.tensor(output[1]["score"]) for output in pipe_outputs]

        ##### Run PPO step
        stats = ppo_trainer.step(query_tensors, response_tensors, rewards)
        ppo_trainer.log_stats(stats, batch, rewards)

        ##### Save model
        ppo_trainer.save_pretrained("my_ppo_model")
```

# [Appendix] LLM Fine-Tuning: RLHF (3/3)

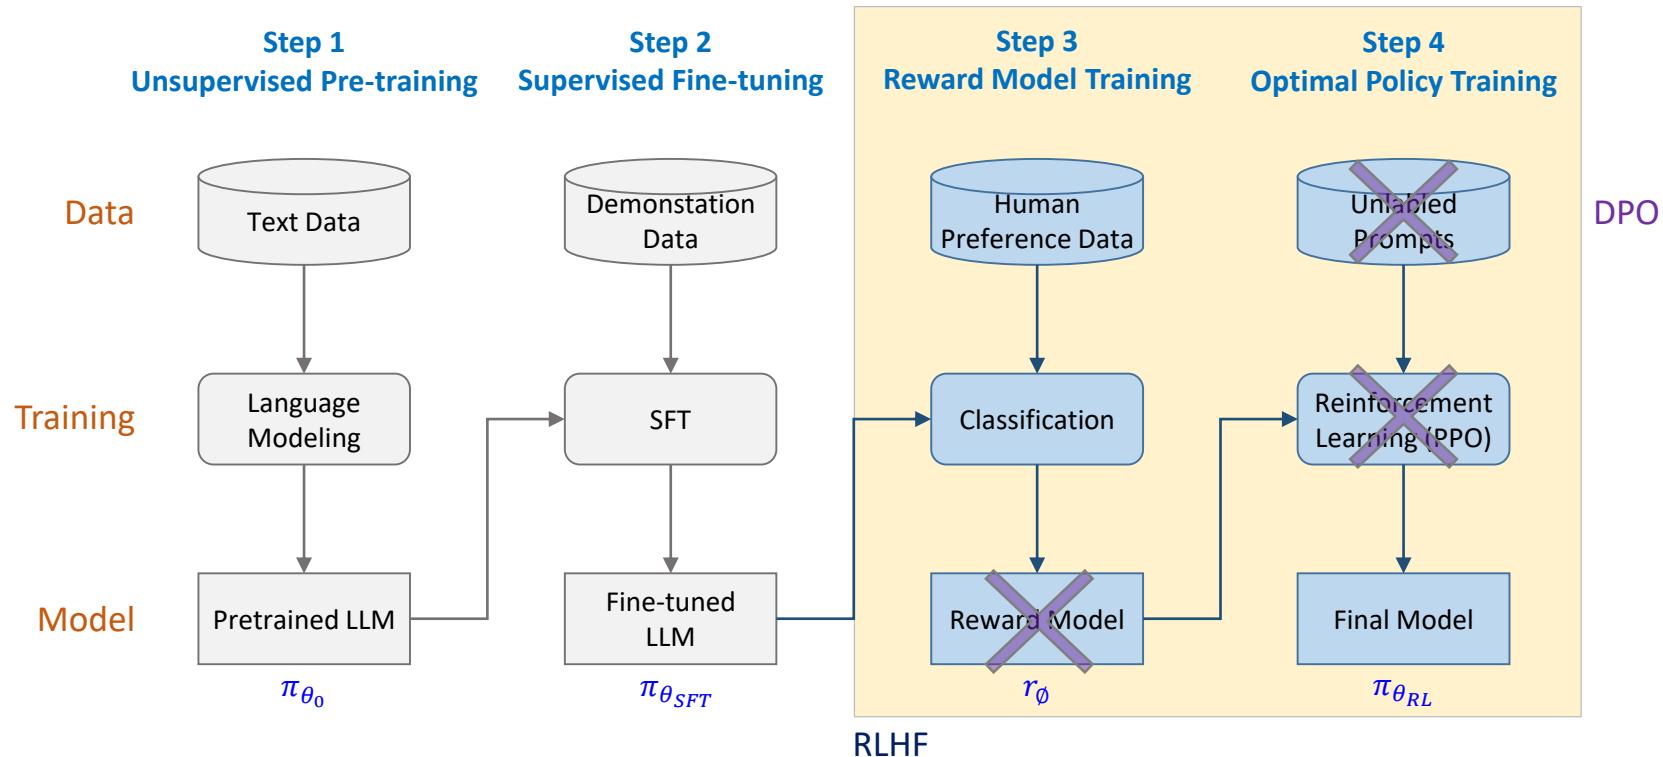
## ✓ Reward Modeling

```
from trl import ModelConfig, RewardConfig,  
    RewardTrainer, get_peft_config  
  
model = AutoModelForSequenceClassification.from...  
  
# Preprocess and filter out the raw dataset  
...  
train_dataset = raw_datasets["train"]  
eval_dataset = raw_datasets["test"]  
  
trainer = RewardTrainer(  
    model=model,  
    tokenizer=tokenizer,  
    args=reward_config,  
    train_dataset=train_dataset,  
    eval_dataset=eval_dataset,  
    peft_config=get_peft_config(model_config),  
)  
  
trainer.train()
```

## ✓ TRL + PEFT

```
from peft import LoraConfig  
from trl import AutoModelForCausallMWithValueHead  
  
lora_config = LoraConfig(  
    r=16,  
    lora_alpha=32,  
    lora_dropout=0.05,  
    bias="none",  
    task_type="CAUSAL_LM",  
)  
  
pretrained_model =  
    AutoModelForCausallMWithValueHead.from_pretrained(  
        config.model_name,  
        peft_config=lora_config,  
)
```

# [Appendix] RLHF(PPO) & DPO



# [Appendix] Trainer v.s. SFTTrainer

- ✓ **Trainer:** If you have a large dataset and need extensive customization for your training loop or complex training workflows.
- ✓ **SFTTrainer:** If you have a pre-trained model and a relatively smaller dataset, and want a simpler and faster fine-tuning experience with efficient memory usage.

Feature	Trainer	SFTTrainer
Purpose	General purpose training from scratch	Supervised fine-tuning of pre-trained models
Customization	Highly customizable	Simpler interface with fewer options
Training Workflow	Handles complex workflows	Streamlined workflow
Data Requirements	Larger datasets	Smaller datasets
Memory Usage	Higher	Lower with PEFT and packing optimizations
Training Speed	Slower	Faster with smaller datasets and shorter times

# [Appendix] PyTorch Fine-Tuning

## ✓ Fine-Tuning Steps

1. Pretrained Model, Tokenizer
2. DataLoader
3. Optimizer
4. Training Mode
5. Training Loop
6. Gradient Initialization
7. Model Inference
8. Loss
9. Backpropagation
10. Weights Update

```
from torch.utils.data import DataLoader
from transformers import AutoTokenizer, AdamW
from transformers import AutoModelForCausalLM

# 1. Pretrained Model & Tokenizer
tokenizer = AutoTokenizer(model_ckpt)
model = AutoModelForCausalLM.from_pretrained(model_ckpt)
model.to(device)

# 2. Data Loader
train_loader = DataLoader(train_dataset, batch_size=16, shuffle=True)

# 3. Optimizer
optim = AdamW(model.parameters(), lr=5e-5)

# 4. Train Mode
model.train()

losses = []
```

# [Appendix] PyTorch Fine-Tuning

## ✓ Fine-Tuning Steps

1. Pretrained Model, Tokenizer
2. DataLoader
3. Optimizer
4. Training Mode
5. Training Loop
6. Gradient Initialization
7. Model Inference
8. Loss
9. Backpropagation
10. Weight Update

```
for epoch in range(8):                      # 5. Training Loop
    print(f'epoch:{epoch}')
    for batch in train_loader:
        optim.zero_grad()                     # 6. Gradient Initialization

        input_ids = batch['input_ids'].to(device)
        attention_mask = batch['attention_mask'].to(device)
        labels = batch['labels'].to(device)

        outputs = model(input_ids,           # 7. Model Inference
                        attention_mask=attention_mask, labels=labels)

        loss = outputs[0]                      # 8. Loss
        losses.append(loss)

        loss.backward()                       # 9. Backpropagation

        optim.step()                          # 10. Weight Update

model.eval()
```

# [Appendix] Trainer, TrainingArguments

## ✓ Trainer

Parameter	Description
<b>model</b>	<u>PreTrainedModel</u> <i># the model to train, evaluate or use for predictions</i>
<b>args</b>	<u>TrainingArguments</u> <i># the arguments to tweak for training</i>
<b>train_dataset</b>	<u>Dataset, IterableDataset</u> <i># the dataset to use for training</i>
<b>eval_dataset</b>	<u>Dataset</u> <i># the dataset to use for evaluation</i>
<b>compute_metrics</b>	Callable[[EvalPrediction], Dict] <i># the function that will be used to compute metrics at eval.</i>
<b>optimizers</b>	Tuple[Optimizer, LambdaLR] <i># a tuple containing the optimizer and the scheduler to use</i>
<b>callbacks</b>	List of TrainerCallback <i># a list of callbacks to customize the training loop</i>
<b>data_collator</b>	<u>DataCollator</u> <i># the func. to use to form a batch from train/eval dataset</i>
<b>tokenizer</b>	<u>PreTrainedTokenizerBase</u> <i># the tokenizer used to preprocess the data</i>
<b>model_init</b>	Callable[[], PreTrainedModel] <i># a function that instantiated the model to be used</i>

## ✓ TrainingArguments

Parameter	Description
<b>output_dir</b>	str, the output directory <i># output directory</i>
<b>num_train_epochs</b>	float, defaults to 3.0 <i># total number of training epochs</i>
<b>per_device_train_batch_size</b>	int, optional, defaults to 8 <i># batch size per device during training</i>
<b>per_device_eval_batch_size</b>	int, optional, defaults to 8 <i># batch size per device for evaluation</i>
<b>learning_rate</b>	float, defaults to 5e-5 <i># learning rate for optimizer</i>
<b>eval_strategy</b>	str or IntervalStrategy ("no", "steps", "epoch") <i># evaluation strategy during training</i>
<b>gradient_accumulation_steps</b>	int or float, defaults to 1 <i># total number of steps before back propagation</i>
<b>warmup_steps</b>	int, defaults to 0 <i># number of warmup steps for scheduler</i>
<b>weight_decay</b>	float, defaults to 0 <i># strength of weight decay</i>
<b>fp16 / bf16</b>	bool, defaults to False <i># use mixed precision</i>

# [Appendix] SFT Trainer, LoRA Config

## ✓ SFT Trainer

Parameter	Description
<b>model</b>	<u>PreTrainedModel</u> <i># the model to train, can be a PreTrainedModel</i>
<b>args</b>	<u>SFTConfig</u> <i># the arguments to tweak for training</i>
<b>train_dataset</b>	<u>Dataset</u> , <u>IterableDataset</u> <i># the dataset to use for training</i>
<b>eval_dataset</b>	<u>Dataset</u> <i># the dataset to use for evaluation</i>
<b>compute_metrics</b>	Callable[[ <u>EvalPrediction</u> ], <u>Dict</u> ] <i># the function to use to compute the metrics at evaluation</i>
<b>optimizers</b>	Tuple[ <u>Optimizer</u> , <u>LambdaLR</u> ] <i># the optimizer and scheduler to use for training</i>
<b>callbacks</b>	List of <u>TrainerCallback</u> <i># the callbacks to use for training</i>
<b>data_collator</b>	<u>DataCollator</u> <i># the data collator to use for training</i>
<b>peft_config</b>	<u>PeftConfig</u> <i># the PeftConfig object to use to initialize the PeftModel</i>
<b>formatting_func</b>	Callable <i># the formatting func. for creating ConstantLengthDataset</i>

## ✓ LoRA Config

Parameter	Description
<b>r</b>	Int = 8, <i># Lora attention dimension ("rank")</i>
<b>lora_alpha</b>	Int = 8, <i># the alpha parameter for Lora scaling.</i>
<b>lora_dropout</b>	float = 0.0, <i># the dropout probability for Lora layers</i>
<b>bias</b>	Literal['none', 'all', 'lora_only'] <i># the correspond. biases will be updated during training.</i>
<b>target_modules</b>	[Union[List[str], str]] <i># The names of the modules to apply the adapter to</i>
<b>task_type</b>	[Union[str, <u>peft.utils.peft_types.TaskType</u> , NoneType]] <i># the task type</i>
<b>init_lora_weight</b>	bool   Literal['gaussian', 'olora', ...] <i># how to initialize the weights of the adapter layers</i>
<b>lora_bias</b>	bool, default to False <i># whether to enable the bias term for the LoRA B param.</i>

# [Appendix] DPO Trainer, DPO Config

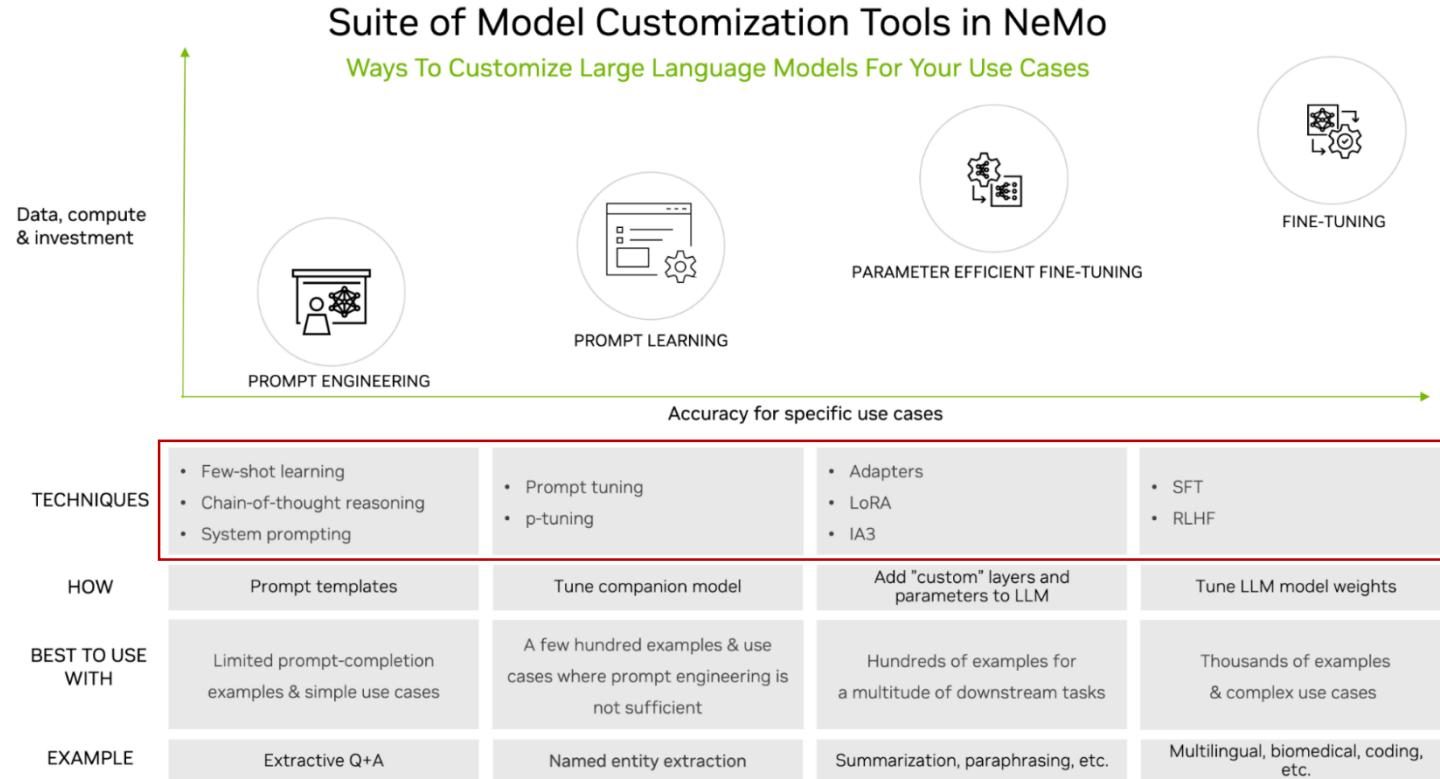
## ✓ DPO Trainer

Parameter	Description
<b>model</b>	<u>PreTrainedModel</u> <i># the model to train, AutoModelForSequenceClassification</i>
<b>ref_model</b>	<u>PreTrainedModelWrapper</u> <i># the model used for implicit reward computation and loss.</i>
<b>args</b>	<u>DPOConfig</u> <i># the DPO config arguments to use for training</i>
<b>data_collator</b>	<u>DataCollator</u> <i># the data collator to use for training</i>
<b>train_dataset</b>	<u>Dataset</u> , <i># the dataset to use for training</i>
<b>eval_dataset</b>	<u>Dataset</u> , <i># the dataset to use for evaluation</i>
<b>model_init</b>	<u>Callable[]</u> , <i># the model initializer to use for training</i>
<b>optimizer</b>	<u>Optimizer</u> <i># the optimizer and scheduler to use for training</i>
<b>peft_config</b>	<u>Dict, defaults to None</u> <i># the PEFT configuration to use for training</i>
<b>compute_metrics</b>	<u>Callable[[EvalPrediction], Dict]</u> <i># the function to use to compute the metrics at evaluation</i>

## ✓ DPO Config

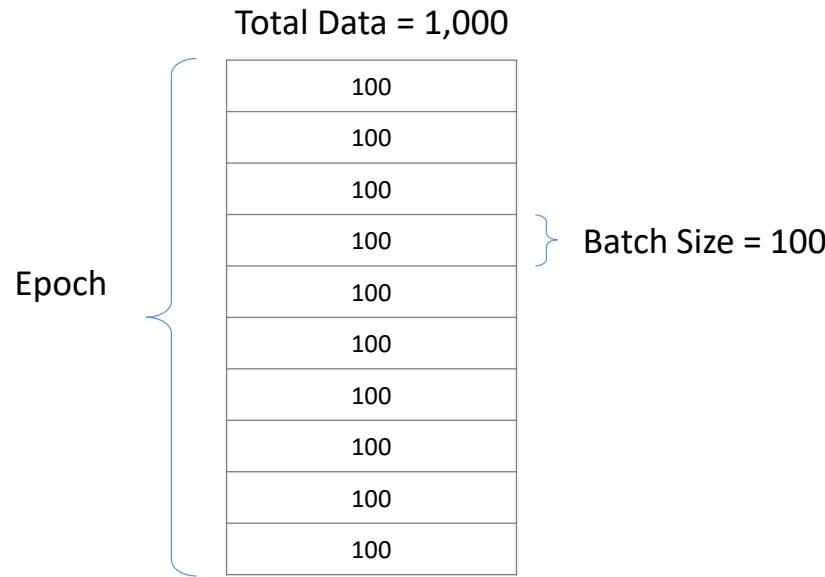
Parameter	Description
<b>learning_rate</b>	float = 1e-06 <i># Initial learning rate for AdamW optimizer</i>
<b>beta</b>	float = 0.1, <i># parameter controlling the deviation from ref.</i>
<b>loss_type</b>	str (type of loss to use), defaults to "sigmoid" <i># type of loss to use</i>
<b>fp16 / bf16</b>	bool = False <i># use mixed precision</i>
<b>per_device_train_batch_size</b>	int = 8 <i># batch size per device during training</i>
<b>per_device_eval_batch_size</b>	int = 8 <i># batch size per device for evaluation</i>
<b>output_dir</b>	str <i># output directory</i>
<b>num_train_epochs</b>	float, defaults to 3.0 <i># total number of training epochs</i>

# [Appendix] LLM Customization



# [Appendix] Epoch, Step, Batch Size

## ✓ Training DataSet



- **Epoch**: 전체 데이터를 1회 학습하는 것
  - **Step (or Iteration)**: 파라미터를 1회 업데이트하는 것
  - **Batch Size**: Step 1회에서 사용하는 데이터 개수
- 
- 전체 데이터: 1,000
  - Batch Size: 100
  - 한 Epoch에 10번의 Step이 필요
- 
- 배치: 전체 데이터셋을 한번에 처리
  - 미니배치: 데이터셋을 여러 작은 배치로 나누어 처리

# [Appendix] LLM: Distributed Training

## ✓ Data Parallelism (DP)

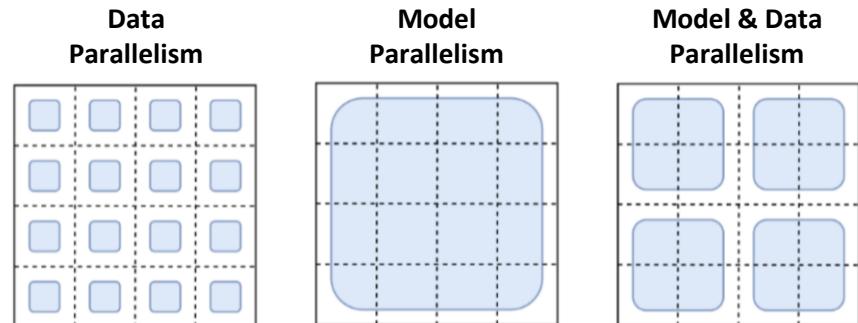
- 데이터를 나눠 여러 GPU에서 연산한 후 각 최종 업데이트 결과에 대해 평균을 내는 방식
- Synchronization: N개의 GPU에서 연산한 업데이트 결과를 모아서 평균을 내고 다시 재분배

## ✓ Model Parallelism (MP)

- 여러 GPU에 모델 Parameter를 나누어 연산
- Tensor Parallelism: Weight Matrix를 나누어 여러 GPU에서 연산 후 결과를 Concatenation
- Pipeline Parallelism: Model의 layer를 여러 GPU에 쪼개서 순차적으로 학습

## ✓ Multi-GPU Training

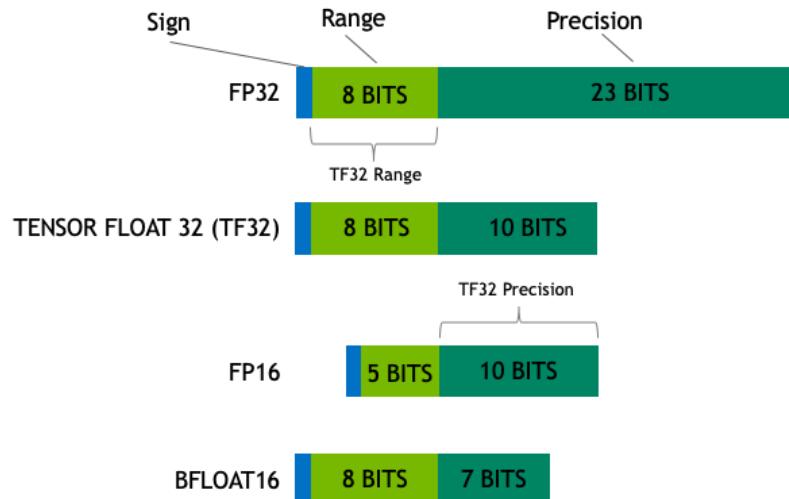
- Accelerator (Huggingface Library)
- DeepSpeed (MS Library for speed and scale)  
: Optimizer State Partitioning + Gradient Partitioning + Parameter Partitioning



# [Appendix] LLM: Data Types

## ✓ Common data types used in Machine Learning

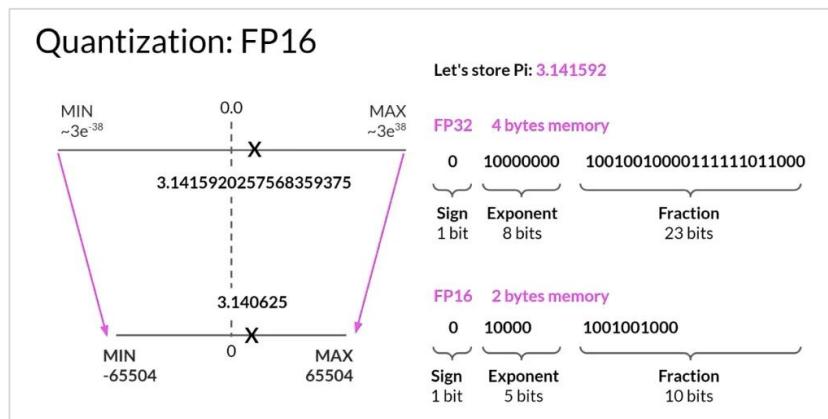
- FLOAT 32, FLOAT 16 or BFLOAT 16



# [Appendix] LLM: Quantization

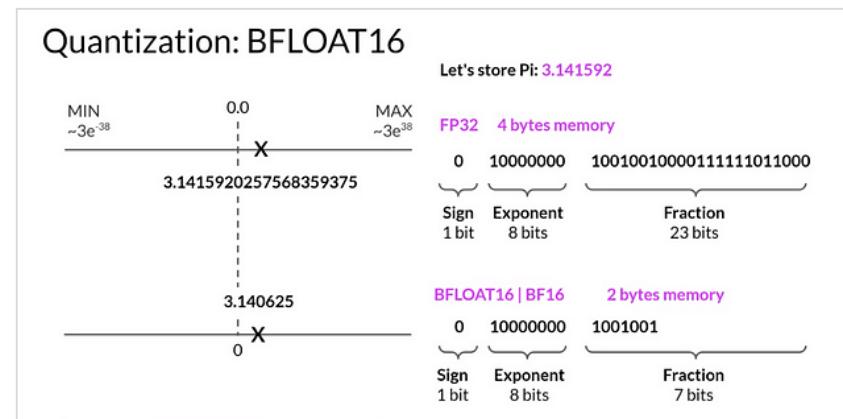
## Quantization: FP16

By employing 16-bit half precision (FP16), the memory requirement can be reduced by 50. Representing the model parameters as 8-bit integers (INT8) can reduce the memory footprint even further representing a total 98.75% reduction compared to full 32-bit precision.



## Quantization: BFLOAT16

Developed by Google Brain, BF16 serves as a hybrid between FP16 and FP32, **capturing the full dynamic range of FP32 with just 16 bits**. This format offers improved training stability and is supported by newer GPUs, such as NVIDIA's A100. While not suitable for integer calculations, BFLOAT16 has become a popular choice in LLMs.



# [Appendix] GPU

## ✓ Tensor Core



2016



2017



2020



2022

# [Appendix] Wav2Vec

## ✓ Wav2Vec Architecture

