

# Quantization for LLM

---

---

## **1. Weight-only Quantization**

1. AWQ

## **2. Weight and Activation Quantization**

1. SmoothQuant
2. QuaRot/SpinQuant

# Setup

---

3

19

- 실습 자료 “Quantization for LLM.ipnyb”을 colab에서 실행해주세요
- Colab 런타임을 **GPU(T4)**로 설정해 주세요
- **Setup** 코드 셀을 실행해 필요한 패키지를 설치해주세요

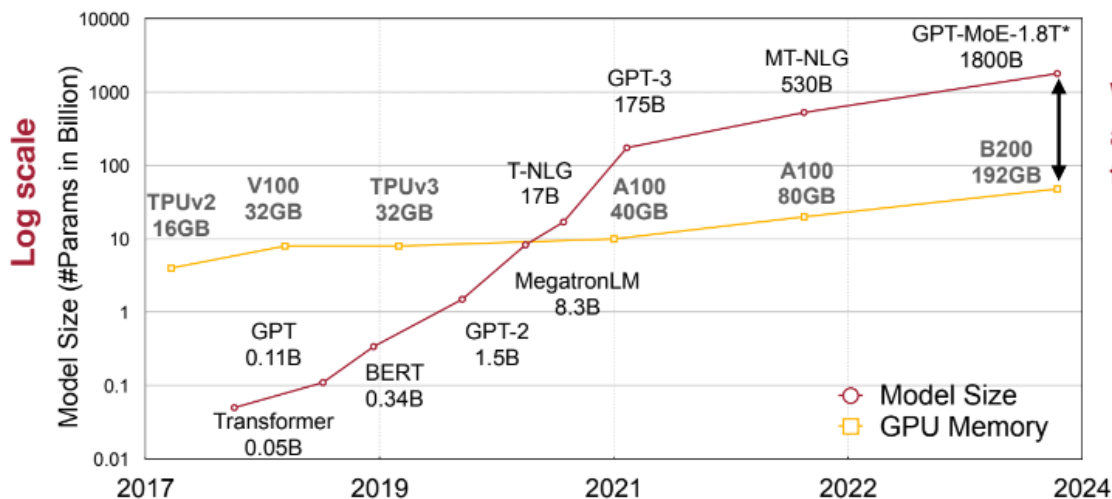
# Challenge for LLM deployment

4

19

## Despite being powerful, LLMs are hard to serve on the edge

- LLM sizes and computation are increasing exponentially.
- Domain-specific accelerator alone is not enough.
- We need model compression techniques and system support to bridge the gap.



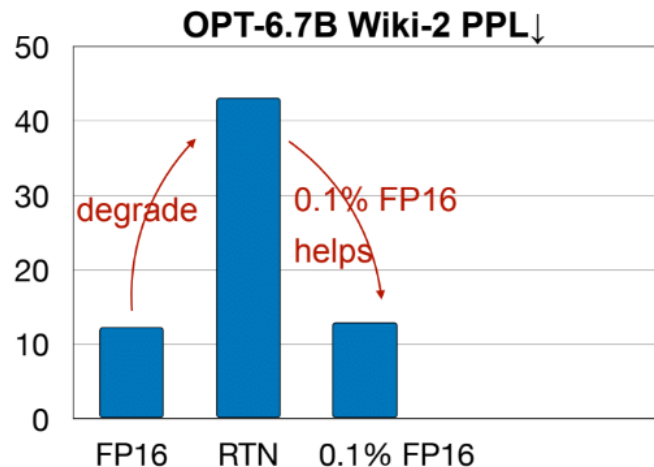
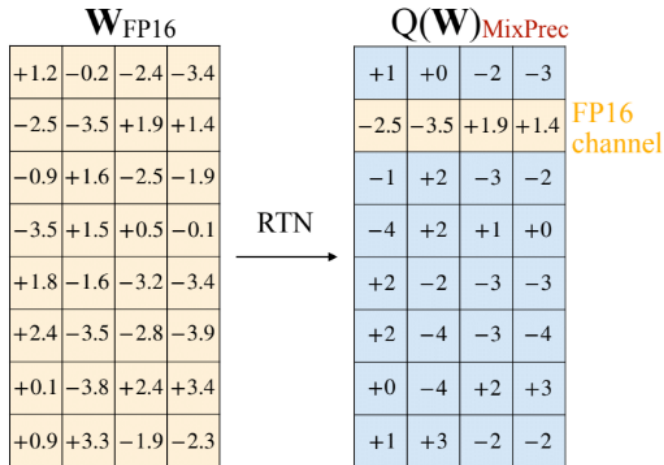
**We need efficient algorithms and systems to bridge the gap.**

# AWQ: **A**ctivation-aware **W**eight **Q**uantization

5

19

Observation: Weights are not equally important; 0.1% salient weights



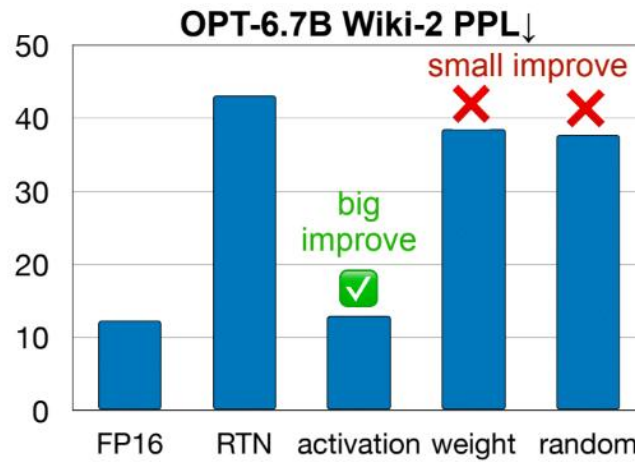
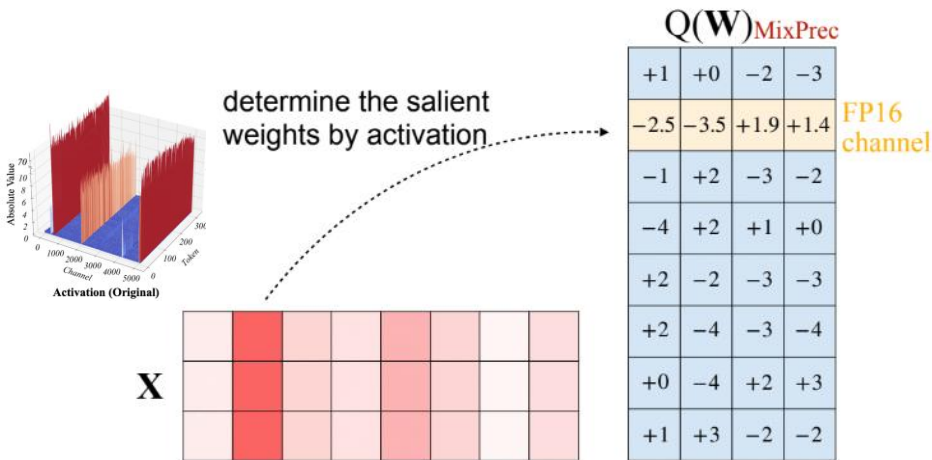
- We find that weights are not equally important, keeping **only 0.1%** of salient weight channels in FP16 can greatly improve perplexity
- But how do we select salient channels? Should we select based on weight magnitude?

# Protect 1% salient channels

6

19

Salient weights are determined by activation distribution, not weight



0.1% FP16 based on

- We find that weights are not equally important, keeping **only 0.1%** of salient weight channels in FP16 can greatly improve perplexity
- But how do we select salient channels? Should we select based on weight magnitude?
- No! We should look for **activation** distribution, but not **weight**!

# Scale 0.1% salient channels

7

19

Protecting salient weights by scaling (no mixed prec.)

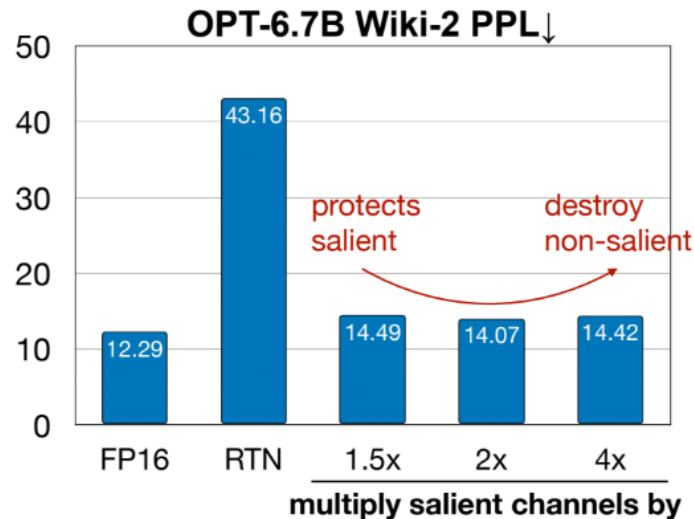
$Q($

W			
+1.2	-0.2	-2.4	-3.4
-2.5	-3.5	+1.9	+1.4
-0.9	+1.6	-2.5	-1.9
-3.5	+1.5	+0.5	-0.1
+1.8	-1.6	-3.2	-3.4
+2.4	-3.5	-2.8	-3.9
+0.1	-3.8	+2.4	+3.4
+0.9	+3.3	-1.9	-2.3

$\times 1$   
 $\times 2$   
 $\times 1$   
 $\times 1$   
 $\times 1$   
 $\times 1$   
 $\times 1$   
 $\times 1$

$\text{fuse to previous op}$

$WX \rightarrow Q(W \cdot s)(s^{-1} \cdot X)$



- Multiplying the salient channels with  $s > 1$  reduces its quantization error
- Why?

# Scale 1% salient channels

8

19

## Protecting salient weights by scaling (no mixed precision)

- Consider a linear layer channel  $\mathbf{y} = \mathbf{w}x$  (from  $\mathbf{W}\mathbf{x}$ ). We care about the quantization error from  $Q(\mathbf{w})x$
- $Q(\mathbf{w}) = \Delta \cdot \text{Round}(\mathbf{w}/\Delta), \quad \Delta = \frac{\max(\mathbf{w})}{2^{N-1}} \quad \text{Error} = \Delta \cdot \text{RoundErr}$
- The scaled version is  $Q(\mathbf{w} \cdot s)(x/s) = \Delta' \cdot \text{Round}(s\mathbf{w}/\Delta') \cdot x \cdot \frac{1}{s} \quad \text{Error}' = \Delta' \cdot \text{RoundErr} \cdot \frac{1}{s}$
- We find that the error from  $\text{Round}()$  is always  $\sim 0.25$  (average from 0-0.5)
- The maximum value in a group “usually” does not change if we just scale up a channel  $\rightarrow \Delta$  not changed
- With  $s > 1$ , the error is scaled down.

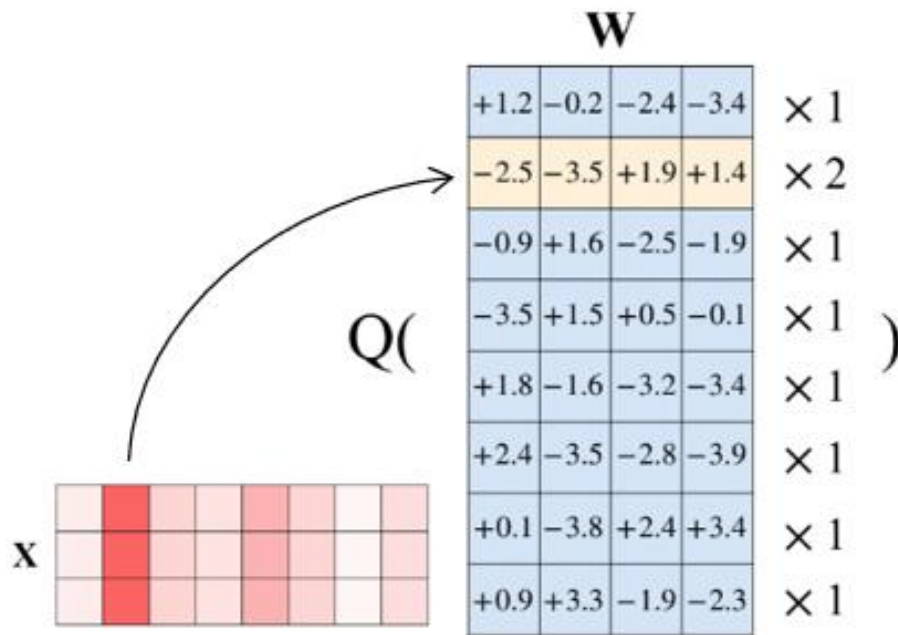
$$q = \text{int}(\text{round}(r/s)) + z$$



# [실습1] Scale 1% salient channels

9

19



$$WX \rightarrow Q(W \cdot s) (s^{-1} \cdot X)$$

```
##### YOUR CODE STARTS HERE #####

# Step 1: importance를 기준으로 1%의 중요한 채널을 찾으세요 (hint: use torch.topk())
# hint : torch.topk() 함수를 사용하세요. torch.topk() 함수는 PyTorch에서 텐서의 값 중 상위 k개의 값과 그들의 인덱스를 반환하는 함수입니다.
outlier_mask = torch.topk(importance, int(len(importance) * 0.01))[1]
assert outlier_mask.dim() == 1

##### YOUR CODE ENDS HERE #####

# 스케일 팩터를 적용하는 것을 시뮬레이션하기 위해, 양자화 전에 스케일 팩터를 곱하고, 양자화 후에 스케일 팩터로 나눕니다.
# scale_factor를 이용해 중요한 가중치 채널의 값을 확대합니다.
m.weight.data[:, outlier_mask] *= scale_factor

m.weight.data = pseudo_quantize_tensor(m.weight.data, n_bit=w_bit, q_group_size=q_group_size)

##### YOUR CODE STARTS HERE #####

# Step 2: scale_factor를 이용해 중요한 가중치 채널의 값을 다시 축소하세요.
m.weight.data[:, outlier_mask] /= scale_factor

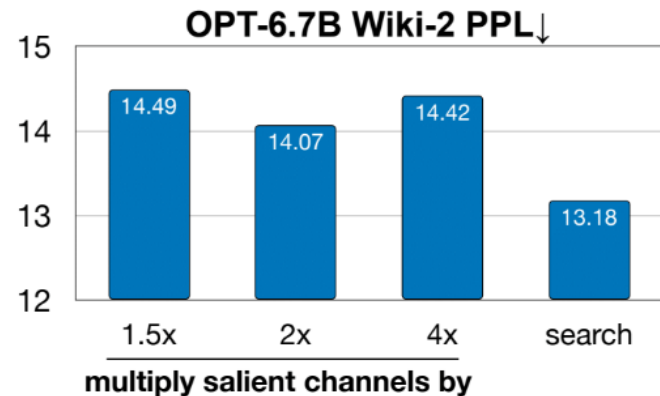
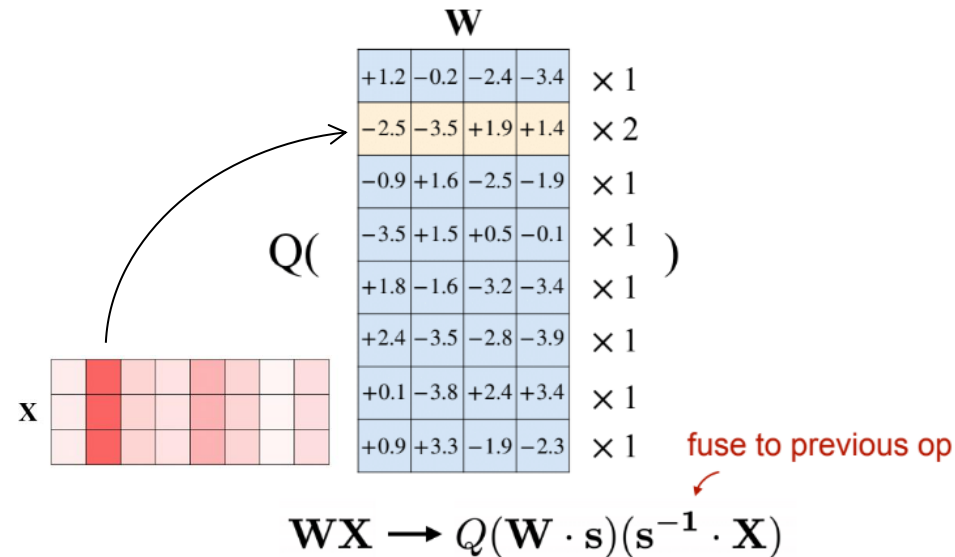
##### YOUR CODE ENDS HERE #####
```

# Scale Factor Search

11

19

Protecting salient weights by scaling (no mixed prec.)



$$\mathcal{L}(s) = \|Q(W \cdot s)(s^{-1} \cdot X) - WX\|$$

$S = S_X^\alpha$  Activation-awareness is important, but not weight-awareness

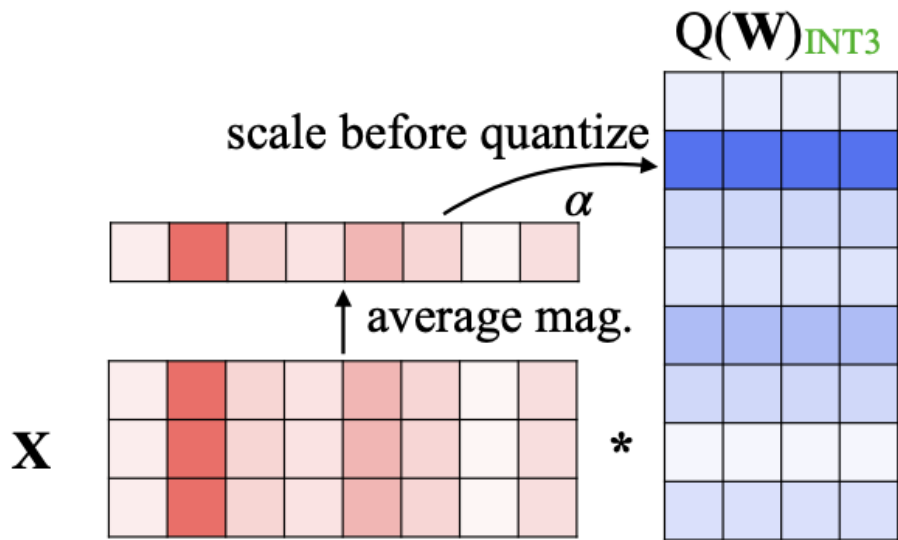
- Multiplying the salient channels with  $s > 1$  reduces its quantization error
- Take a data-driven approach with a fast **grid search**

**L1 norm of Activation**

# [실습2] Scale Factor Search

12

19



$$S_{X^\alpha} = \|X\|_1, \quad \|X\|_1 = \sum |X_i|$$

$S = S_{X^\alpha}$  **Activation-awareness** is important,  
but not weight-awareness

- Multiplying the salient channels with  $s > 1$  reduces its quantization error
- Take a data-driven approach with a fast **grid search**

```
##### YOUR CODE STARTS HERE #####

# Step 2: 공식에 따라 스케일 계산: scales = s_x^ratio
scales = s_x ** ratio # must clip the s_x, otherwise will get nan later

assert scales.shape == s_x.shape

##### YOUR CODE ENDS HERE #####

scales = scales / (scales.max() + scales.min()).sqrt().view(1, -1)

for fc in linears2scale:

    scales = scales.to(fc.weight.device)

    # scale_factor를 이용해 중요한 가중치 채널의 값을 확대합니다.
    fc.weight.mul_(scales)

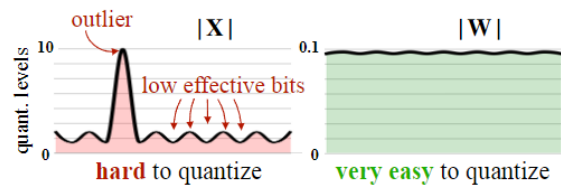
    fc.weight.data = pseudo_quantize_tensor(fc.weight.data, w_bit, q_group_size)

##### YOUR CODE STARTS HERE #####

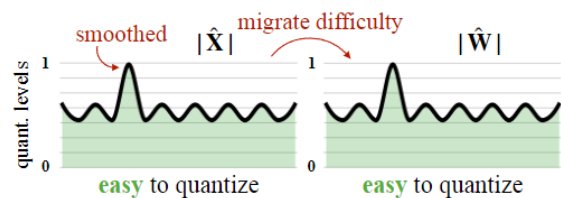
# Step 3: scale_factor를 이용해 중요한 가중치 채널의 값을 다시 축소하세요.
fc.weight.data /= scales

##### YOUR CODE ENDS HERE #####
```

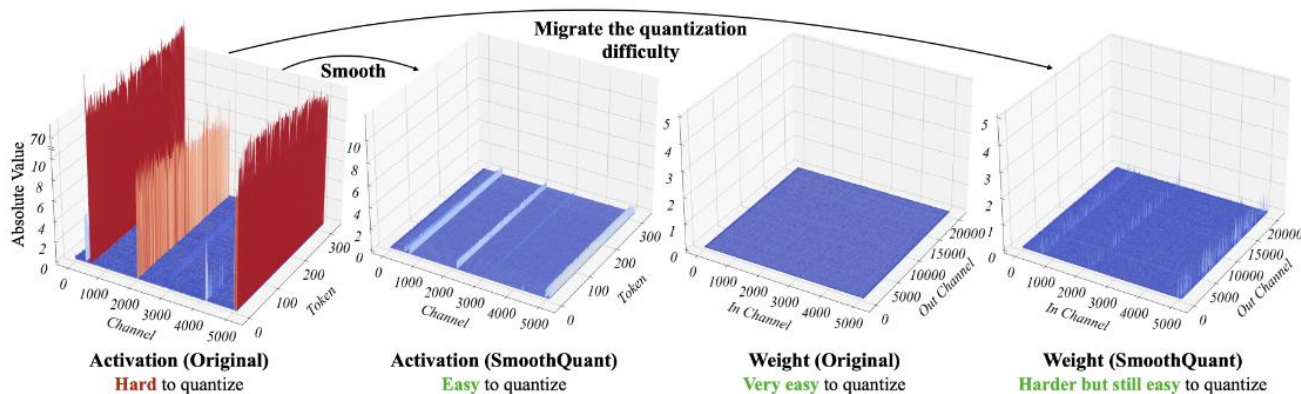
- SmoothQuant's intuition
  - Migrate **scale** from activations to weights  $W$  before quantization



(a) Original



(b) SmoothQuant



- Main idea of SmoothQuant**

$$\mathbf{Y} = \mathbf{X} \cdot \mathbf{W}, \mathbf{Y} \in \mathbb{R}^{T \times C_o}, \mathbf{X} \in \mathbb{R}^{T \times C_i}, \mathbf{W} \in \mathbb{R}^{C_i \times C_o},$$

$$s_j = \max(|\mathbf{X}_j|)^\alpha / \max(|\mathbf{W}_j|)^{1-\alpha}$$

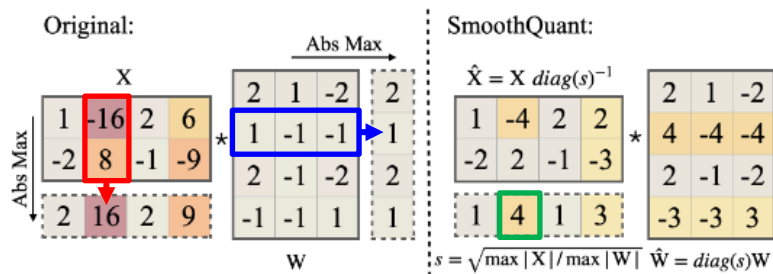


Figure 5: Main idea of SmoothQuant when  $\alpha$  is 0.5. The smoothing factor  $s$  is obtained on calibration samples and the entire transformation is performed offline. At runtime, the activations are smooth without scaling.

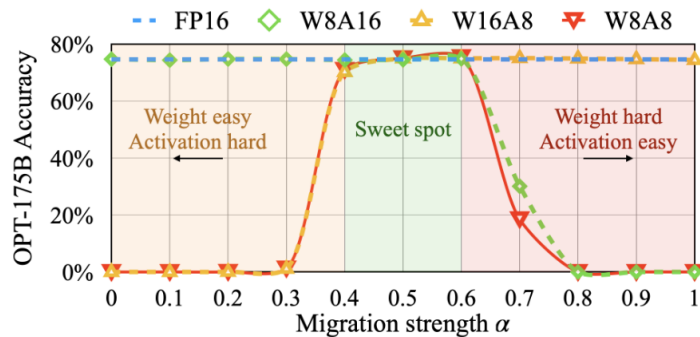
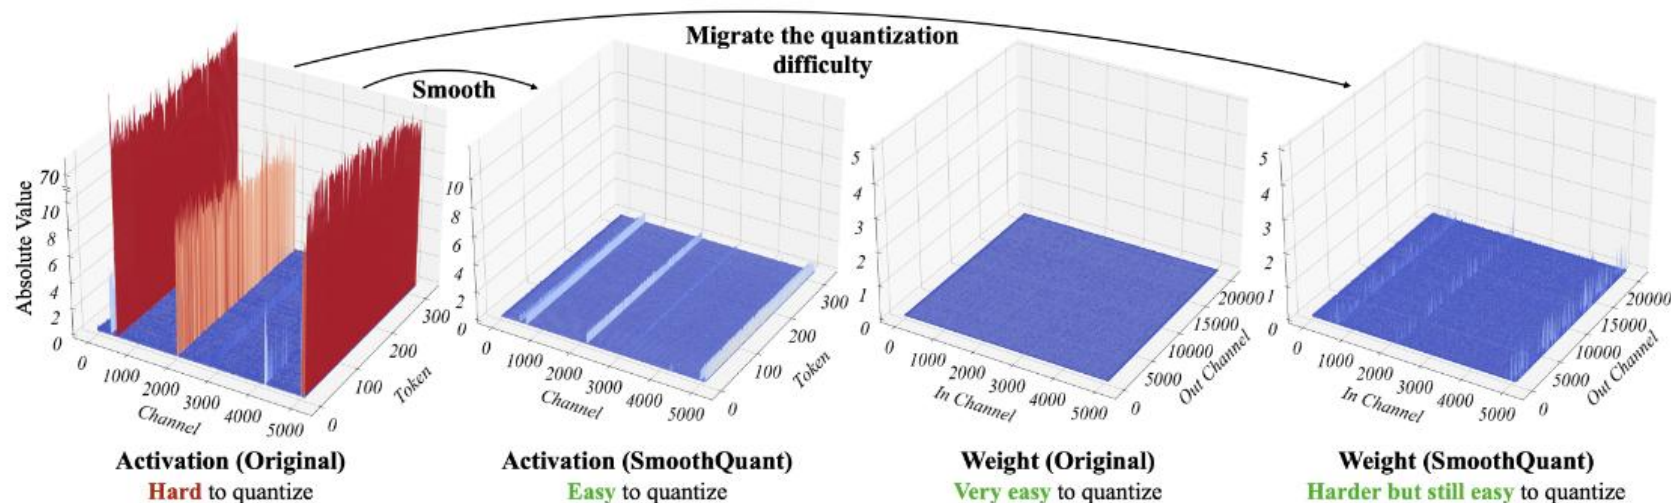


Figure 10: A suitable migration strength  $\alpha$  (sweet spot) makes both activations and weights easy to quantize. If the  $\alpha$  is too large, weights will be hard to quantize; if too small, activations will be hard to quantize.

# [실습3] Quantization Difficulty Migration

16

19





```
def smooth_ln_fcs_by_scale(ln, fcs, scale):
    if not isinstance(fcs, list):
        fcs = [fcs]
    assert isinstance(ln, nn.LayerNorm)
    for fc in fcs:
        assert isinstance(fc, nn.Linear)
        ##### YOUR CODE STARTS HERE #####
        # Step 1: layernorm의 weight와 bias를 scale로 나누어주세요. (hint: div_()함수를 통해 tensor 전체를 특정한 값으로 나누어 줄 수 있습니다.)
        ln.weight.div_(scale)
        ln.bias.div_(scale)
        ##### YOUR CODE ENDS HERE #####

    for fc in fcs:
        ##### YOUR CODE STARTS HERE #####
        # Step 2: fc의 weight에 scale을 곱해주세요. (hint: mul_()함수를 통해 tensor 전체에 특정한 값을 곱해 줄 수 있습니다.)
        fc.weight.mul_(scale)
        ##### YOUR CODE ENDS HERE #####
```

# [실습4] Scale Factor Search

18

19

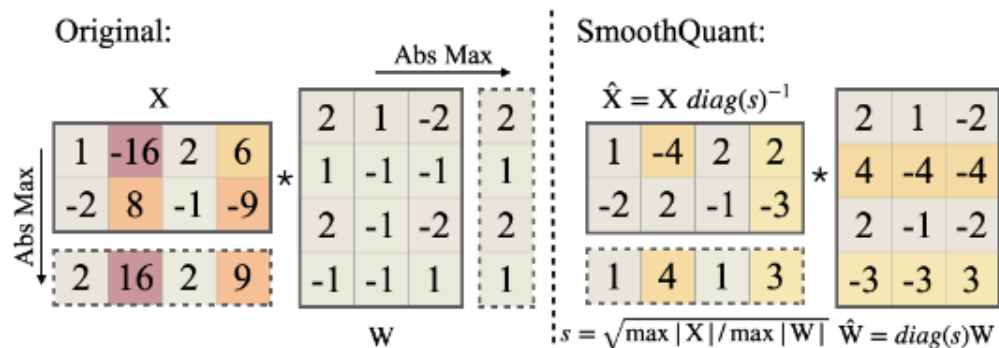


Figure 5: Main idea of SmoothQuant when  $\alpha$  is 0.5. The smoothing factor  $s$  is obtained on calibration samples and the entire transformation is performed offline. At runtime, the activations are smooth without scaling.

```
scales = (  
    ##### YOUR CODE STARTS HERE #####  
    #Activation Scales 값과 Weight Scales 값에 alpha를 적절히 거듭제곱해주어야 합니다.  
    #Hint: pow()함수를 통해서 거듭제곱을 사용할 수 있습니다.  
    (act_scales.pow(alpha) / weight_scales.pow(1 - alpha))  
    ##### YOUR CODE ENDS HERE #####  
)
```

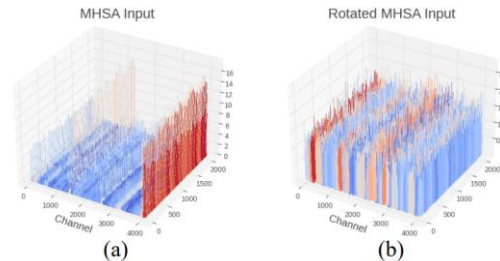
# Rotation based Quantization

20

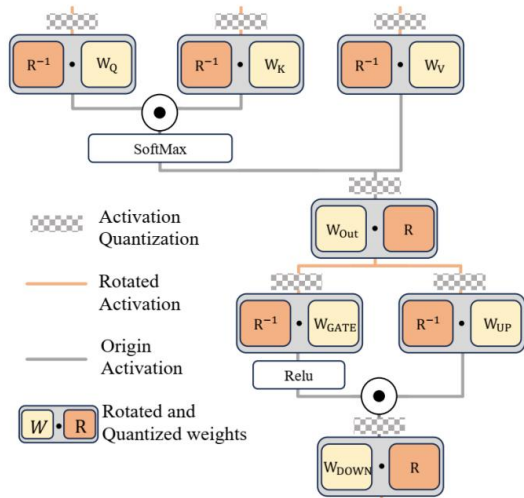
19

$$Y = (XR)(R^{-1}W^{-1}) = XW^T$$

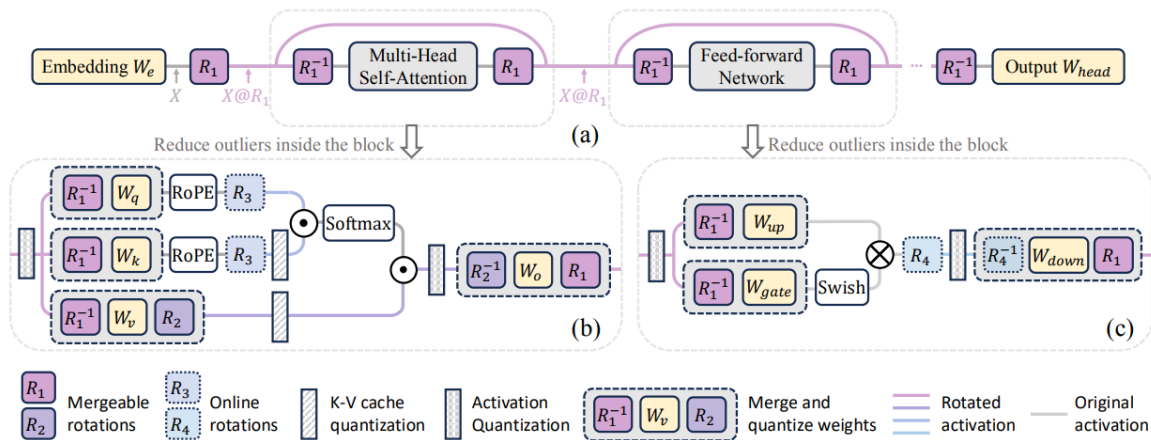
A rotation matrix is an orthogonal matrix  $R$  satisfied  $RR^T = 1$  and  $|R| = 1$



## QuaRot



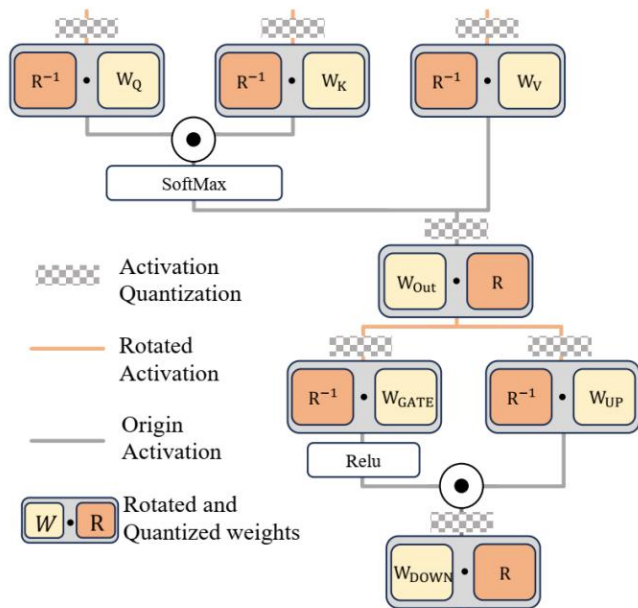
## SpinQuant



# [실습5] Rotate Matrix 적용

21

19



- @  $\rightarrow$  Dot Product
- nn.Linear 연산은  $W^T$  형태로 저장
- Embedding Parameter Shape : (Num\_Tokens, Hidden\_dim)
- Linear Parameter Shape : (Output\_Channel, Input\_Channel)
- Roation Matrix Shape : (Hidden\_dim, Hidden\_dim)

```
##### YOUR CODE STARTS HERE #####
# Pytorch에서 @ 연산이 Dot Product 임을 사용하기 바랍니다.
# nn.Linear 연산의 Parameter는 W^T 형태로 저장되어 있다는 것을 유의하시기 바랍니다.
# Embedding Parameter Shape : (Num_Tokens, Hidden_dim)
# Linear Parameter Shape : (Output_Channel, Input_Channel)
# Roation Matrix Shape : (Hidden_dim, Hidden_dim)

if isinstance(m, nn.Embedding):
    W_ = m.weight.data
    m.weight.data = W_ @ R1

if isinstance(m, nn.Linear):
    if "out_proj" in n or "fc2" in n:
        # Att Out Proj, FFN Down Proj
        W_ = m.weight.data
        m.weight.data = R1.T @ W_ ➡  $x \cdot W \cdot R$ 
    else:
        # QKV Proj, FFN Up Proj, FFN Gate Proj
        W_ = m.weight.data
        m.weight.data = W_ @ R1 ➡  $x \cdot R^T \cdot W$ 

##### YOUR CODE ENDS HERE #####
```