

# Time Series Forecasting: Advanced

**Jaemin Yoo**

School of Electrical Engineering  
Kim Jaechul Graduate School of AI



# Outline

1. **State space models**
2. From linear to deep models
3. Convolutional neural networks
4. Encoder-decoder structure
5. Summary

# Error-Correction Function

- **Q:** What are alternatives to linear regression models?
- **Idea:** Create a prediction by improving the previous prediction:

$$\hat{z}_{t+1} = \underbrace{\hat{z}_t}_{\text{previous forecast}} + \alpha \underbrace{(z_t - \hat{z}_t)}_{\text{error in previous forecast}},$$

- Also known as an **error-correction function**.
- Since it corrects the error caused from the previous forecast.
- $\alpha$  is called a smoothing parameter.

# Simple Exponential Smoothing

- We can rewrite the error-correction function as follows:


$$\begin{aligned}\hat{z}_3 &= \hat{z}_2 + \alpha(z_2 - \hat{z}_2) \\ &= \alpha z_2 + (1 - \alpha)\hat{z}_2 \\ &= \alpha z_2 + \alpha(1 - \alpha)z_1 + (1 - \alpha)^2\hat{z}_1 \\ &\dots\end{aligned}$$

- The model is called **ETS** (Simple **ExponenTial S**moothing):

$$\hat{z}_{T+h} = \alpha z_T + \alpha(1 - \alpha)z_{T-1} + \alpha(1 - \alpha)^2 z_{T-2} + \dots + (1 - \alpha)^T \hat{z}_1$$

# Simple Exponential Smoothing

- Larger  $\alpha$  puts more emphasis on the recent observations.



# State and Prediction


- Let's separate ETS into two parts: **state** and **prediction**.
  - $l_t$  is a state at time  $t$ .
  - $l_t$  is used to create the prediction  $\hat{z}_t$  at time  $t$ .
  - $l_t$  is updated to  $l_{t+1}$  for the next time step.

Forecast equation

Error Correction

$$\hat{z}_t = l_{t-1}$$

$$l_t = l_{t-1} + \alpha(z_t - \hat{z}_t),$$



# General Exponential Smoothing

- **Idea:** Let's make the state  $\mathbf{l}_t$  contain more information.

Forecast equation

$$\hat{z}_t = \mathbf{a}_t^T \mathbf{l}_{t-1}$$

Error Correction

$$\mathbf{l}_t = \mathbf{F}_t \mathbf{l}_{t-1} + \mathbf{g}_t (z_t - \hat{z}_t),$$

- $\mathbf{l}_t$  is now a vector, and is not equivalent to the prediction  $\hat{z}_t$ .
- Parameter  $\mathbf{a}_t$  maps  $\mathbf{l}_t$  to  $\hat{z}_t$  through the dot product.
- Parameters  $\mathbf{F}_t$  and  $\mathbf{g}_t$  are used to update  $\mathbf{l}_t$  to  $\mathbf{l}_{t+1}$ .

# State Space Models

- **Q:** Do we really need the error-correction part  $z_t - \hat{z}_t$ ?
  - Maybe not. Let's model the error  $z_t - \hat{z}_t$  as a random variable  $\epsilon_t$ .
- **State space model (SSM)** simplifies the previous model.
  - Consist of the measurements and state transition parts.
  - Add white noise to both parts, replacing the error-correction function.

**Measurements**  $z_t = \mathbf{a}_t^T \mathbf{l}_{t-1} + \epsilon_t, \quad \epsilon_t \sim N(0, \sigma^2)$

**State transition**  $\mathbf{l}_t = \mathbf{F}_t \mathbf{l}_{t-1} + \mathbf{g}_t \epsilon_t, \quad \mathbf{l}_0 \sim N(\boldsymbol{\mu}_0, \text{diag}(\sigma_0^2)).$



# Linear State Space Models


- Let's combine the SSM with (feature-based) linear regression:

$$z_t \sim P(z_t|y_t)$$

$$y_t = \mathbf{a}_t^T \mathbf{l}_{t-1} + \mathbf{w}^T \mathbf{x}_t$$

$$\mathbf{l}_t = \mathbf{F}_t \mathbf{l}_{t-1} + \mathbf{g}_t \epsilon_t$$

- Pros:** Show the strength of both models.
- Cons:** More parameters to learn.



# Linear State Space Models


Linear State Space Model part:

$$u_t = \mathbf{a}_t^T \mathbf{l}_{t-1}$$

Feature-based part:

$$b_t = \mathbf{w}^T \mathbf{x}_t$$

Probabilistic model for data (likelihood):  $z_t \sim P(z_t | u_t + b_t)$




# Outline

1. State space models
2. **From linear to deep models**
3. Convolutional neural networks
4. Encoder-decoder structure
5. Summary


# Generalizing Linear Regression

- Recall that **linear regression** works as  $z_t = w^T x_t$  for forecasting.



# Generalizing Linear Regression


- We can generalize the linear mapping using a **deep neural network**.



$$z_t = \sigma(\mathbf{w}_l^T (\sigma(W_{l-1}^T (\sigma(W_{l-2}^T (\cdots W_0^T \mathbf{x}_t)))))) := \text{DEEP-NET}(\mathbf{x}_t)$$


# Multi-layer Perceptrons

- **Multi-layer perceptrons (MLPs):**
  - The most basic deep learning architecture.
  - Each neuron in a hidden layer computes an affine function of the previous layer.
  - It is then followed by an *activation function*:
$$h_{l,j} = \sigma(\mathbf{w}_{l,j}^T \mathbf{h}_{l-1} + b_{l,j}).$$
- MLPs are flexible function estimators.
  - More layers  $\rightarrow$  more complex functions.




# Multi-layer Perceptrons

- **Advantages:** MLPs can learn complex input-output relationships.
  - Less manual feature engineering.
- **Disadvantages:** More data are needed.
  - Careful tuning (e.g., regularization, learning rate, etc.) is necessary.
  - The model is sensitive to scaling of inputs.



# Recap: MLPs for Forecasting


- **Question:** How can we model the sequential relationship?





# Recap: State Space Models


- **Question:** Can we do the same with neural networks?



$$l_t = l_{t-1} + \alpha \cdot \epsilon_t$$
$$z_t = v^T l_t + w^T x_t + \epsilon_t$$

# From Feed-forward to Recurrent Models



- We add the concept of **state** to the deep forecasting model.
  - The previous predictions affect the current one.




# Toward Recurrent Neural Networks

- **Recurrent neural networks:**

- Current state  $h_t$  combines
  - Previous state  $h_{t-1}$ .
  - Input features  $x_t$ .
  - Activation function  $\sigma$ .




$$l_t = l_{t-1} + \alpha \cdot \epsilon_t$$
$$z_t = \mathbf{v}^T \mathbf{l}_t + \mathbf{w}^T \mathbf{x}_t + \epsilon_t$$



$$h_t = \sigma(\theta_0 h_{t-1} + \theta_1 x_t)$$
$$z_t = \sigma(\theta h_t)$$

# Long Short-Term Memory (LSTM)

- **LSTM** uses two states  $C_t$  and  $h_t$  with a **forget gate**:
  - The forget gate is similar to the exponential smoothing from ETS.



[HTTP://COLAH.GITHUB.IO/POSTS/2015-08-UNDERSTANDING-LSTMS/](http://colah.github.io/posts/2015-08-understanding-lstms/)


$$C_t = \alpha_t \cdot C_{t-1} + \beta_t \times \sigma(\theta_0 h_{t-1} + \theta_1 x_t)$$

# Outline

1. State space models
2. From linear to deep models
3. **Convolutional neural networks**
4. Encoder-decoder structure
5. Summary

# Convolutional Neural Networks

- **Convolutional neural networks (CNNs):**
  - Neural networks that use convolutional layers.
  - CNNs with 2D convolutions are successful in CV.
    - The idea is to encode *spatial invariance*.
  - 1D convolutions are a promising alternative to RNNs for sequential data.
    - Encode temporal invariance, e.g., stationarity.
    - Often more lightweight and effective than RNNs.

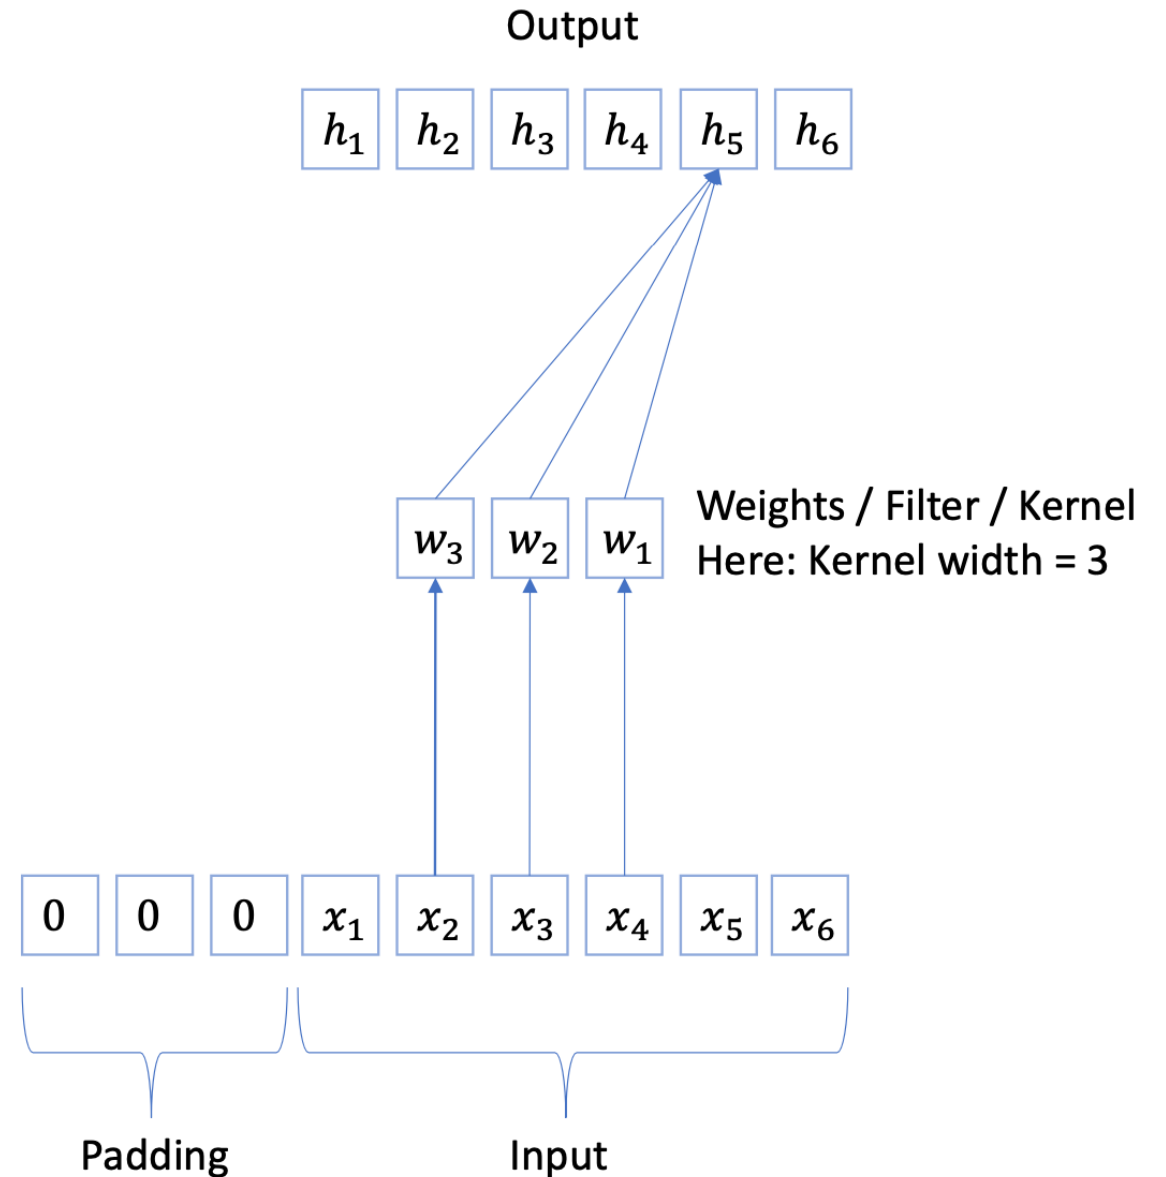


# Convolutional Layers

- The output  $h_j$  in a convolution layer is a discrete convolution of the inputs  $\mathbf{x}$  with weights  $\mathbf{w}$ .
- For a 1-dimensional convolution with a kernel with width  $D$ :

$$h_j = \sum_{d=1}^D w_d x_{j-d}.$$

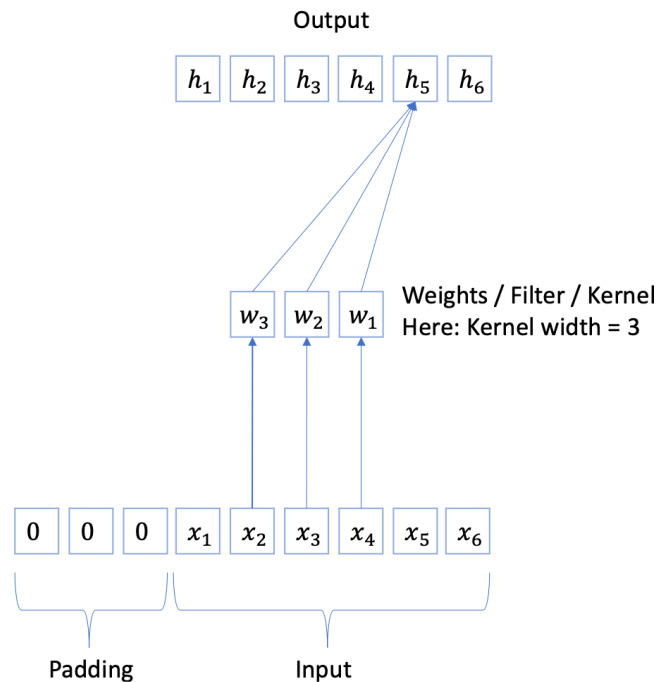
- *Padding* is usually added to the first part of the sequence.




# Causal vs. Non-causal Convolution

- Non-causal convolution is used mostly for timeseries-level tasks.

## Causal Convolution

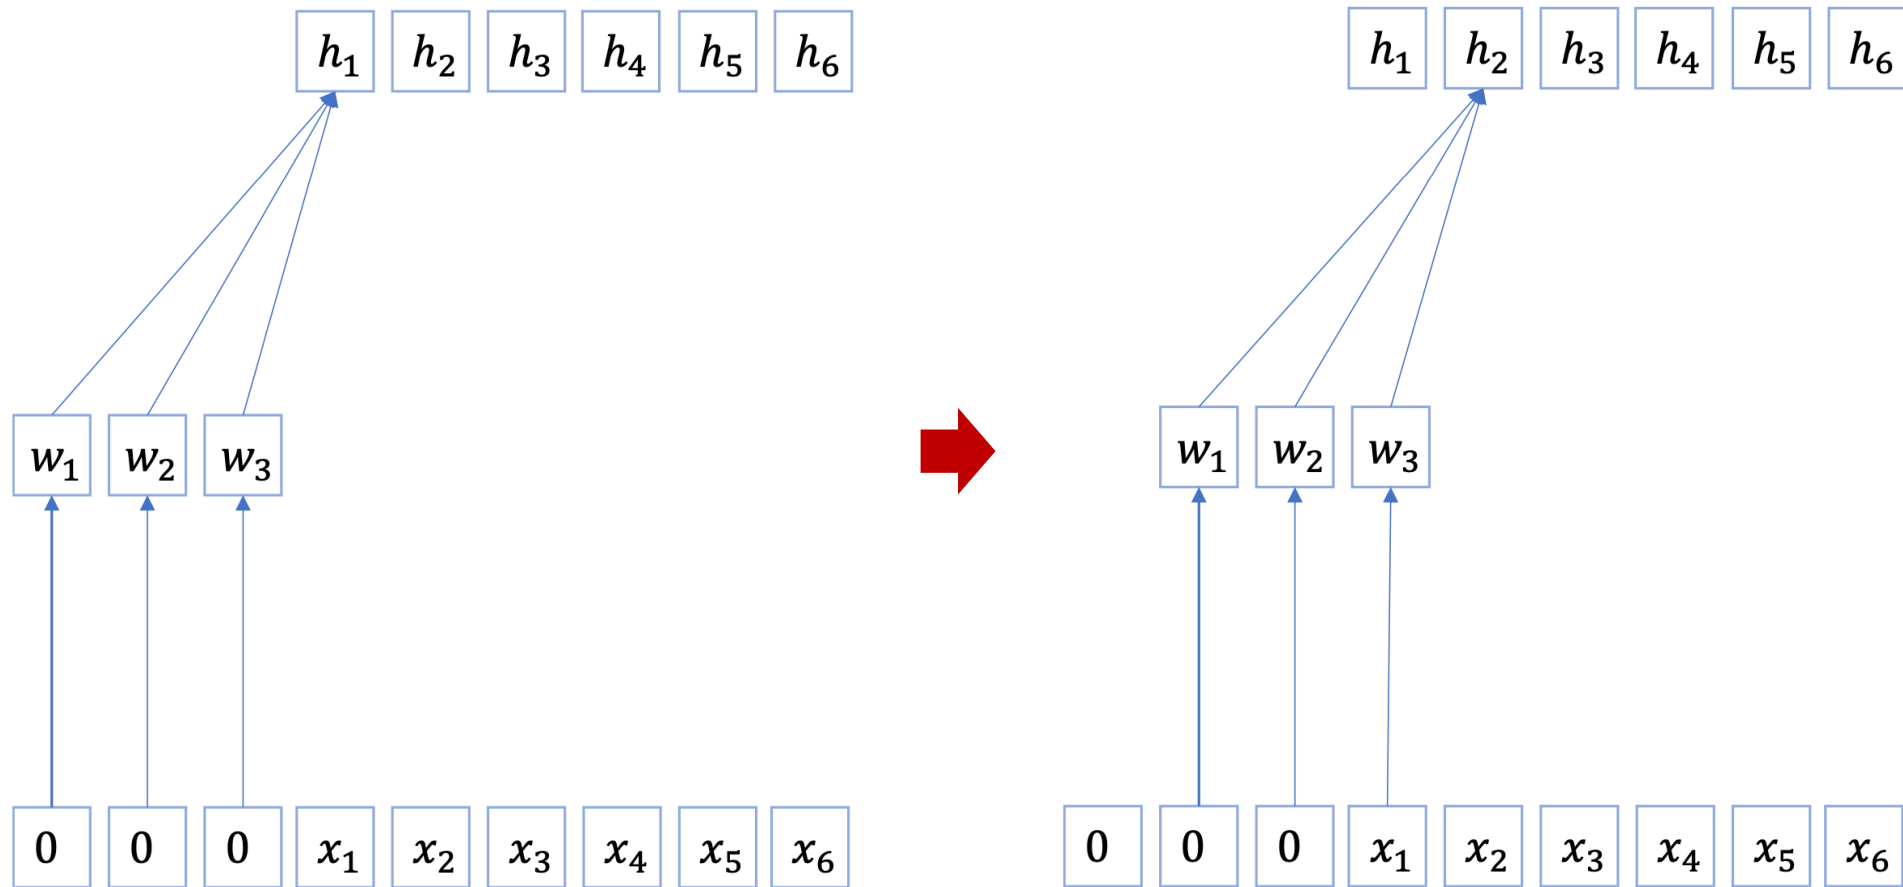


## Non-Causal Convolution






# 1D Causal Convolution



# Canonical Models: Dilated Convolution

- **Dilation** quickly increases receptive field through multiple layers.
  - Forecast is generated in an autoregressive fashion.



# Outline

1. State space models
2. From linear to deep models
3. Convolutional neural networks
4. **Encoder-decoder structure**
5. Summary

# Encoder-Decoder Structure


- Many deep forecasting models rely on additional features  $X$ .
- **Q:** What if  $X$  is given only for the current  $T$ , not for the future?
  - Autoregressive prediction is no longer possible.
- **Solution:** Generalize the model into **encoder-decoder** structure.

# Encoder-Decoder Structure


- **Idea:** Our prediction model consists of an **encoder** and a **decoder**.
  - **Encoder** takes  $\mathbf{z}_{1:T}$  and  $\mathbf{X}_{1:T}$  and summarizes them as a state  $\mathbf{h}_T$ .
  - **Decoder** takes  $\mathbf{h}_T$  and generates predictions without more input.
- We choose a linear mapping  $\hat{y} = \mathbf{W}\mathbf{h}_T$  as a decoder in many cases.
  - Meaning that we map a representation  $\mathbf{h}_T$  into a scalar prediction.

# Encoder-Decoder Structure

- What if we use a sequential model (e.g., an RNN) as a decoder?
- We can generate a sequence, from  $\mathbf{h}_T$ , without further input!



Canonical (One-to-One)



Seq2Seq (Many-to-Many)

# Encoder Part

- Encoder is a general representation learner of time series.
  - That is, the encoded output  $\mathbf{h}_T$  can be used for other tasks as well.
- Encoder is the same in one-to-one and many-to-many cases.
  - **One-to-one:** The output  $\mathbf{h}_T = f_{\text{encoder}}(\cdot)$  is used to predict  $z_{T+1}$ .
  - **Seq2seq:** The output  $\mathbf{h}_T = f_{\text{encoder}}(\cdot)$  is used to predict  $z_{T+1} : z_{T+h}$ .

$$f_{\text{encoder}} : \{z_1, \dots, z_{T_e}\} \mapsto \mathbf{h}_{T_e}$$

$$f_{\text{decoder}} : \mathbf{h}_{T_e} \mapsto \{z_{T_e+1}, \dots, z_{T_e+T_d}\}$$

# Decoder Part

- Decoder should be able to work in an autoregressive way.
  - That is, it should create a sequence without any input.
- Possible decoder models:
  - MLP with  $h$  output neurons; it creates  $h$  outputs at the same time.
  - RNN that assumes dummy (= meaningless) inputs.


$$f_{\text{encoder}} : \{z_1, \dots, z_{T_e}\} \mapsto \mathbf{h}_{T_e}$$

$$f_{\text{decoder}} : \mathbf{h}_{T_e} \mapsto \{z_{T_e+1}, \dots, z_{T_e+T_d}\}$$




# Example: RNN-RNN

- We can use an RNN as both an encoder and a decoder.
  - Decoder RNN uses a dummy input  $u$ .




# Example: RNN-MLP

- We can combine different model structures as well.
  - Decoder MLP generates a sequence altogether.




# Attention RNN

- Suppose we use an RNN encoder in the Seq2Seq structure.
- **Natural approach:** We pass the last state  $h_T$  to the decoder.
  - Why? The last cell creates a good summary of all observations.



# Attention RNN

- **Limitation 1:** The encoder is likely to forget early observations.
  - Especially when the window size is large.
- **Limitation 2:** Single state  $\mathbf{h}_T$  may not be enough for the decoder.
  - Decoder needs different information at different locations.



# Attention RNN

- **Solution:** Attention mechanism.


- Let  $\mathbf{h}_1, \dots, \mathbf{h}_T$  be the state vectors created from RNN cells.
- Let  $\mathbf{h}_{T+k}$  be the state vector for the  $k$ -th decoding step.
- Create a weighting function  $f$  such that

$$f(\mathbf{h}_{T+k}, \mathbf{h}_1, \dots, \mathbf{h}_T) = \frac{\exp(\mathbf{h}_{T+k}^\top \mathbf{h}_i)}{\sum_{i=1}^T \exp(\mathbf{h}_{T+k}^\top \mathbf{h}_i)}.$$

- $\mathbf{h}_{T+k}$  is the query of attention; it determines which  $i$  is more important.
- $\mathbf{h}_i$  is a key of attention; it contains the property of time step  $i$ .

# Properties of Attention

- **Property 1:** All scores are positive (thanks to  $\exp(\cdot)$ ).
- **Property 2:** The sum of scores is always one.
- **Property 3:** Higher  $\mathbf{h}_{T+k}^\top \mathbf{h}_i$  leads to a higher score.
  - It means more related.



# Attention RNN

- How an attention RNN works:
  - Compute the scores of all encoder states using the scoring function.
  - Compute the weighted average of these states:

$$\vec{\mathbf{h}}_{T+k} = \sum_{i=1}^T \alpha_i \mathbf{h}_i.$$

- Pass  $\vec{\mathbf{h}}_{T+k}$  as an input to the  $k$ -th decoding step.
- **Note:** The state changes at every decoding step by the query  $\mathbf{h}_{T+k}$ .

# Outline

1. State space models
2. From linear to deep models
3. Convolutional neural networks
4. Encoder-decoder structure
5. **Summary**



# Summary

- MLPs are robust baseline methods.
- RNNs were the de facto standard model for sequence modeling.
- Later research shows that CNNs are more effective than RNNs.
  - Although some people favor advanced versions of RNNs.
- Transformer methods show SOTA performance in some cases.
  - They still have various limitations, e.g., quadratic complexity.
  - Currently an active area of research in time series analysis.