

Time Series Practice

Data AI Lab

School of Electrical Engineering

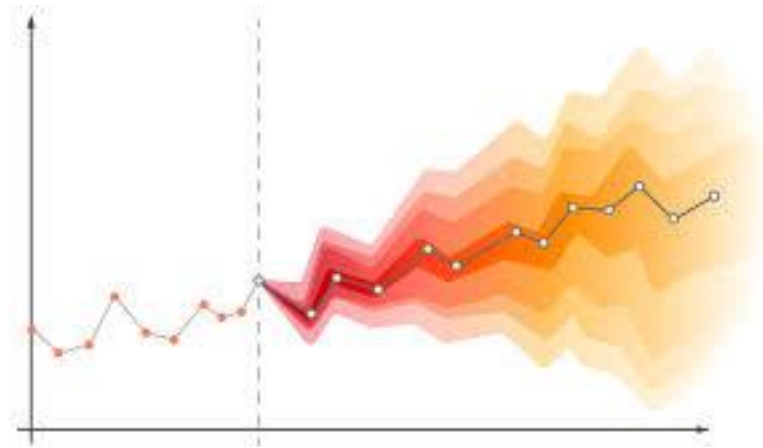


Outline

1. **Introduction**
2. Data Processing
3. Background of Practice Model
4. Practice
5. (Optional) Encoder-decoder structure

Time Series Forecasting

- Predict future values based on historical data.
 - e.g., weather forecasting, traffic prediction, sales forecasting ...



Time Series and Deep Learning

- Learns complex, non-linear patterns.
- Popular Models
 - **RNN/LSTM**: Sequential dependencies over time.
 - **CNN**: Efficiently captures local patterns and stationarity.
 - **Transformer**: Handles global dependencies with attention mechanisms.

Time Series Practice

- Build a model for time series forecasting.
- Steps:
 - Data Collection
 - Data Preprocessing
 - Model Training
 - Evaluation

Outline

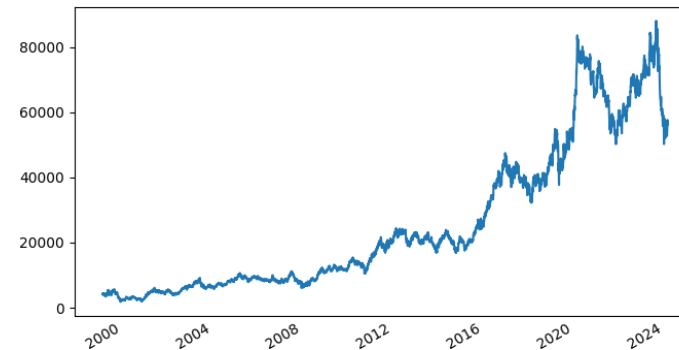
1. Introduction
- 2. Data Processing**
3. Background of Practice Model
4. Practice
5. (Optional) Encoder-decoder structure

Introduction

- **Objective:**
 - Forecast financial time series data using deep learning models.
- **Data Source:** Yahoo Finance (yFinance API)
 - It offers a Pythonic way to fetch financial data from Yahoo!® finance.
 - Documentation website: <https://ranaroussi.github.io/yfinance/index.html>
 - **[Caveat] It is intended for research and educational purposes.**



yfinance is licensed under the Apache License, Version 2.0



Experimental Setup and Environment

- **Install** required packages and configure GPU settings

```
# Math and data preprocessing libraries
import math
import pandas as pd
import numpy as np
from sklearn.preprocessing import MinMaxScaler

# For handling dataset
import yfinance as yf
from datetime import date

# For visualization
import seaborn as sns
import matplotlib.pyplot as plt
```

```
# For deep learning
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, TensorDataset

# For evaluation
from sklearn.metrics \
import root_mean_squared_error, mean_absolute_percentage_error

device = torch.device('cuda' if \
                      torch.cuda.is_available() else 'cpu')

print(device)
```


Dataset Preparation

- **Download** the market data
 - Try other stock tickers: 'AAPL', 'NVDA', '005930.KS'

```
start_date = '2020-01-01'
end_date = '2024-12-31'

df = yf.download('GOOG', start=start_date, end=end_date)

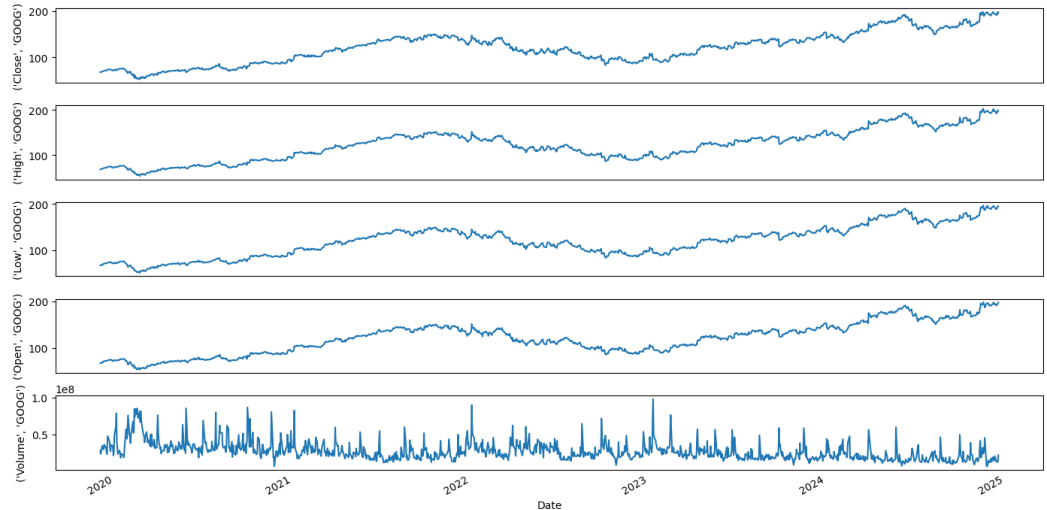
# Inspect the data
print()
print(df.head())
print(df.info())
```

Price Ticker Date	Close GOOG	High GOOG	Low GOOG	Open GOOG	Volume GOOG
2020-01-02	68.123726	68.162086	66.837348	66.837348	28132000
2020-01-03	67.789429	68.379312	67.036336	67.151721	23728000
2020-01-06	69.460922	69.575007	67.258334	67.258334	34646000
2020-01-07	69.417572	69.898343	69.270099	69.646752	30054000
2020-01-08	69.964615	70.326314	69.293024	69.354799	30560000

Data Visualization

- **Draw line plots** for each feature:
 - There are five features used (Open, High, Low, Close, Volume)

```
ncols = 1
nrows = int(round(df.shape[1] / ncols, 0))
fig, ax = plt.subplots(nrows=nrows, ncols=ncols, \
                        sharex=True, figsize=(14, 7))
for i, ax in enumerate(fig.axes):
    sns.lineplot(data=df.iloc[:, i], ax=ax)
    ax.tick_params(axis="x", rotation=30, \
                  labelsize=10, length=0)
fig.tight_layout()
plt.show()
```

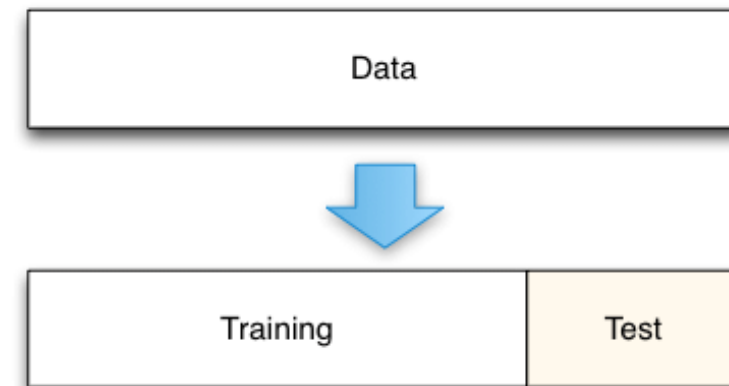


Data Preprocessing: Train-Test Split

- Why do we split the data?
 - Ensures that the model is evaluated on unseen data.
 - Prevents overfitting by testing on separate data which is not used during training.
- **Split** into training and test data:

```
# Train test split
train_ratio = 0.8
training_data_len = math.ceil(len(df) * train_ratio)

# Splitting the dataset
train_data = df[:training_data_len][['Open']]
test_data = df[training_data_len:][['Open']]
print(train_data.shape)
print(test_data.shape)
```



Data Preprocessing: Scaling

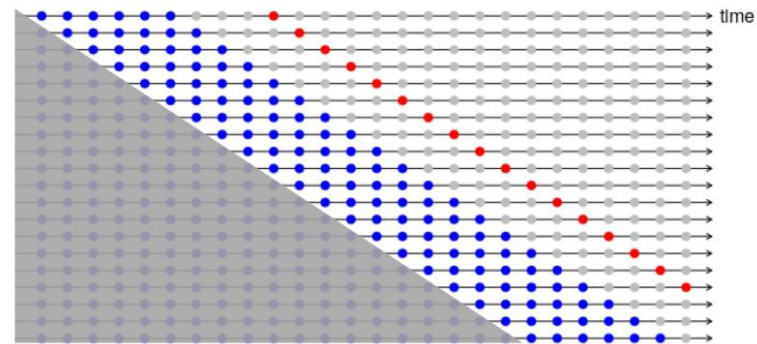
- Why do we scale the data?
 - Prevents features with larger magnitudes from dominating the training process.
- **Scale** the data to normalize values between 0 and 1:

```
scaler = MinMaxScaler(feature_range=(0, 1))  
train_scaled = scaler.fit_transform(train_data.values)  
test_scaled = scaler.transform(test_data.values)
```

Data Preprocessing: Sliding Window

- Why do we convert (time-series) data into sequences?
 - Time series models require sequences to capture temporal dependencies.
 - The model can learn specific patterns for predicting future values.
 - Preprocess the data to enable time-dependent sequential input.
 - Sequence length defines the look-back window for predicting future values.
- **Create labeled training pairs** by sliding window:

```
def convert_data_into_tensors(data_seq):  
    features, labels = [], []  
    for i in range(len(data_seq) - sequence_length):  
        features.append(data_seq[i:i + sequence_length])  
        labels.append(data_seq[i + sequence_length, 0])  
    features, labels = np.array(features), np.array(labels)  
  
    features = torch.tensor(features, dtype=torch.float32)  
    labels = torch.tensor(labels, dtype=torch.float32)  
    return features, labels
```



Data Preprocessing: Batching

- Why do we make batch data?
 - If using the entire training dataset to update model parameters, the training process can become **slow**, and the entire dataset may **not fit into memory**.
- **Create data loaders** for efficient batch processing:

```
batch_size = 32

def to_loader(x, y, batch_size, shuffle):
    dataset = torch.utils.data.TensorDataset(x, y)
    return torch.utils.data.DataLoader(dataset, batch_size, shuffle)

train_loader = to_loader(X_train, y_train, batch_size, shuffle=True)
test_loader = to_loader(X_test, y_test, batch_size, shuffle=False)
```

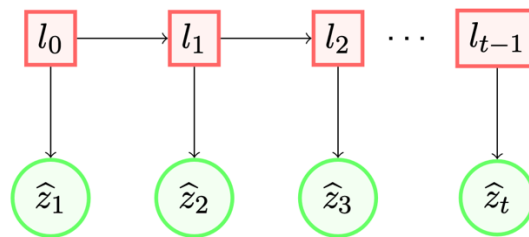
Outline

1. Introduction
2. Data Processing
3. **Background of Practice Model**
4. Practice
5. (Optional) Encoder-decoder structure

Recap: State Space Models

- **State Space Model (SSM)**

- **Input:** last state \mathbf{l}_{t-1}
- **Update:** $\mathbf{l}_t = \mathbf{F}_t \mathbf{l}_{t-1} + \mathbf{g}_t \epsilon_t$
 - White noise ϵ_t , parameters $\{\mathbf{F}_t, \mathbf{g}_t\}$
- **Output:** $z_t = \mathbf{a}_t^\top \mathbf{l}_{t-1} + \epsilon_t$
 - Parameter \mathbf{a}_t



Measurements $z_t = \mathbf{a}_t^\top \mathbf{l}_{t-1} + \epsilon_t, \quad \epsilon_t \sim N(0, \sigma^2)$

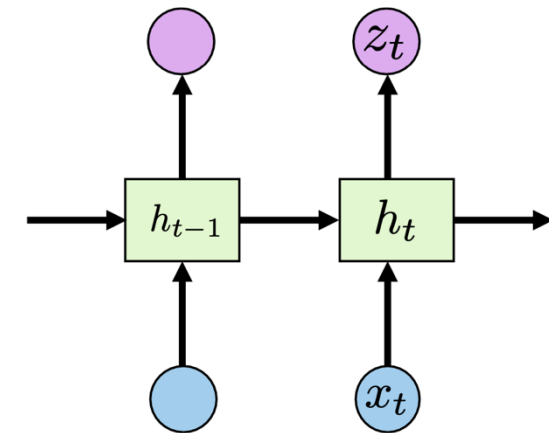
State transition $\mathbf{l}_t = \mathbf{F}_t \mathbf{l}_{t-1} + \mathbf{g}_t \epsilon_t, \quad \mathbf{l}_0 \sim N(\boldsymbol{\mu}_0, \text{diag}(\sigma_0^2)).$

Recap: Recurrent Neural Networks (RNN)

- **RNN(Recurrent Neural Network)**

- **Input:** last state h_{t-1} , current feature x_t
- **Update:** $h_t = \sigma(\theta_0 h_{t-1} + \theta_1 x_t)$
 - Activation function σ , learnable parameters $\{\theta_0, \theta_1\}$
- **Output:** $z_t = \sigma(\theta h_t)$
 - Learnable parameter θ

RECURRENT NEURAL NETWORK



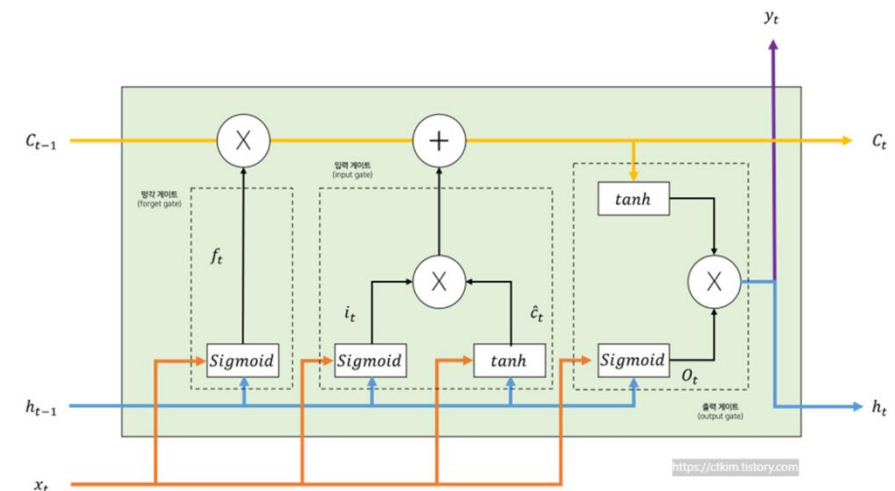
$$h_t = \sigma(\theta_0 h_{t-1} + \theta_1 x_t)$$

$$z_t = \sigma(\theta h_t)$$

Recap: Long Short-Term Memory (LSTM)

- **LSTM(Long Short Term Memory)**

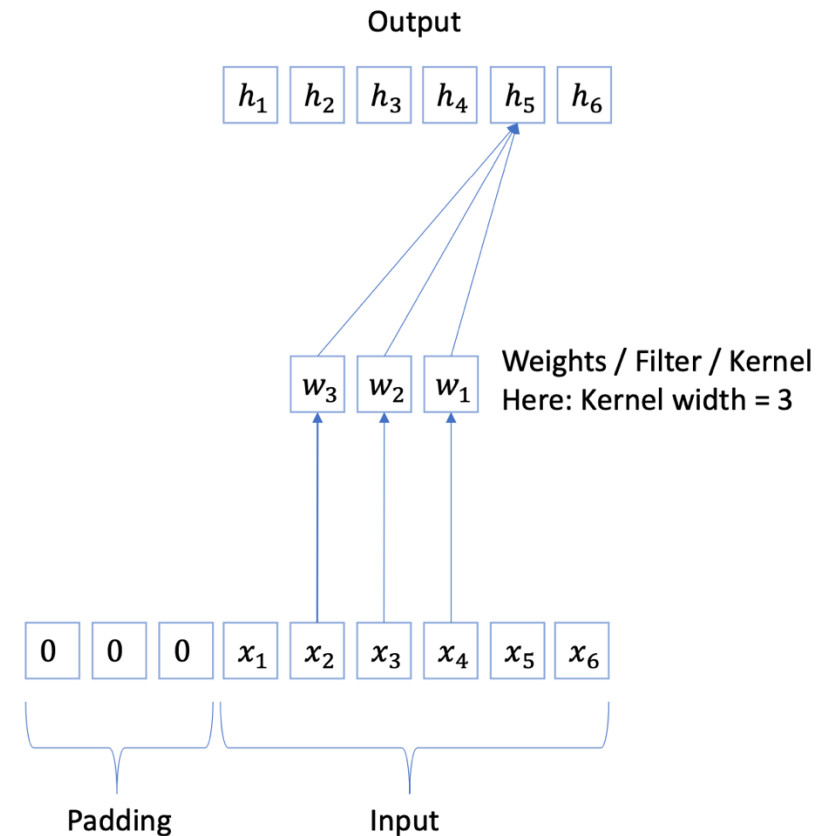
- **Input:** last short-term state h_{t-1} , last long-term state C_{t-1} , current feature x_t
- **Update:** $C_t = f_t \odot C_{t-1} + i_t \odot \hat{C}_t$, $h_t = o_t \odot \tanh(C_t)$
 - Forget Gate: $f_t = \sigma(\mathbf{W}_{hf}h_{t-1} + \mathbf{W}_{xf}x_t)$
 - Input Gate: $i_t = \sigma(\mathbf{W}_{hi}h_{t-1} + \mathbf{W}_{xi}x_t)$, $\hat{C}_t = \tanh(\mathbf{W}_{hc}h_{t-1} + \mathbf{W}_{xc}x_t)$
 - Output Gate: $o_t = \sigma(\mathbf{W}_{ho}h_{t-1} + \mathbf{W}_{xo}x_t)$
- **Output:** $z_t = h_t$



Recap: Convolutional Neural Networks (CNN)

- **CNN(Convolutional Neural Networks)**

- **Input:** previous features $\{x_i\}_{i=t-D}^{t-1}$
- **Output :** $h_t = \sum_{d=1}^D w_d x_{t-d}$
 - Kernel $W = [w_1, \dots, w_D]$
- After convolution layer, passing more layers like pooling, flatten, fully-connected layer



Outline

1. Introduction
2. Data Processing
3. Background of Practice Model
- 4. Practice**
5. (Optional) Encoder-decoder structure

LSTM Model Architecture

- **Objective:** Build an LSTM model for time series forecasting
- **TODO:**
 - Sequential data passes through the LSTM layer.
 - The fully connected layer maps the last LSTM output to a single value.

```
class LSTMModel(nn.Module):
    def __init__(self, input_size, hidden_size, num_layers):
        super(LSTMModel, self).__init__()
        self.lstm = nn.LSTM(input_size, hidden_size, num_layers, batch_first=True)
        # (lstm): LSTM(1, 64, num_layers=2, batch_first=True)

        self.linear = nn.Linear(hidden_size, 1)
        # (linear): Linear(in_features=64, out_features=1, bias=True)

    def forward(self, x):
        out, _ = self.lstm(x)
        return self.linear(out)

model = LSTMModel(input_size, hidden_size, num_layers).to(device)
print(model)
```

Train

- **Objective:** Implement a training loop
- **TODO:**
 - Compute predictions
 - Calculate the train loss
 - Perform optimization

```
# train
model.train()
for batch_x, batch_y in train_loader:
    # (1) Move input and target tensors to the device (e.g., GPU)
    batch_x, batch_y =

    # (2) Pass the input (batch_x) through the model
    #     The model outputs shape [batch_size, sequence_length, 1]
    #     Take the prediction at the last timestep (index -1) and feature index 0
    #     → model(batch_x)[: , -1, 0]
    pred =

    # (3) Compute the loss between pred and batch_y by using loss_fn
    loss =

    # (4) Clear previous gradients to avoid accumulation
    optimizer.zero_grad_()

    # (5) Perform backpropagation to compute gradients
    loss.backward()

    # (6) Update the model parameters using the optimizer
    optimizer.step()

    total_train_loss += loss.item()

avg_loss = total_train_loss / len(train_loader)
train_hist.append(avg_loss)
```

Test

- **Objective:** Implement evaluation
- **TODO:**
 - Compute predictions
 - Calculate the test loss

```
# evaluate
model.eval()
with torch.no_grad():
    for test_x, test_y in test_loader:
        # TODO
        # (1) Move input and target tensors to the device (e.g., GPU)
        test_x, test_y =

        # (2) Pass the input (test_x) through the model
        #     The model outputs shape [batch_size, sequence_length, 1]
        #     Take the prediction at the last timestep (index -1) and feature index 0
        #     → model(test_x)[:,-1, 0]
        test_pred =

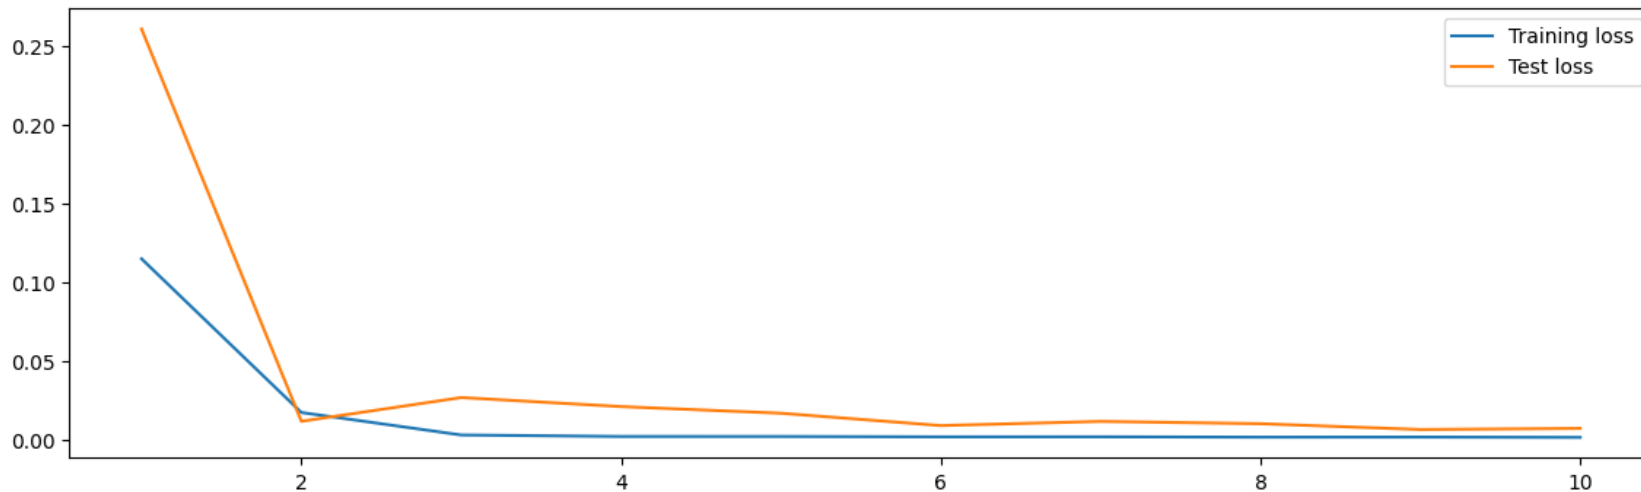
        # (3) Compute the loss between test_pred and test_y by using loss_fn
        test_loss =

        total_test_loss += test_loss.item()

avg_test_loss = total_test_loss / len(test_loader)
test_hist.append(avg_test_loss)
```

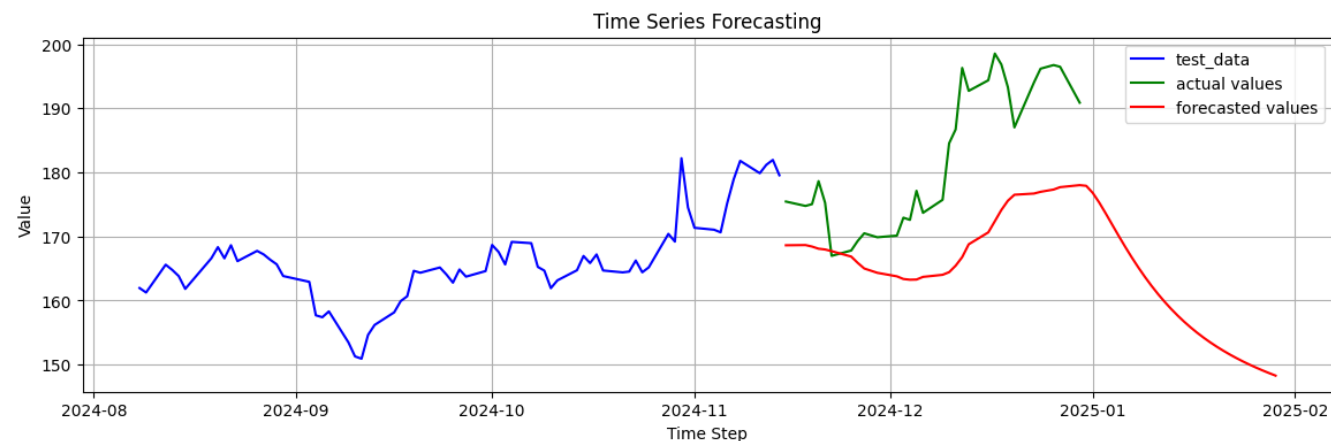
Loss Visualization

- **Objective:** Compare training loss and test loss over the epochs
- **Interpretation:**
 - **Training loss:** Stabilizes after a few epochs, suggesting convergence.
 - **Test loss:** Follows a similar trend as training loss but remains slightly higher



Forecasting Visualization

- **Objective:** Compare actual values with the model's forecasted values
- **Interpretation:**
 - **Test data:** unseen test data that the model is being evaluated on
 - **Actual values:** ground truth values used to validate the model's forecasts
 - **Forecasted values:** the model's predicted values for the future



Model Evaluation

- **Objective:** Evaluate the performance of the LSTM model
 - Root Mean Squared Error (RMSE)
 - Mean Absolute Percentage Error (MAPE)
- **TODO:**
 - Measure RMSE and MAPE

```
def test(model, X_test, y_test):  
    model.eval()  
    with torch.no_grad():  
        test_predictions = []  
        for batch_X_test in X_test:  
            batch_X_test = batch_X_test.to(device).unsqueeze(0)  
            test_predictions.append(model(batch_X_test) \  
                                   .cpu().numpy().flatten()[0])  
  
    test_predictions = np.array(test_predictions)  
    y_test = y_test.cpu().numpy()  
  
    rmse = root_mean_squared_error(, ) # TODO  
    mape = mean_absolute_percentage_error(, ) # TODO  
  
    return rmse, mape
```

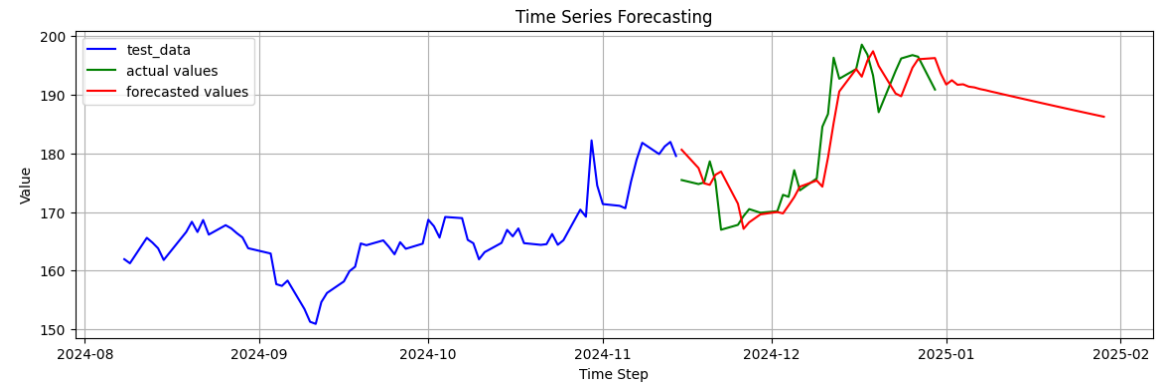
Other Models

- **Objective:** Compare the performance of Conv1D and RNN models
 - Conv1D focuses on capturing local temporal patterns
 - RNNs are designed for sequential dependencies over time

```
class Conv1DModel(nn.Module):
    def __init__(self, input_size, hidden_size):
        super(Conv1DModel, self).__init__()
        # TODO
        self.conv1d = nn.Conv1d(in_channels=, out_channels=, \
                                   kernel_size=2, stride=1)

        self.fc = nn.Linear(, 1)

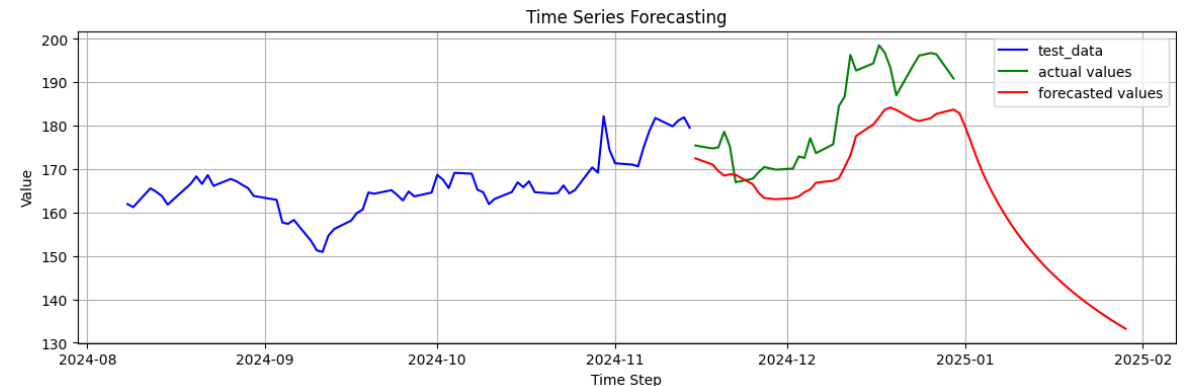
    def forward(self, x):
        x = x.transpose(1, 2)
        x = self.conv1d(x)
        x = x.transpose(1, 2)
        return self.fc(x)
```



Other Models

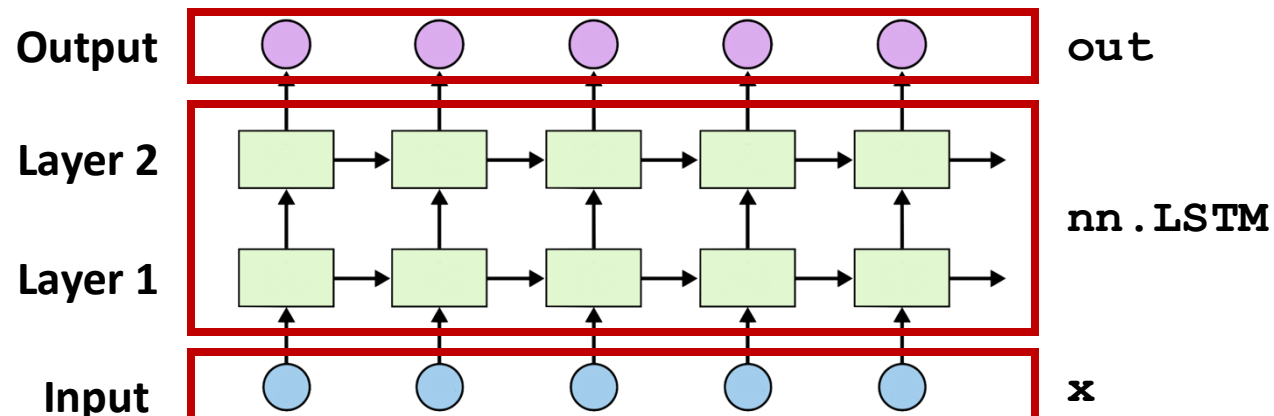
- **Objective:** Compare the performance of Conv1D and RNN models
 - Conv1D focuses on capturing local temporal patterns
 - RNNs are designed for sequential dependencies over time

```
class RNNModel(nn.Module):  
    def __init__(self, input_size, hidden_size, num_layers):  
        super(RNNModel, self).__init__()  
        # TODO  
        self.rnn = nn.RNN(, , , batch_first=True)  
        self.fc = nn.Linear(, 1)  
  
    def forward(self, x):  
        # TODO
```



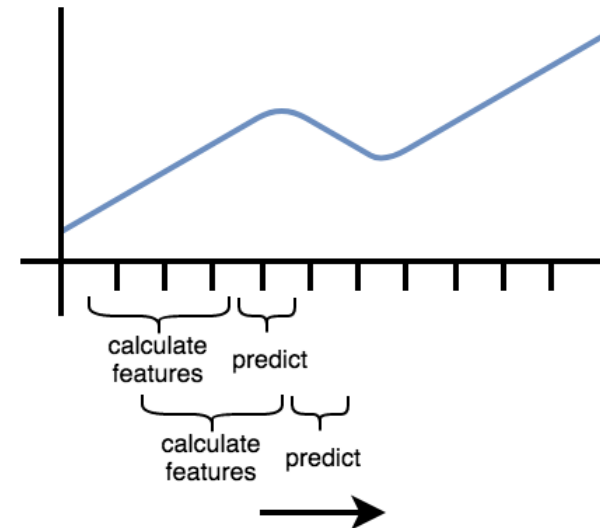
Code Explanation: LSTM Architecture

- `out, (hn, cn) = nn.LSTM(x)`
 - `x`: [batch_size, sequence_length, num_features] # [32, 50, 1]
 - `out`: [batch_size, sequence_length, hidden_size] # [32, 50, 64]
 - `hn`: [num_layers, batch_size, hidden_size] # [2, 32, 64]
 - `cn`: [num_layers, batch_size, hidden_size] # [2, 32, 64]
 - Only the output features from the final LSTM layer are typically required.



Code Explanation: Rolling Forecast

- `input_data = np.roll(input_data, shift=-1)`
 - For example, `[10, 20, 30, 40, 50] → [20, 30, 40, 50, 10]`
 - Then replace the last value (10).
 - Used to shift the input window for autoregressive forecasting.
 - If ground-truth future value is available, insert it at the end of the input window.
 - Otherwise, insert the model's predicted value.

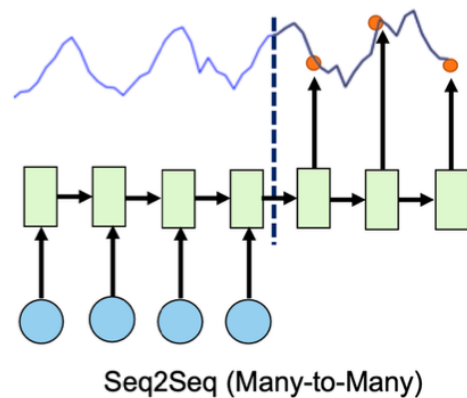


Outline

1. Introduction
2. Data processing
3. Background of practice model
4. Practice
5. **(Optional) Encoder-decoder structure**

Recap: Encoder-Decoder Structure

- **Encoder:** Captures the sequence of temporal dependencies from the past observations.
- **Decoder:** Uses the encoded information from the encoder to produce the future predictions.



$$f_{\text{encoder}} : \{z_1, \dots, z_{T_e}\} \mapsto \mathbf{h}_{T_e}$$

$$f_{\text{decoder}} : \mathbf{h}_{T_e} \mapsto \{z_{T_e+1}, \dots, z_{T_e+T_d}\}$$

Multi-Step Forecasting

- **Forecasting multi-step:** The model is trained on sequences of past observations and sequences of future values.

```
sequence_length = 50
target_len = 10

def create_enc_dec_sequences(data):
    features, labels = [], []
    for i in range(len(data) - sequence_length - target_len):
        features.append(data[i:i + sequence_length])
        labels.append(data[i + sequence_length : \
                        i + sequence_length + target_len])

    features = np.array(features, dtype=np.float32)
    labels = np.array(labels, dtype=np.float32)

    features = torch.tensor(features, dtype=torch.float32)
    labels = torch.tensor(labels, dtype=torch.float32)
    return features, labels
```

Functions for Training

- In multi-step forecasting, the format for creating training pairs changes.
 - Unlike one-step forecasting, **multi-step forecasting requires output sequences** that predict several future time steps at once.

```
def train_(model, train_loader, test_loader):
```

```
def plot_forecasting_(model, X_test, y_test):
```

```
def test_(model, X_test, y_test):
```

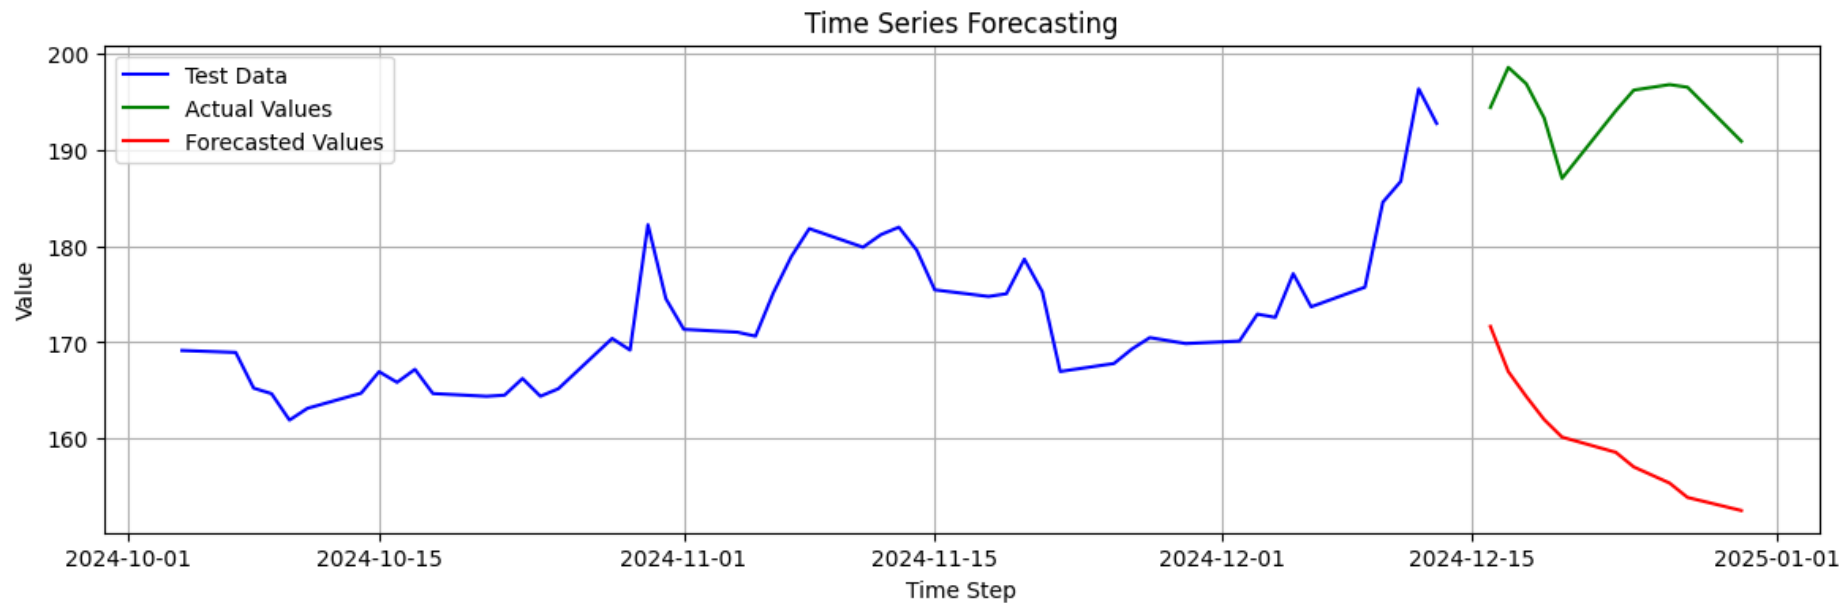
RNN-RNN Model Architecture

- **Objective:** Build an RNN-RNN architecture
- **TODO:**
 - Implement the `__init__` method
 - Complete the forward function

```
class EncoderRNN(nn.Module):  
    def __init__(self, input_size, hidden_size, num_layers):  
        # TODO  
  
    def forward(self, x):  
        # TODO  
  
class DecoderRNN(nn.Module):  
    def __init__(self, input_size, hidden_size, num_layers):  
        # TODO  
  
    def forward(self, x, h):  
        # TODO
```

Forecasting Visualization

- **Objective:** Display the test data, actual values, and the forecasted values to evaluate the model's performance.



References

- [NIPS '14] Sequence to Sequence Learning with Neural Networks
 - <https://arxiv.org/abs/1409.3215>